

Intelligent Agents

Chapter 2

TB Artificial Intelligence




Slides from AIMA — <http://aima.cs.berkeley.edu>

Outline

- ▶ Agents and environments
- ▶ Rationality
- ▶ PEAS (Performance measure, Environment, Actuators, Sensors)
- ▶ Environment types
- ▶ Agent types

Agents and environments



► **Agents** include humans, robots, softbots, thermostats, etc.

► The **agent function** maps from percept histories to actions:

$$f : \mathcal{P}^* \rightarrow \mathcal{A}$$

► The **agent program** runs on the physical **architecture** to produce f

Vacuum-cleaner world

- ▶ Percepts: location and contents, e.g., $[A, \textit{Dirty}]$
- ▶ Actions: *Left*, *Right*, *Suck*, *NoOp*

A vacuum-cleaner agent

Percept sequence	Action
<i>[A, Clean]</i>	<i>Right</i>
<i>[A, Dirty]</i>	<i>Suck</i>
<i>[B, Clean]</i>	<i>Left</i>
<i>[B, Dirty]</i>	<i>Suck</i>
<i>[A, Clean], [A, Clean]</i>	<i>Right</i>
<i>[A, Clean], [A, Dirty]</i>	<i>Suck</i>
<i>⋮</i>	<i>⋮</i>

function REFLEX-VACUUM-AGENT([*location,status*]) **returns** an action

if *status* = *Dirty* **then return** *Suck*
else if *location* = *A* **then return** *Right*
else if *location* = *B* **then return** *Left*

- ▶ What is the **right** function?
- ▶ Can it be implemented in a small agent program?

Rationality

- ▶ Fixed **performance measure** evaluates the **environment sequence**
 - ▶ one point per square cleaned up in time T ?
 - ▶ one point per clean square per time step, minus one per move?
 - ▶ penalize for $> k$ dirty squares?
 - ▶ A **rational agent** chooses whichever action maximizes the **expected** value of the performance measure **given the percept sequence to date**
 - ▶ Rational \neq omniscient
 - ▶ percepts may not supply all relevant information
 - ▶ Rational \neq clairvoyant
 - ▶ action outcomes may not be as expected
- Hence, rational \neq successful
- ▶ Rational \implies exploration, learning, autonomy

PEAS

To design a rational agent, we must specify the **task environment**

Consider, e.g., the task of designing an automated taxi:

- ▶ Performance measure??
- ▶ Environment??
- ▶ Actuators??
- ▶ Sensors??

To design a rational agent, we must specify the **task environment**

Consider, e.g., the task of designing an automated taxi:

- ▶ Performance measure?? safety, destination, profits, legality, comfort, ...
- ▶ Environment?? streets/freeways, traffic, pedestrians, weather, ...
- ▶ Actuators?? steering, accelerator, brake, horn, speaker/display, ...
- ▶ Sensors?? video, accelerometers, gauges, engine sensors, keyboard, GPS, ...

Internet shopping agent

- ▶ Performance measure??
- ▶ Environment??
- ▶ Actuators??
- ▶ Sensors??

Internet shopping agent

- ▶ Performance measure?? price, quality, appropriateness, efficiency
- ▶ Environment?? current and future WWW sites, vendors, shippers
- ▶ Actuators?? display to user, follow URL, fill in form
- ▶ Sensors?? HTML pages (text, graphics, scripts)

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u> <u>Deterministic??</u> <u>Episodic??</u> <u>Static??</u> <u>Discrete??</u> <u>Single-agent??</u>				

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic??</u>				
<u>Episodic??</u>				
<u>Static??</u>				
<u>Discrete??</u>				
<u>Single-agent??</u>				

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic??</u>	Yes	No	Partly	No
<u>Episodic??</u>				
<u>Static??</u>				
<u>Discrete??</u>				
<u>Single-agent??</u>				

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic??</u>	Yes	No	Partly	No
<u>Episodic??</u>	No	No	No	No
<u>Static??</u>				
<u>Discrete??</u>				
<u>Single-agent??</u>				

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic??</u>	Yes	No	Partly	No
<u>Episodic??</u>	No	No	No	No
<u>Static??</u>	Yes	Semi	Semi	No
<u>Discrete??</u>				
<u>Single-agent??</u>				

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic??</u>	Yes	No	Partly	No
<u>Episodic??</u>	No	No	No	No
<u>Static??</u>	Yes	Semi	Semi	No
<u>Discrete??</u>	Yes	Yes	Yes	No
<u>Single-agent??</u>				

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic??</u>	Yes	No	Partly	No
<u>Episodic??</u>	No	No	No	No
<u>Static??</u>	Yes	Semi	Semi	No
<u>Discrete??</u>	Yes	Yes	Yes	No
<u>Single-agent??</u>	Yes	No	Yes (except auctions)	No

Environment types

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Observable??</u>	Yes	Yes	No	No
<u>Deterministic??</u>	Yes	No	Partly	No
<u>Episodic??</u>	No	No	No	No
<u>Static??</u>	Yes	Semi	Semi	No
<u>Discrete??</u>	Yes	Yes	Yes	No
<u>Single-agent??</u>	Yes	No	Yes (except auctions)	No

The environment type largely determines the agent design

The real world is (of course) partially observable, stochastic, sequential, dynamic, continuous, multi-agent

Agent types

Four basic types in order of increasing generality:

- ▶ simple reflex agents
- ▶ reflex agents with state
- ▶ goal-based agents
- ▶ utility-based agents

All these can be turned into learning agents

Simple reflex agents

Example

function REFLEX-VACUUM-AGENT([*location,status*]) **returns** an action

if *status* = *Dirty* **then return** *Suck*
else if *location* = *A* **then return** *Right*
else if *location* = *B* **then return** *Left*

```
(setq joe (make-agent :name 'joe :body (make-agent-body)
                      :program (make-reflex-vacuum-agent-program)))
```

```
(defun make-reflex-vacuum-agent-program ()
  #'(lambda (percept)
      (let ((location (first percept)) (status (second percept)))
        (cond ((eq status 'dirty) 'Suck)
              ((eq location 'A) 'Right)
              ((eq location 'B) 'Left))))))
```

Reflex agents with state

Example

function REFLEX-VACUUM-AGENT([*location,status*]) **returns** an action
static: *last_A, last_B*, numbers, initially ∞
if *status* = *Dirty* **then** ...

```
(defun make-reflex-vacuum-agent-with-state-program ()
  (let ((last-A infinity) (last-B infinity))
    #'(lambda (percept)
      (let ((location (first percept)) (status (second percept)))
        (incf last-A) (incf last-B)
        (cond
         ((eq status 'dirty)
          (if (eq location 'A) (setq last-A 0) (setq last-B 0))
          'Suck)
         ((eq location 'A) (if (> last-B 3) 'Right 'NoOp))
         ((eq location 'B) (if (> last-A 3) 'Left 'NoOp)))))))
```

Goal-based agents

Utility-based agents

Learning agents

Summary

- ▶ Agents interact with environments through actuators and sensors
- ▶ The agent function describes what the agent does in all circumstances
- ▶ The performance measure evaluates the environment sequence
- ▶ A perfectly rational agent maximizes expected performance
- ▶ Agent programs implement (some) agent functions
- ▶ PEAS descriptions define task environments
- ▶ Environments are categorized along several dimensions:
observable? deterministic? episodic? static? discrete? single-agent?
- ▶ Several basic agent architectures exist:
reflex, reflex with state, goal-based, utility-based