

# Distributed Constraint Optimization for the Internet-of-Things

Pierre Rust Gauthier Picard

Orange Labs  
MINES Saint-Étienne  
LaHC UMR CNRS 5516

[pierre.rust@orange.com](mailto:pierre.rust@orange.com)



# Internet-of-Things (IoT) and its Control

- Huge (marketing ?) trend today
- 25 billion of connected objects in 2020 ? (Gartner)
- Hardware and communication is cheaper and cheaper
- Constrained devices
  - ▶ limited cpu and memory resources
  - ▶ limited communication capabilities
- Connected things' actions should be **coordinated**
- Current approach: centralizing decisions
  - ▶ Communications
  - ▶ Resilience
  - ▶ Scalability
  - ▶ Cost



# Distributed Coordination and Decision Making

Autonomous and spontaneous

Coordinating objects to achieve objectives

- Coordination
  - ▶ Decentralized
  - ▶ Spontaneous
  - ▶ Autonomous
- No central point
- Self-adaption to environmental changes
- Self-repair in case of one component failure



# About decisions

$x_i$  ?

## About decisions

$x_i$  ?

s.t. "I'm happy with  $x_i$ "

## About decisions

$x_i$  ?

s.t. "I'm happy with  $x_i$ "

$x_j$  ?

s.t. "agent  $i$  is fine with  $x_j$ "

## About decisions

$x_i$  ?

s.t. "I'm happy with  $x_i$ "

$x_j$  ?

s.t. "agent  $i$  is fine with  $x_j$ "

How can agents autonomously make their decisions  
in a coordinated way, without external control ?

## About decisions

$x_i$  ?

s.t. "I'm happy with  $x_i$ "

$x_j$  ?

s.t. "agent  $i$  is fine with  $x_j$ "

How can agents autonomously make their decisions  
in a coordinated way, without external control ?

⇒ Decentralized decision making

## About decisions

$x_i$  ?

s.t. "I'm happy with  $x_i$ "

$x_j$  ?

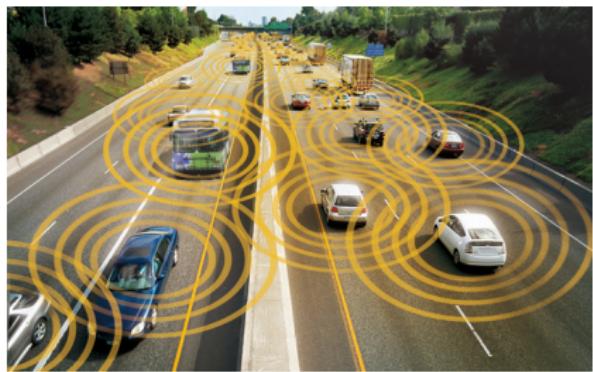
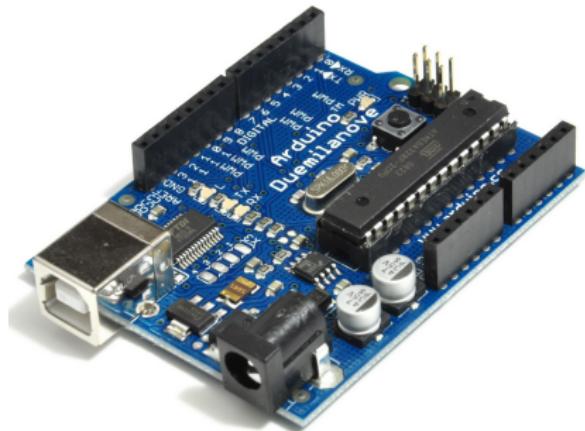
s.t. "agent  $i$  is fine with  $x_j$ "

How can agents autonomously make their decisions  
in a coordinated way, without external control ?

⇒ Decentralized decision making

- Agents have to coordinate to perform best actions
- Agents form a team → best actions for the team

# Application Domains



# Menu

DCOP Framework

Focus on Some Solution Methods

Hands on PyDCOP I

Focus on Smart Environment Configuration Problems

Distributing Computations

Hands on PyDCOP II

Dynamic DCOPs

Conclusion

# Menu

DCOP Framework

Focus on Some Solution Methods

Hands on PyDCOP I

Focus on Smart Environment Configuration Problems

Distributing Computations

Hands on PyDCOP II

Dynamic DCOPs

Conclusion

# DCOP<sup>1</sup>

## Distributed Constraints Optimization Problem

### Definition (DCOP)

A DCOP is a tuple  $\langle \mathcal{A}, \mathcal{X}, \mathcal{D}, \mathcal{C}, \mu \rangle$ , where:

- $\mathcal{A} = \{a_1, \dots, a_{|\mathcal{A}|}\}$  is a set of agents
- $\mathcal{X} = \{x_1, \dots, x_n\}$  are variables
- $\mathcal{D} = \{\mathcal{D}_{x_1}, \dots, \mathcal{D}_{x_n}\}$  is a set of finite domains, for the  $x_i$  variables
- $\mathcal{C} = \{f_1, \dots, f_m\}$  is a set of soft constraints, where each  $c_i$  defines a cost  $\in \mathbb{R} \cup \{\infty\}$  for each combination of assignments to a subset of variables
- $\mu$  is a function mapping variables to their associated agent

### Definition (Solution)

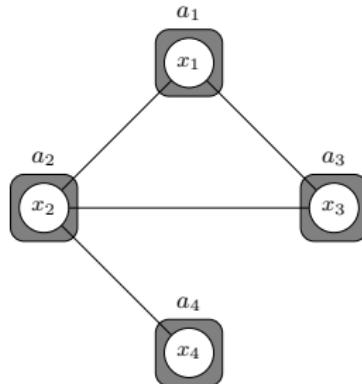
A *solution* to the DCOP is an assignment  $\mathcal{A}$  to all variables that minimizes  $\sum_i f_i$

---

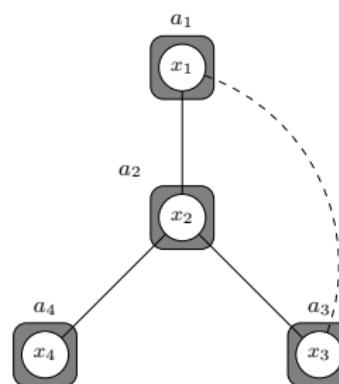
<sup>1</sup>Some contents taken from OPTMAS 2011 and OPTMAS-DCR 2014

# DCOP

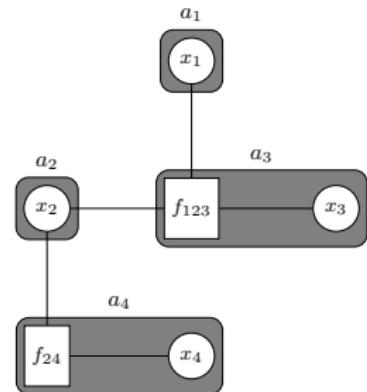
## Example and Graphical Representation



(a)



(b)



(c)

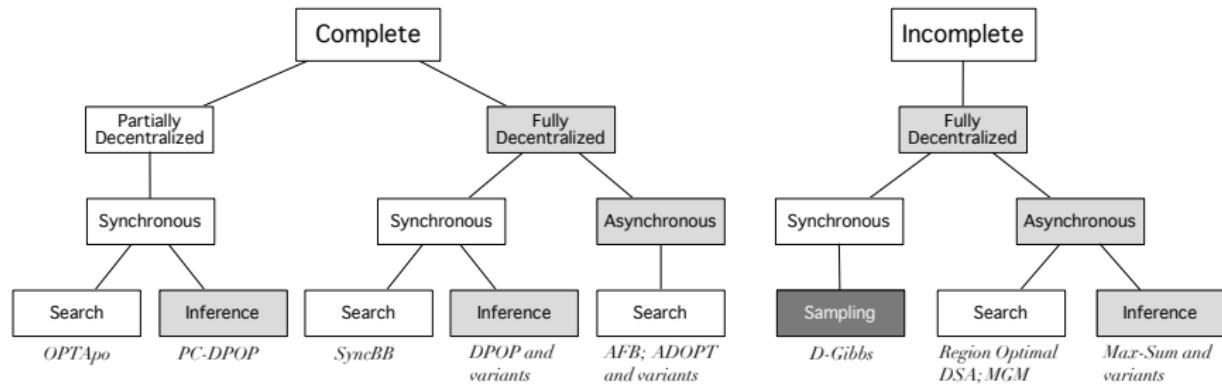
## Objective Function

$$F(\mathcal{A}) = \sum_{x_i, x_j \in X} f_{ij} \quad \text{where } f_{ij} = (\mathbf{x}_i + \mathbf{x}_j + 1) \bmod 3$$

In figure (a):

- $F(\{(x_1, 0), (x_2, 0), (x_3, 0), (x_4, 0)\}) = 4$
- $F(\{(x_1, 1), (x_2, 1), (x_3, 1), (x_4, 1)\}) = 0$

# DCOP Algorithms



(FIORETTI et al., 2018)

# Menu

DCOP Framework

Focus on Some Solution Methods

DPOP

Max-Sum

DSA

MGM

Hands on PyDCOP I

Focus on Smart Environment Configuration Problems

Distributing Computations

Hands on PyDCOP II

Dynamic DCOPs

Conclusion

# Distributed Pseudotree Optimization Procedure (DPOP)

(PETCU and FALTING, 2005)

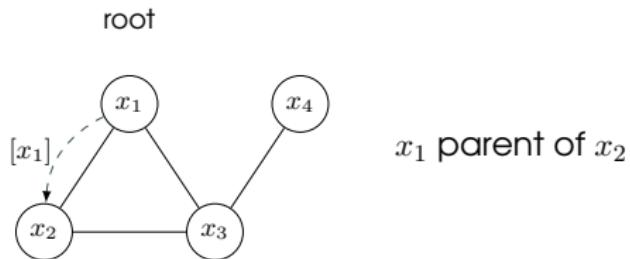
3-phase distributed algorithm

PHASES	MESSAGES
1. DFS Tree construction	token passing
2. Utility phase: from leaves to root	<b>util</b> (child → parent, constraint table (-child))
3. Value phase: from root to leaves	<b>value</b> (parent → children ∪ pseudochildren, parent value)

## DFS Tree Phase

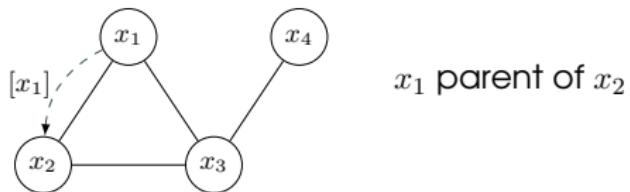
- **Distributed DFS graph traversal:** token, ID,  $\text{neighbors}(X)$ 
  1.  $X$  owns the token: adds its own  $ID$  and sends it in turn to each of its neighbors, which become children
  2.  $Y$  receives the token from  $X$ : it marks  $X$  as visited. First time  $Y$  receives the token then  $\text{parent}(Y) = X$ . Other IDs in token which are also  $\text{neighbors}(Y)$  are **pseudoparent**. If  $Y$  receives token from neighbor  $W$  to which it was never sent,  $W$  is pseudochild.
  3. When all  $\text{neighbors}(X)$  visited,  $X$  removes its ID from token and sends it to  $\text{parent}(X)$ .
- A node is selected as root, which starts
- When all neighbors of root are visited, the DFS traversal ends

## DFS Tree Phase: Example

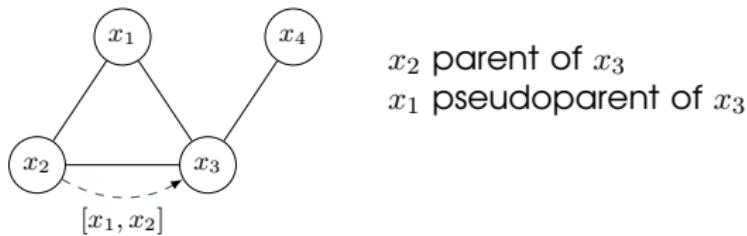


## DFS Tree Phase: Example

root



$x_1$  parent of  $x_2$

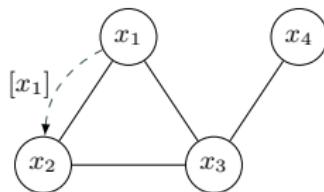


$x_2$  parent of  $x_3$

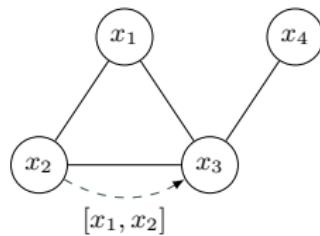
$x_1$  pseudoparent of  $x_3$

## DFS Tree Phase: Example

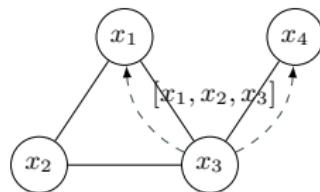
root



$x_1$  parent of  $x_2$



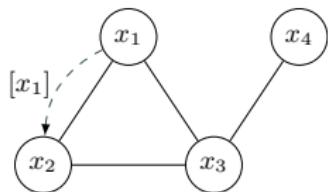
$x_2$  parent of  $x_3$   
 $x_1$  pseudoparent of  $x_3$



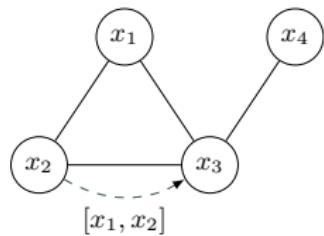
$x_3$  parent of  $x_4$   
 $x_3$  pseudoparent of  $x_1$

## DFS Tree Phase: Example

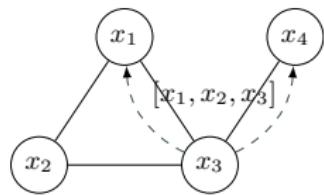
root



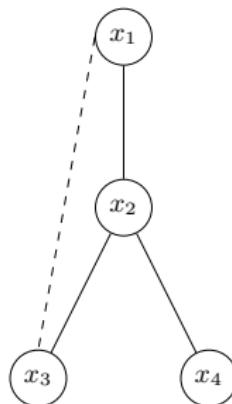
$x_1$  parent of  $x_2$



$x_2$  parent of  $x_3$   
 $x_1$  pseudoparent of  $x_3$



$x_3$  parent of  $x_4$   
 $x_3$  pseudoparent of  $x_1$



## Util Phase

Agent  $X$ :

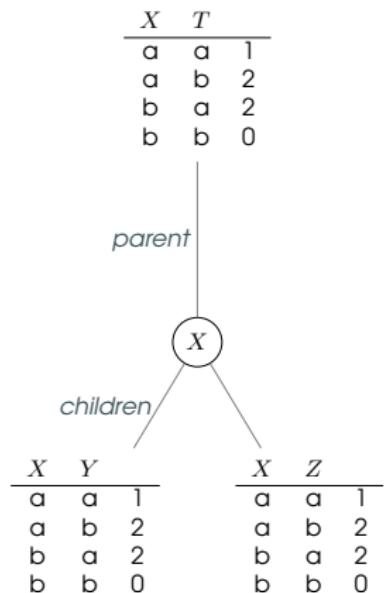
- receives from each child  $Y_i$  a cost function:  $C(Y_i)$
- combines (adds, joins) all these cost functions with the cost functions with  $\text{parent}(X)$  and  $\text{pseudoparents}(X)$
- projects  $X$  out of the resulting cost function, and sends it to  $\text{parent}(X)$

From the leaves to the root

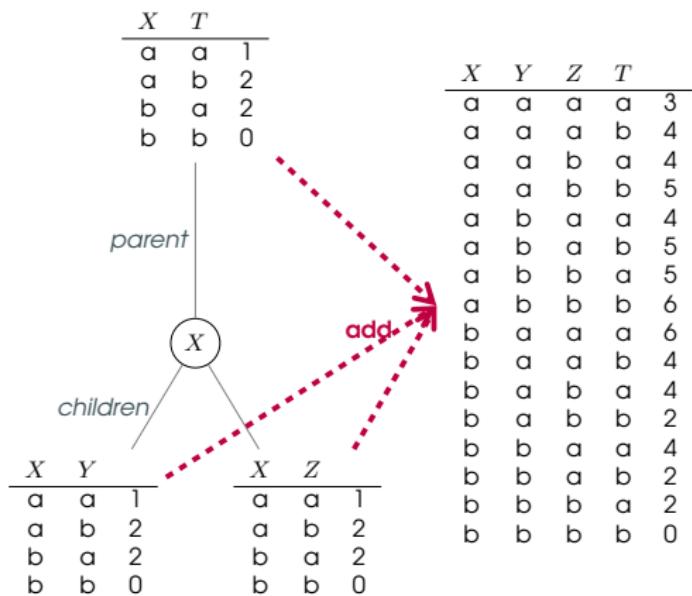
## Util Phase: Example



## Util Phase: Example



## Util Phase: Example



All value combinations  
Costs are the sum of applicable costs

## Util Phase: Example

$X$	$T$	
a	a	
a	b	
b	a	
b	b	0

parent

children

$X$	$Y$	$T$
a	a	1
a	b	2
b	a	2
b	b	0

$X$	$Z$	$T$
a	a	1
a	b	2
b	a	2
b	b	0

$X$	$Y$	$Z$	$T$
a	a	a	3
a	a	b	4
a	a	a	4
a	a	b	5
a	b	a	4
a	b	a	5
a	b	b	5
a	b	a	5
a	b	b	6
b	a	a	6
b	a	b	4
b	a	b	4
b	a	b	2
b	b	a	4
b	b	b	2
b	b	a	2
b	b	b	0

All value combinations  
Costs are the sum of applicable costs

$X$	$Y$	$Z$	$T$
c	a	a	3
c	a	b	4
c	a	a	4
c	a	b	5
c	b	a	4
c	b	a	5
c	b	b	6
t	a	a	6
b	a	a	4
b	a	b	4
b	a	b	2
b	b	a	4
b	b	b	2
b	b	a	2
b	b	b	0

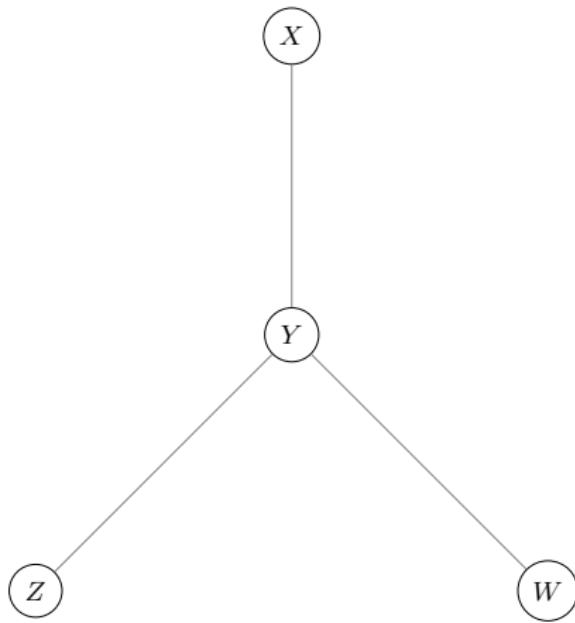
Remove  $X$   
Remove duplicates  
Keep the min cost

## Value Phase

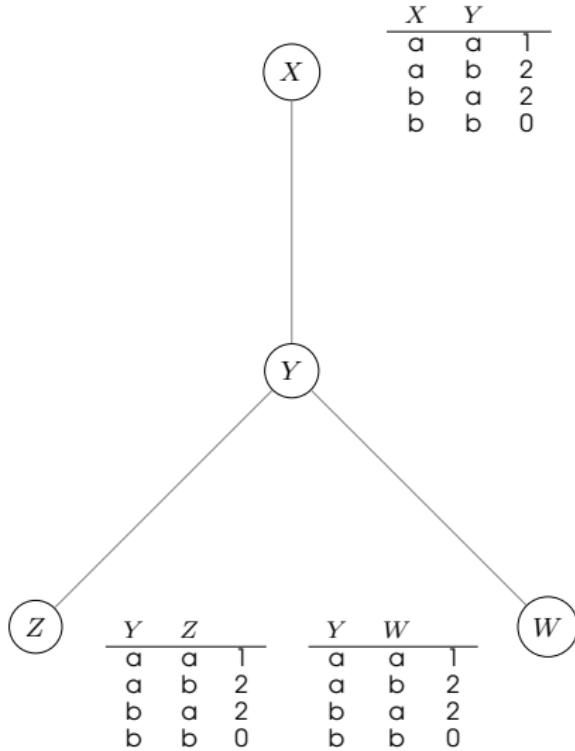
1. The root finds the **value that minimizes the received cost function** in the util phase, and informs its descendants (children  $\cup$  pseudochildren)
2. Each agent **waits to receive** the value of its parent / pseudoparents
3. Keeping fixed the value of parent/pseudoparents, finds **the value that minimizes the received cost function** in the Util phase
4. Informs of this value to its children/pseudochildren

This process starts at the root and ends at the leaves

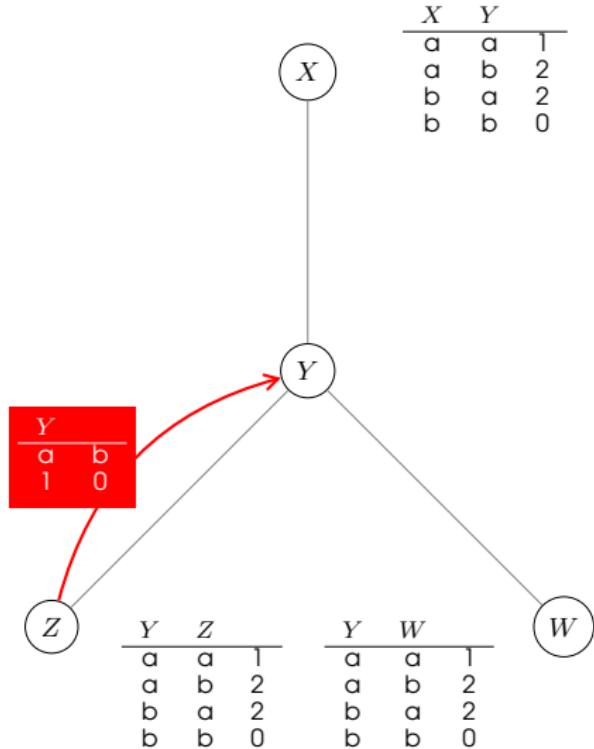
## DTREE : DPOP for DCOPs without backedges



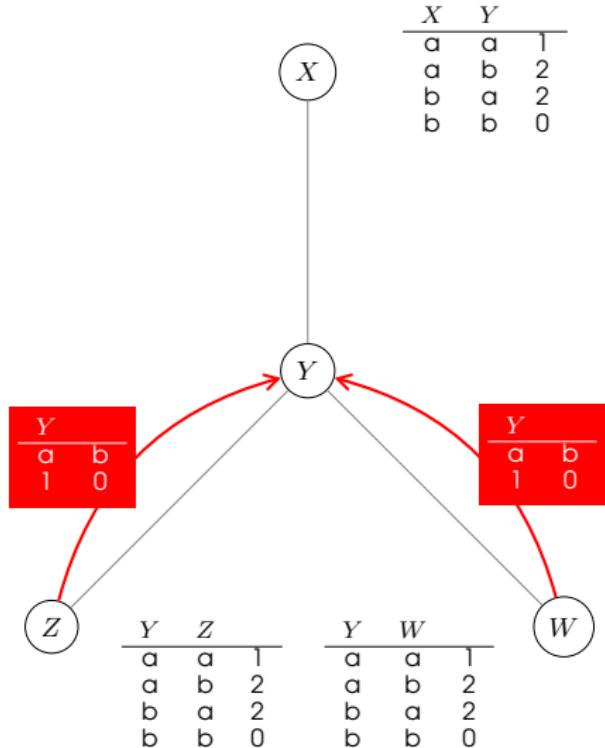
## DTREE : DPOP for DCOPs without backedges



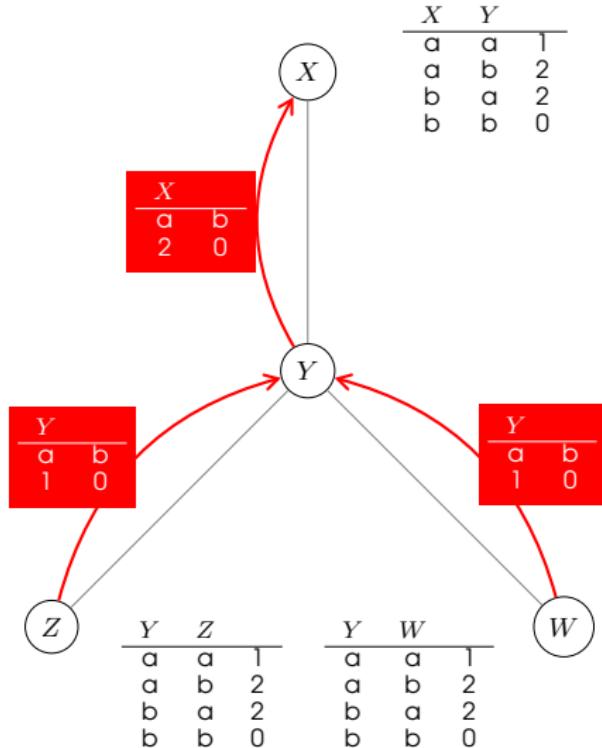
## DTREE : DPOP for DCOPs without backedges



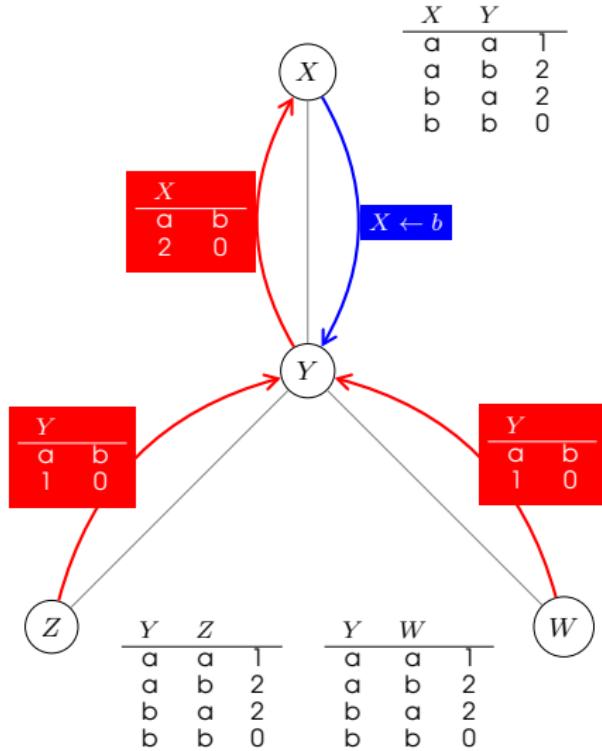
## DTREE : DPOP for DCOPs without backedges



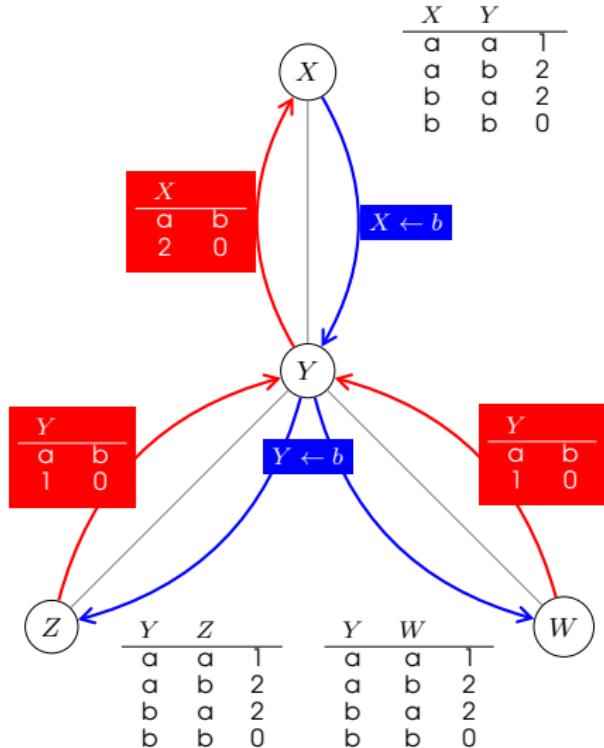
## DTREE : DPOP for DCOPs without backedges



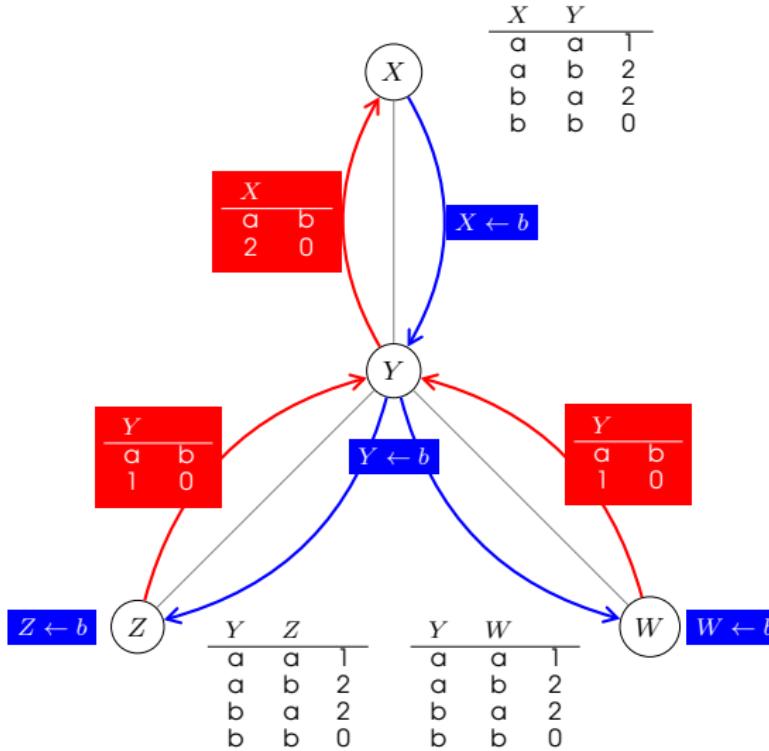
## DTREE : DPOP for DCOPs without backedges



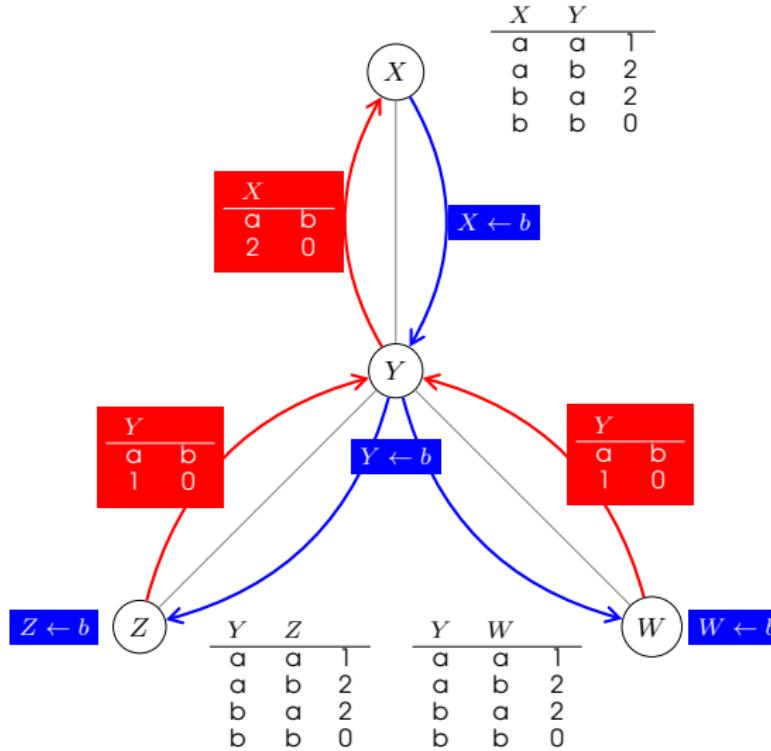
## DTREE : DPOP for DCOPs without backedges



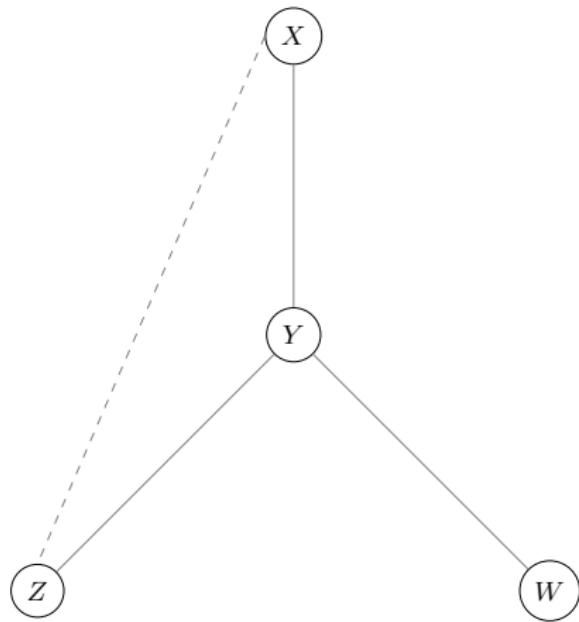
## DTREE : DPOP for DCOPs without backedges



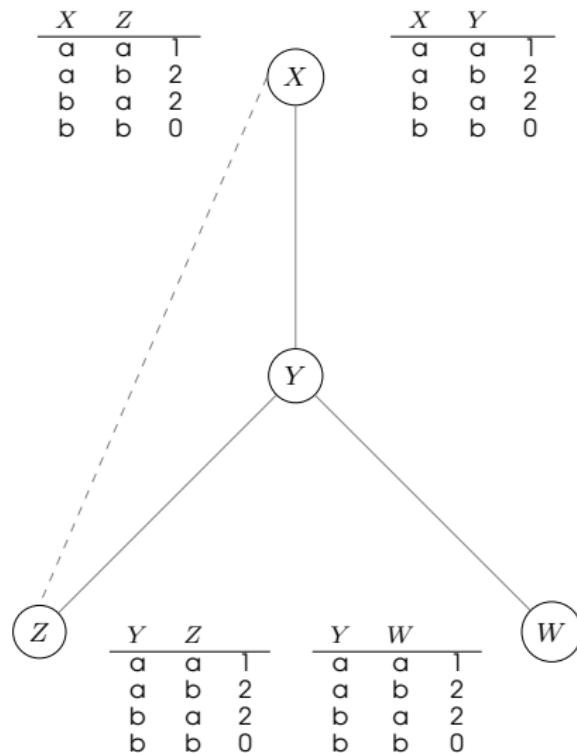
# DTREE : DPOP for DCOPs without backedges



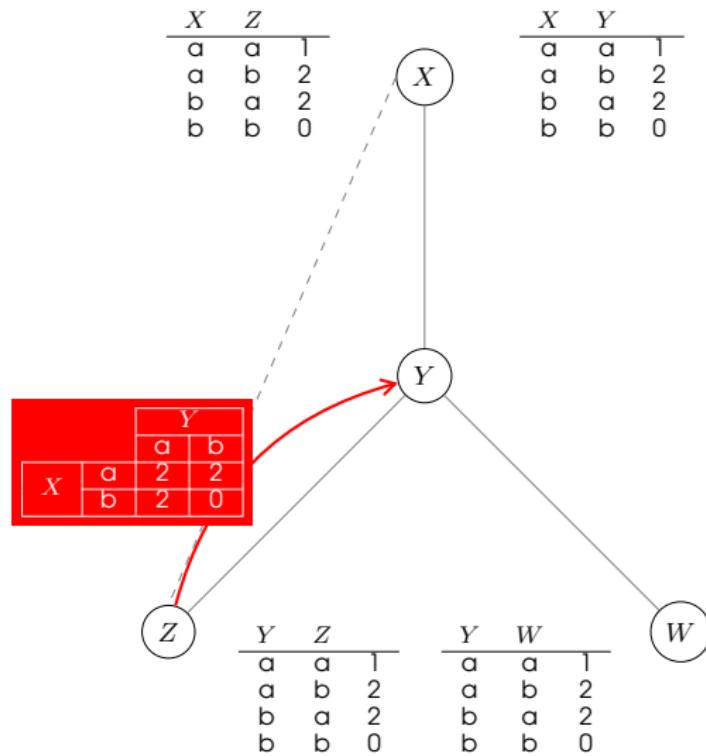
## DPOP for any DCOP



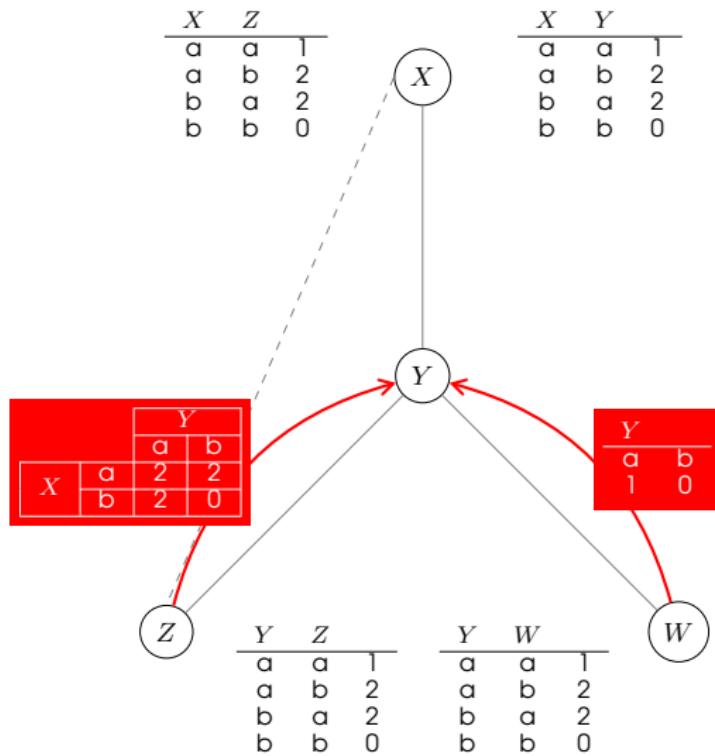
## DPOP for any DCOP



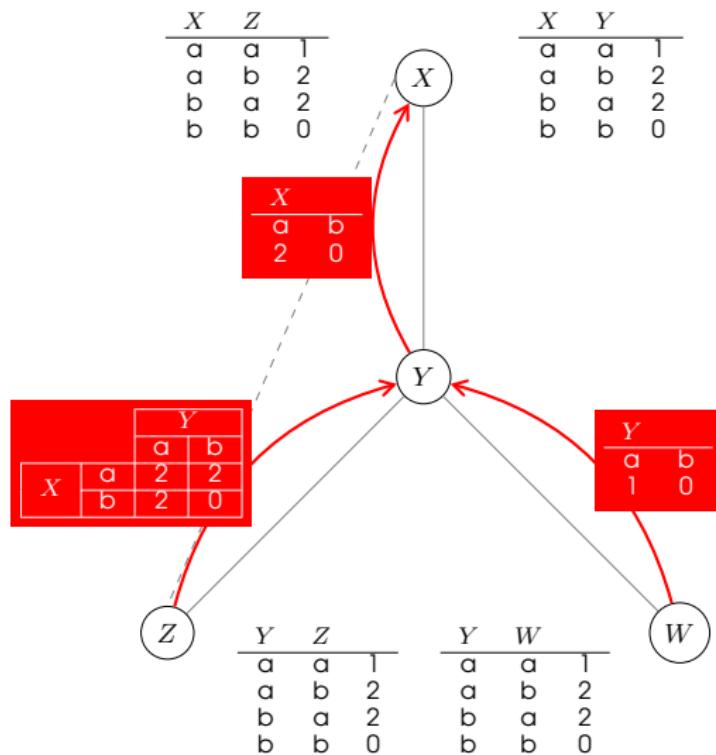
## DPOP for any DCOP



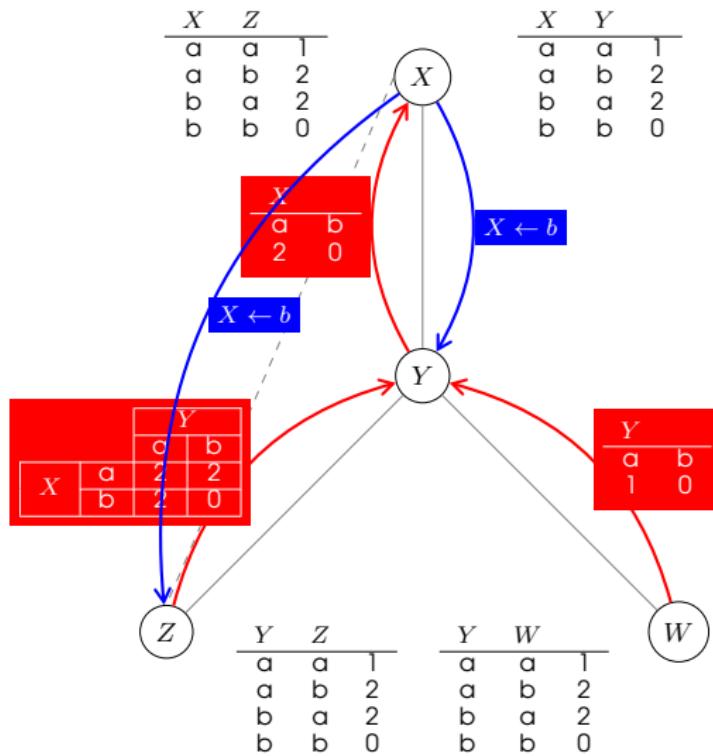
## DPOP for any DCOP



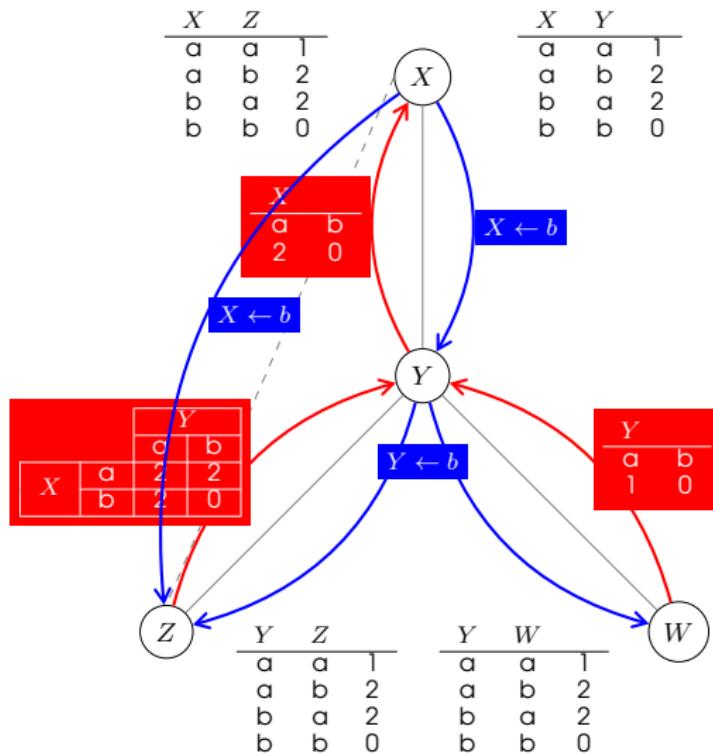
# DPOP for any DCOP



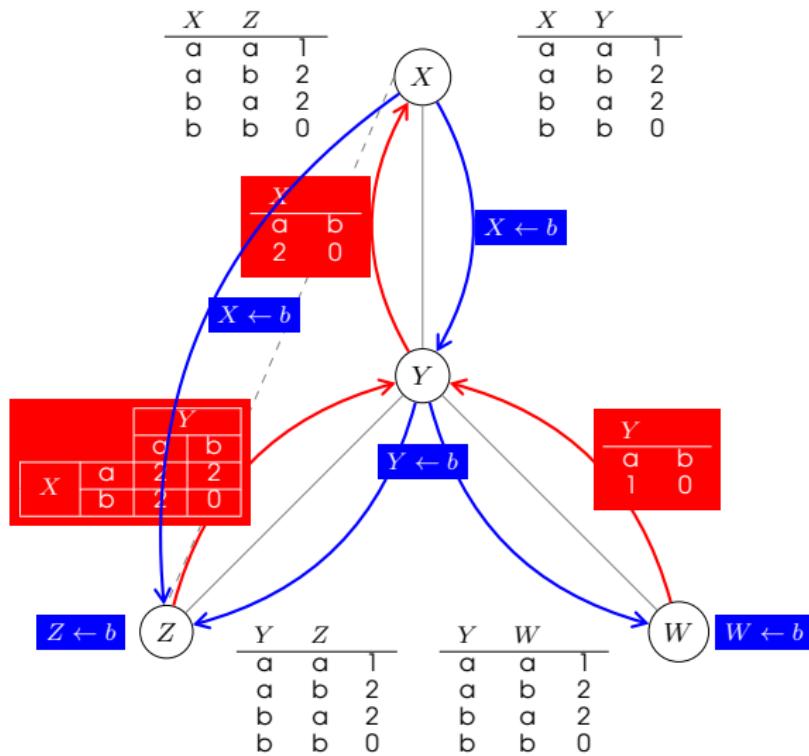
# DPOP for any DCOP



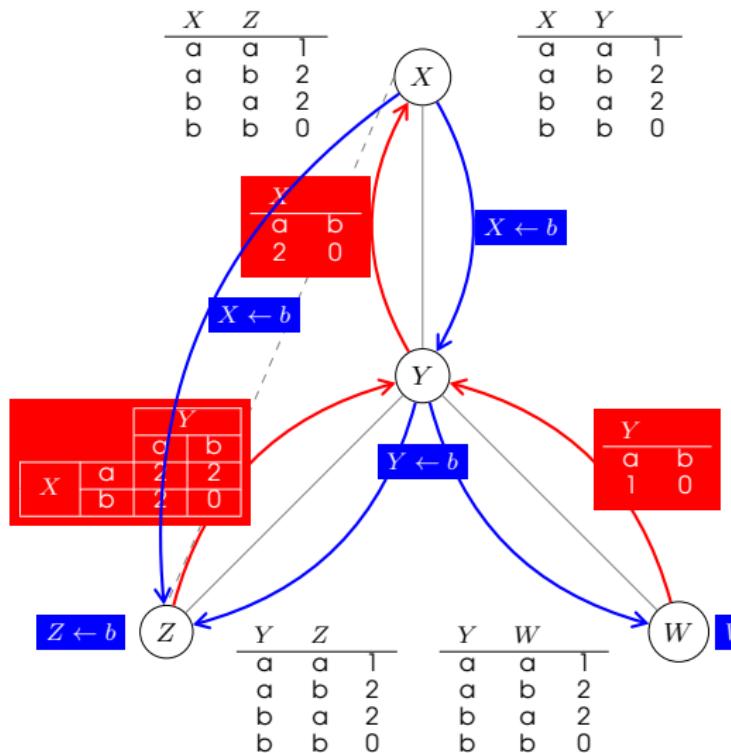
# DPOP for any DCOP



# DPOP for any DCOP



# DPOP for any DCOP



Optimal solution:

- linear number of messages
- message size: exponential

# GDL-based approaches

## ■ Generalized Distributive Law (AJI and McELIECE, 2000)

- ▶ Unifying framework for inference in Graphical models
- ▶ Builds on basic mathematical properties of semi-rings
- ▶ Widely used in Info theory, Statistical physics, Probabilistic models

## ■ Max-sum

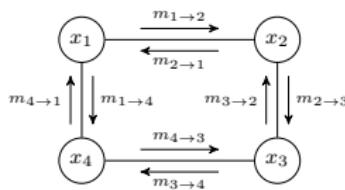
- ▶ DCOP settings: maximise social welfare

	$K$	“(+, 0)”	“(·, 1)”	short name
1.	$A$	(+, 0)	(·, 1)	
2.	$A[x]$	(+, 0)	(·, 1)	
3.	$A[x, y, \dots]$	(+, 0)	(·, 1)	
4.	$[0, \infty)$	(+, 0)	(·, 1)	sum-product
5.	$(0, \infty]$	(min, ∞)	(·, 1)	min-product
6.	$[0, \infty)$	(max, 0)	(·, 1)	max-product
7.	$(-\infty, \infty]$	(min, ∞)	(+, 0)	min-sum
8.	$[-\infty, \infty)$	(max, -∞)	(+, 0)	max-sum
9.	$\{0, 1\}$	(OR, 0)	(AND, 1)	Boolean
10.	$2^S$	( $\cup$ , $\emptyset$ )	( $\cap$ , $S$ )	
11.	$\Lambda$	( $\vee$ , 0)	( $\wedge$ , 1)	
12.	$\Lambda$	( $\wedge$ , 1)	( $\vee$ , 0)	

# Max-Sum

(FARINELLI et al., 2008)

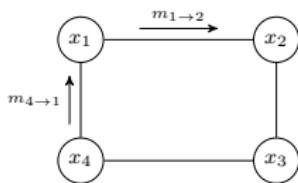
Agents iteratively computes local functions that depend only on the variable they control



# Max-Sum

(FARINELLI et al., 2008)

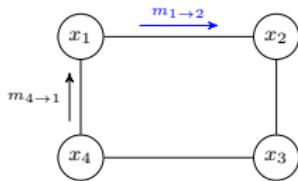
Agents iteratively computes local functions that depend only on the variable they control



# Max-Sum

(FARINELLI et al., 2008)

Agents iteratively computes local functions that depend only on the variable they control

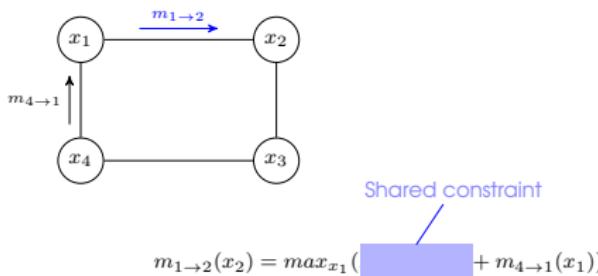


$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (F_{12}(x_1, x_2) + m_{4 \rightarrow 1}(x_1))$$

# Max-Sum

(FARINELLI et al., 2008)

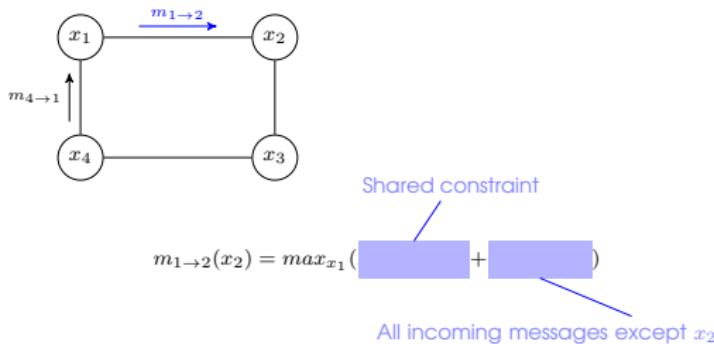
Agents iteratively computes local functions that depend only on the variable they control



# Max-Sum

(FARINELLI et al., 2008)

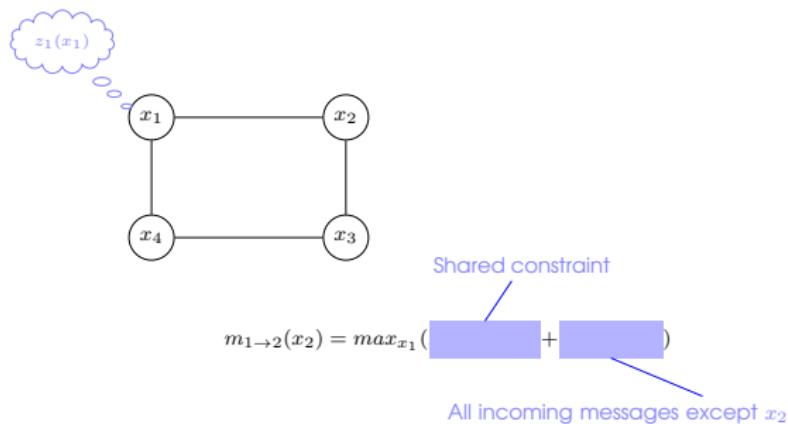
Agents iteratively computes local functions that depend only on the variable they control



# Max-Sum

(FARINELLI et al., 2008)

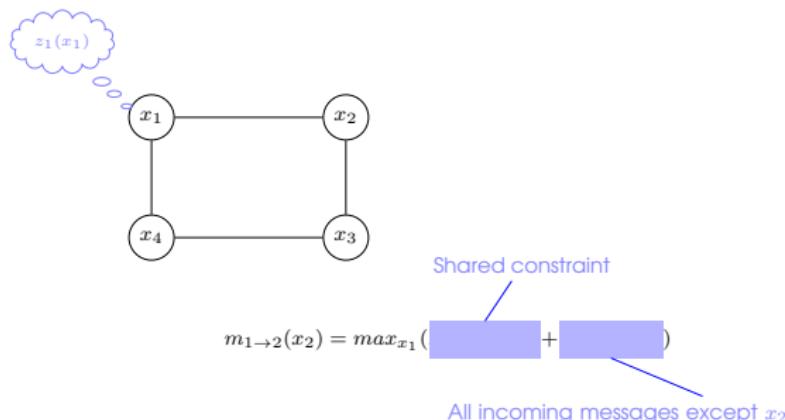
Agents iteratively computes local functions that depend only on the variable they control



# Max-Sum

(FARINELLI et al., 2008)

Agents iteratively computes local functions that depend only on the variable they control

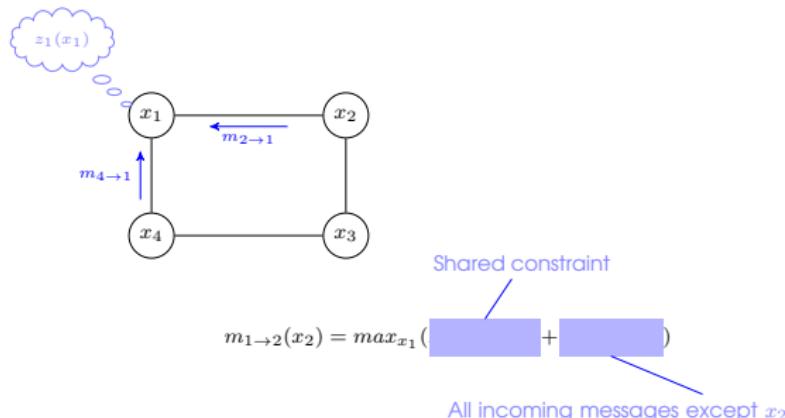


$$z_1(x_1) = m_{4 \rightarrow 1}(x_1) + m_{2 \rightarrow 1}(x_1)$$

# Max-Sum

(FARINELLI et al., 2008)

Agents iteratively computes local functions that depend only on the variable they control

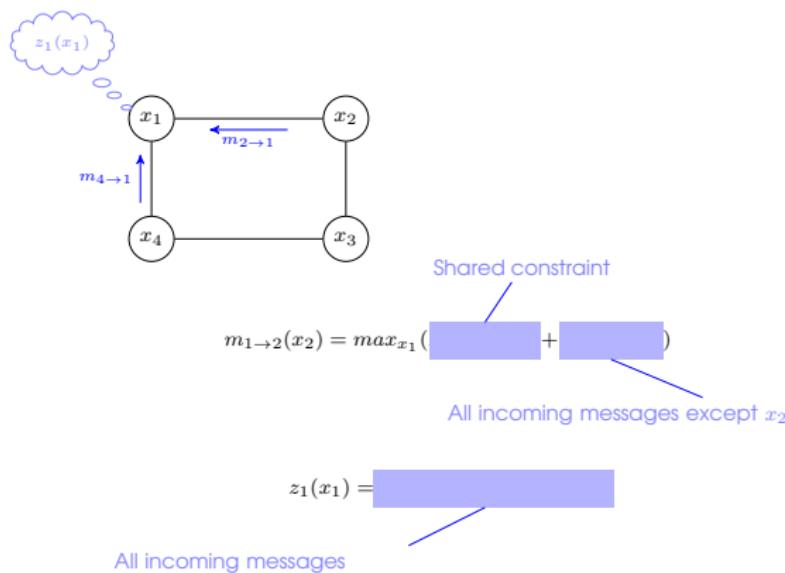


$$z_1(x_1) = m_{4 \rightarrow 1}(x_1) + m_{2 \rightarrow 1}(x_1)$$

# Max-Sum

(FARINELLI et al., 2008)

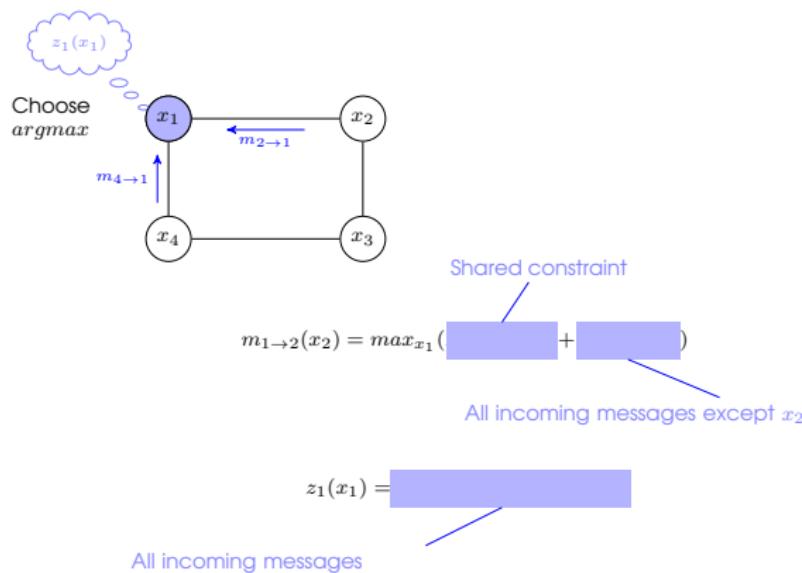
Agents iteratively computes local functions that depend only on the variable they control



# Max-Sum

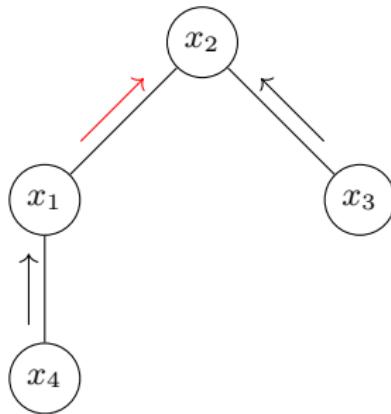
(FARINELLI et al., 2008)

Agents iteratively computes local functions that depend only on the variable they control



# Max-Sum on acyclic graphs

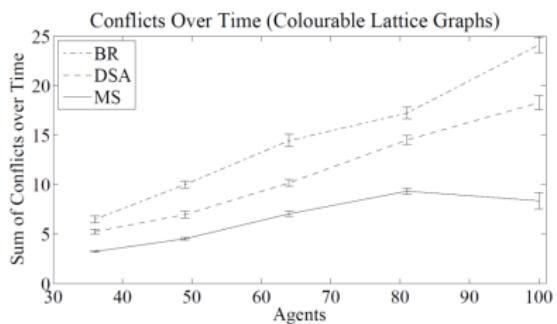
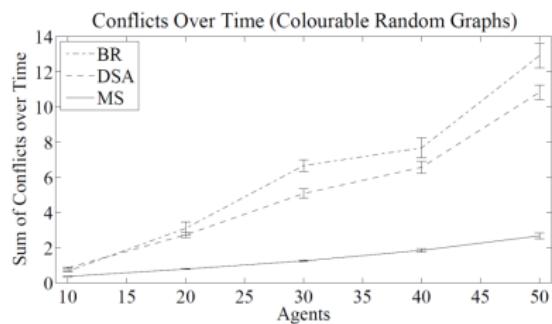
- Max-sum Optimal on acyclic graphs
  - ▶ Different branches are independent
  - ▶ Each agent can build a correct estimation of its contribution to the global problem ( $z$  functions)
- Message equations very similar to Util messages in DPOP
  - ▶ Sum messages from children and shared constraint
  - ▶ Maximize out agent variable
  - ▶ GDL generalizes DPOP (VINYALS et al., 2011)



$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (F_{12}(x_1, x_2) + m_{4 \rightarrow 1}(x_1))$$

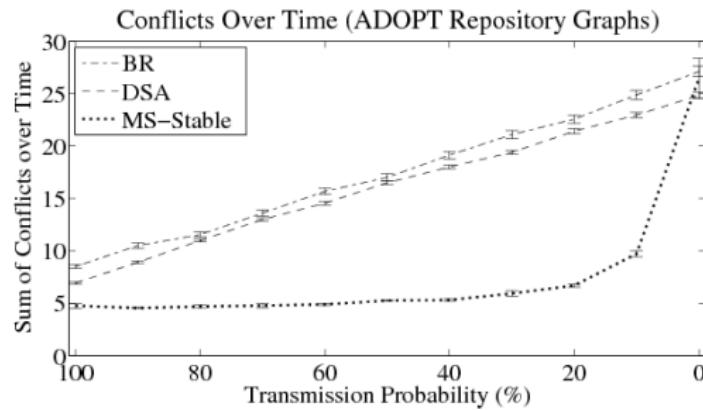
# Max-Sum Performance

- Good performance on loopy networks (FARINELLI et al., 2008)
  - ▶ When it converges very good results
    - ▶ Interesting results when only one cycle (WEISS, 2000)
  - ▶ We could remove cycle but pay an exponential price (see DPOP)



# Max-Sum for low power devices

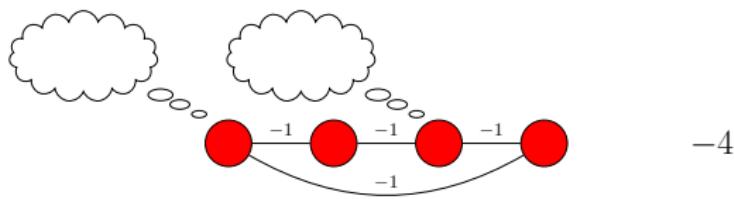
- Low overhead
  - ▶ Msgs number/size
- Asynchronous computation
  - ▶ Agents take decisions whenever new messages arrive
- Robust to message loss



# Local Greedy Approaches

## ■ Greedy local search

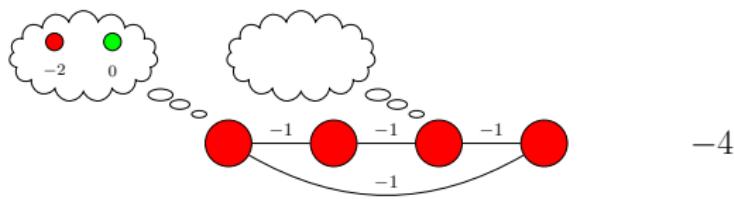
- ▶ Start from random solution
- ▶ Do local changes if global solution improves
- ▶ Local: change the value of a subset of variables, usually one



# Local Greedy Approaches

## ■ Greedy local search

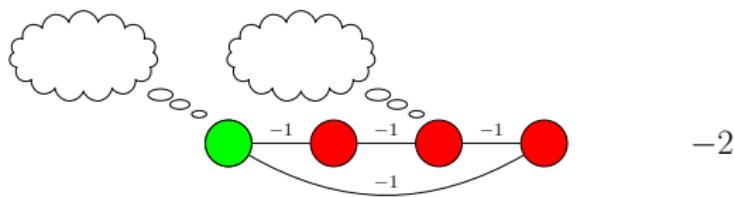
- ▶ Start from random solution
- ▶ Do local changes if global solution improves
- ▶ Local: change the value of a subset of variables, usually one



# Local Greedy Approaches

## ■ Greedy local search

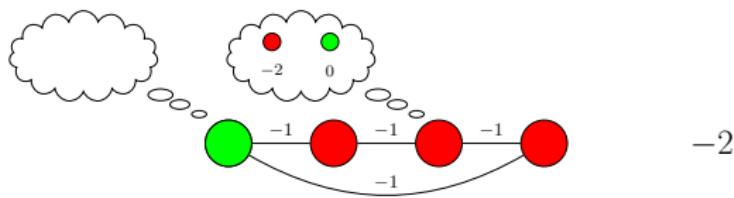
- ▶ Start from random solution
- ▶ Do local changes if global solution improves
- ▶ Local: change the value of a subset of variables, usually one



# Local Greedy Approaches

## ■ Greedy local search

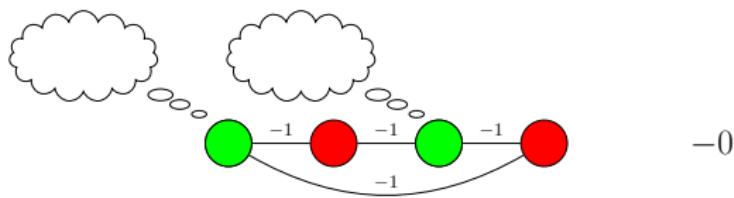
- ▶ Start from random solution
- ▶ Do local changes if global solution improves
- ▶ Local: change the value of a subset of variables, usually one



# Local Greedy Approaches

## ■ Greedy local search

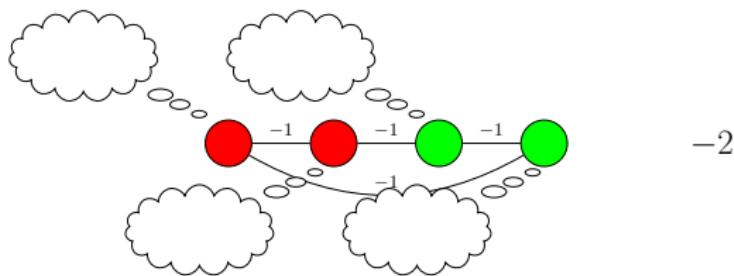
- ▶ Start from random solution
- ▶ Do local changes if global solution improves
- ▶ Local: change the value of a subset of variables, usually one



# Local Greedy Approaches

## ■ Problems

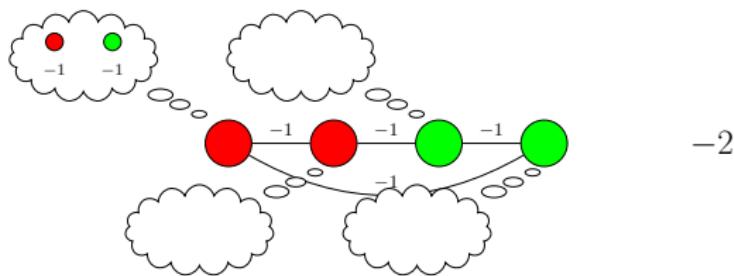
- ▶ Local minima
- ▶ Standard solutions: Random Walk, Simulated Annealing



# Local Greedy Approaches

## ■ Problems

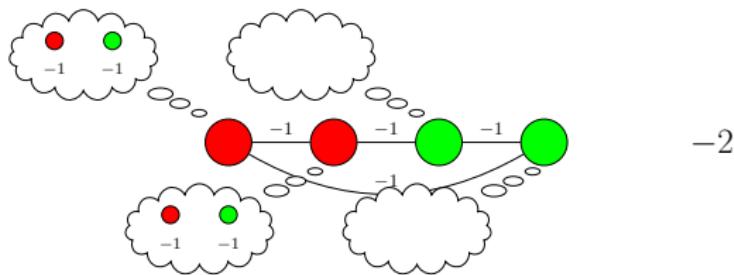
- ▶ Local minima
- ▶ Standard solutions: Random Walk, Simulated Annealing



# Local Greedy Approaches

## ■ Problems

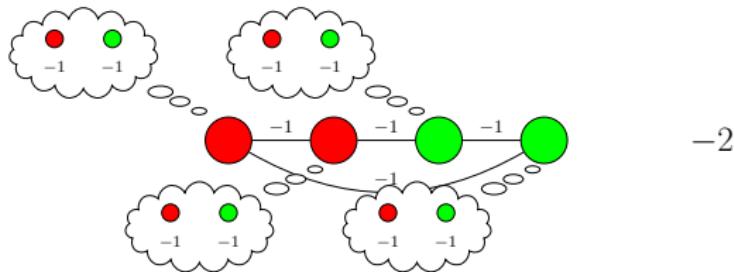
- ▶ Local minima
- ▶ Standard solutions: Random Walk, Simulated Annealing



# Local Greedy Approaches

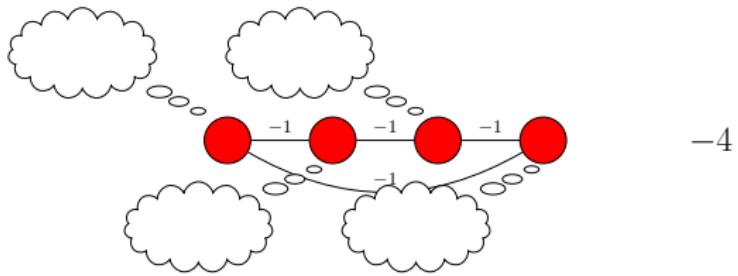
## ■ Problems

- ▶ Local minima
- ▶ Standard solutions: Random Walk, Simulated Annealing



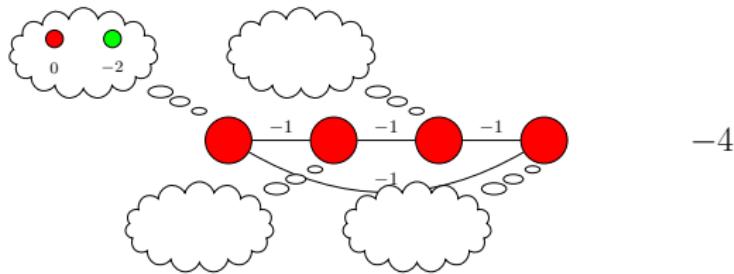
# Distributed Local Greedy approaches

- Local knowledge
- Parallel execution
  - ▶ A greedy local move might be harmful/useless
  - ▶ Need coordination



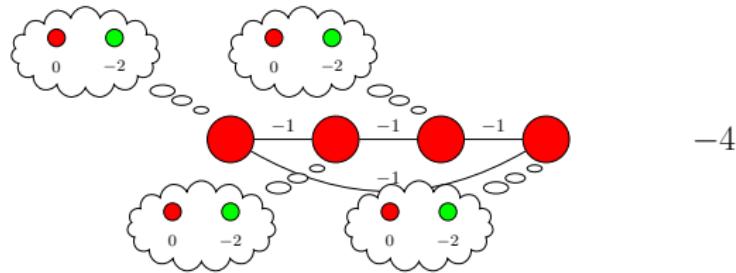
# Distributed Local Greedy approaches

- Local knowledge
- Parallel execution
  - ▶ A greedy local move might be harmful/useless
  - ▶ Need coordination



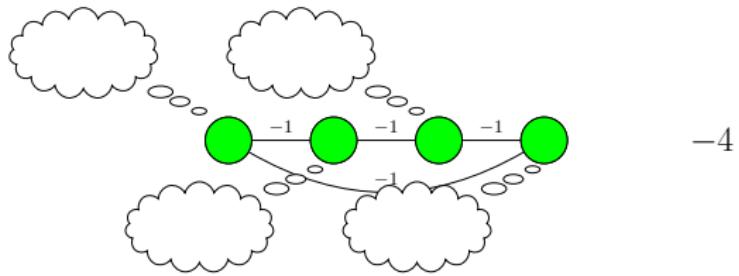
# Distributed Local Greedy approaches

- Local knowledge
- Parallel execution
  - ▶ A greedy local move might be harmful/useless
  - ▶ Need coordination



# Distributed Local Greedy approaches

- Local knowledge
- Parallel execution
  - ▶ A greedy local move might be harmful/useless
  - ▶ Need coordination

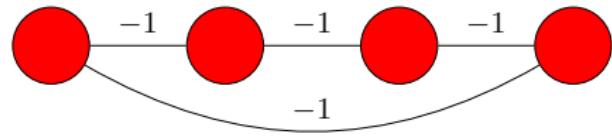


# Distributed Stochastic Search Algorithm (DSA)

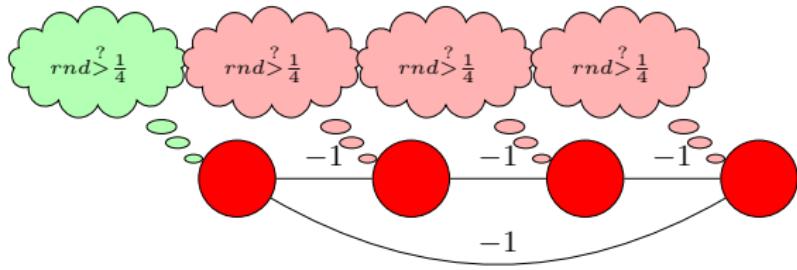
(ZHANG et al., 2005)

- Greedy local search with activation probability to mitigate issues with parallel executions
- DSA-1: change value of one variable at time
- Initialize agents with a random assignment and communicate values to neighbors
- Each agent:
  - ▶ Generates a random number and execute only if rnd less than activation probability
  - ▶ When executing changes value maximizing local gain
  - ▶ Communicate possible variable change to neighbors

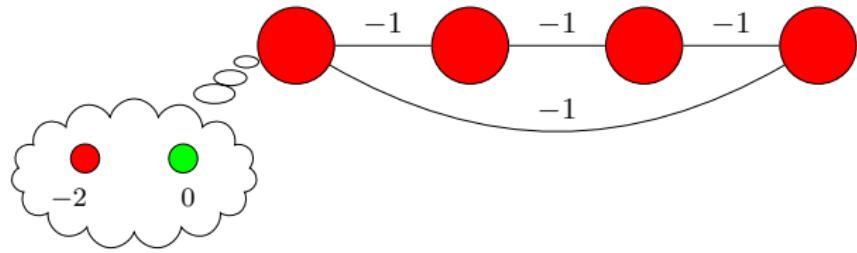
## DSA-1: Execution Example



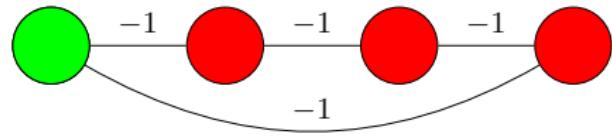
## DSA-1: Execution Example



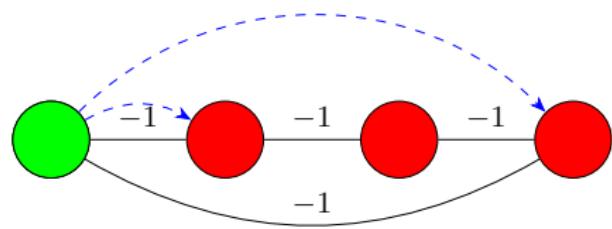
## DSA-1: Execution Example



## DSA-1: Execution Example



## DSA-1: Execution Example



# DSA-1: Discussion

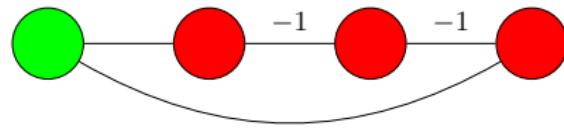
- Extremely “cheap” (computation/communication)
- Good performance in various domains
  - ▶ e.g. target tracking (FITZPATRICK and MEERTENS, 2003; ZHANG et al., 2003)
  - ▶ Shows an anytime property (not guaranteed)
  - ▶ Benchmarking technique for coordination
- Problems
  - ▶ Activation probability must be tuned (ZHANG et al., 2003)
  - ▶ No general rule, hard to characterise results across domains

# Maximum Gain Message (MGM-1)

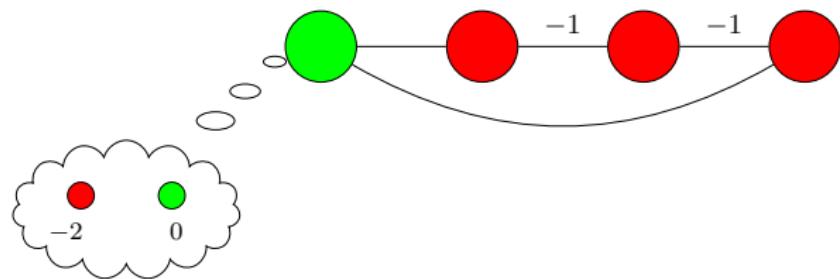
(MAHESWARAN et al., 2004)

- Coordinate to decide who is going to move
  - ▶ Compute and exchange possible gains
  - ▶ Agent with maximum (positive) gain executes
- Analysis
  - ▶ Empirically, similar to DSA
  - ▶ More communication (but still linear)
  - ▶ **No Threshold to set**
  - ▶ **Guaranteed to be monotonic** (Anytime behavior)

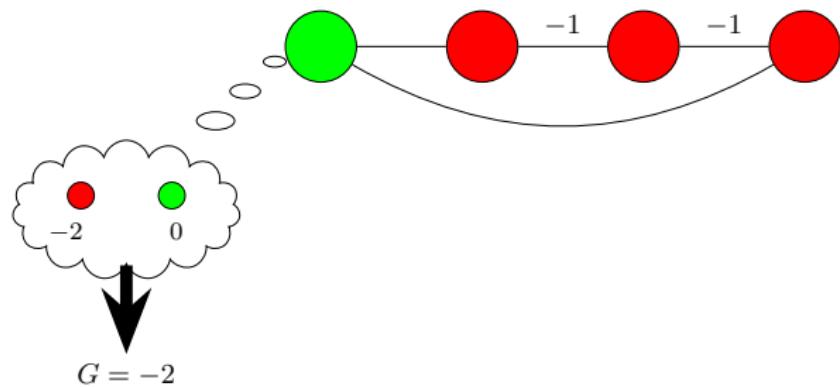
## MGM-1: Example



## MGM-1: Example

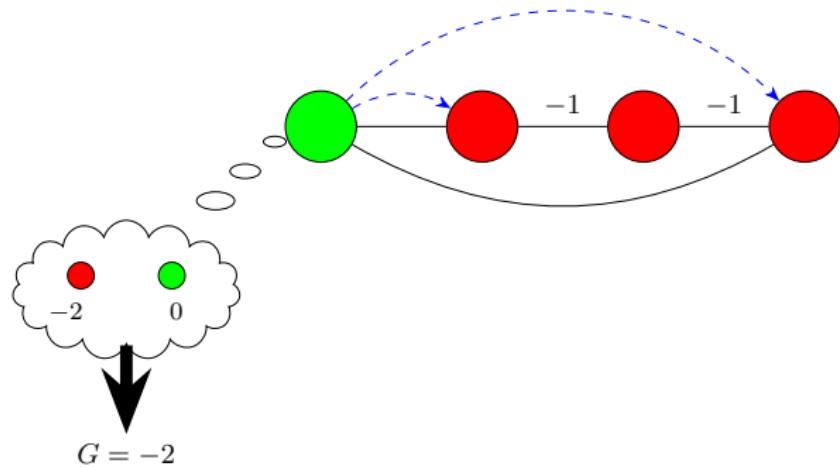


## MGM-1: Example

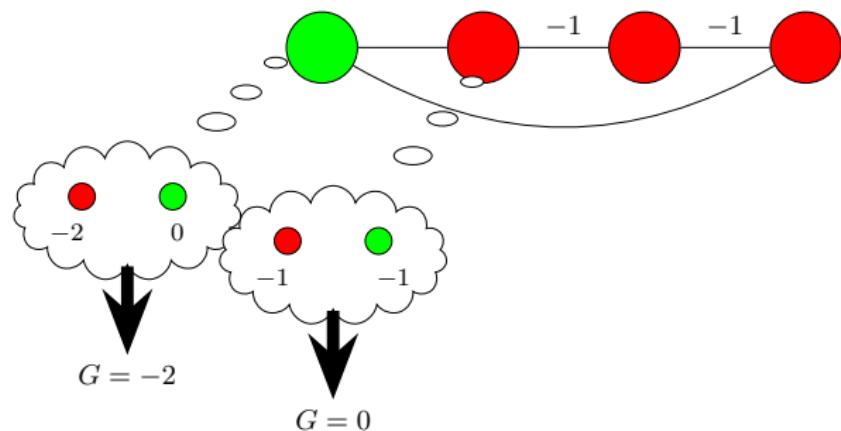


$$G = -2$$

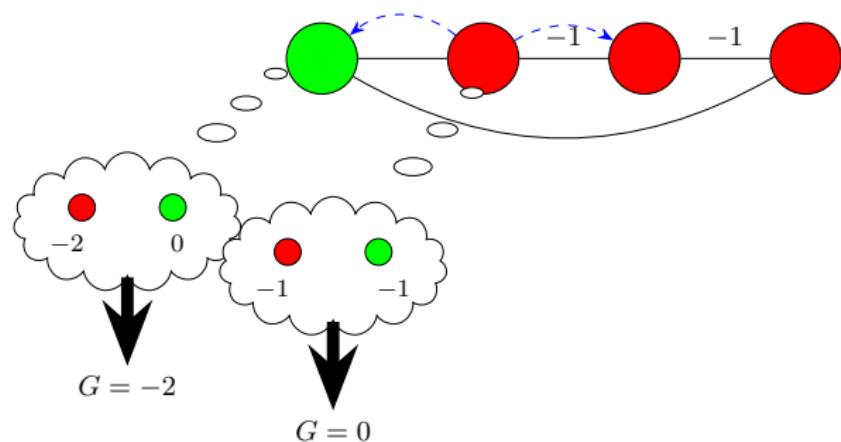
## MGM-1: Example



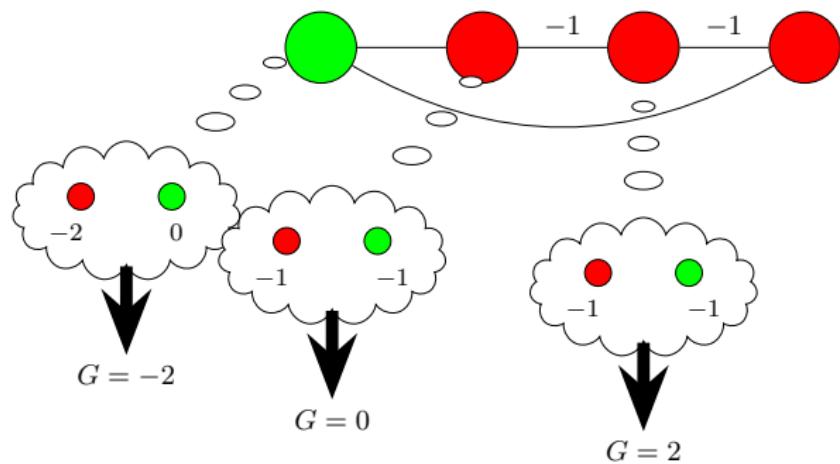
## MGM-1: Example



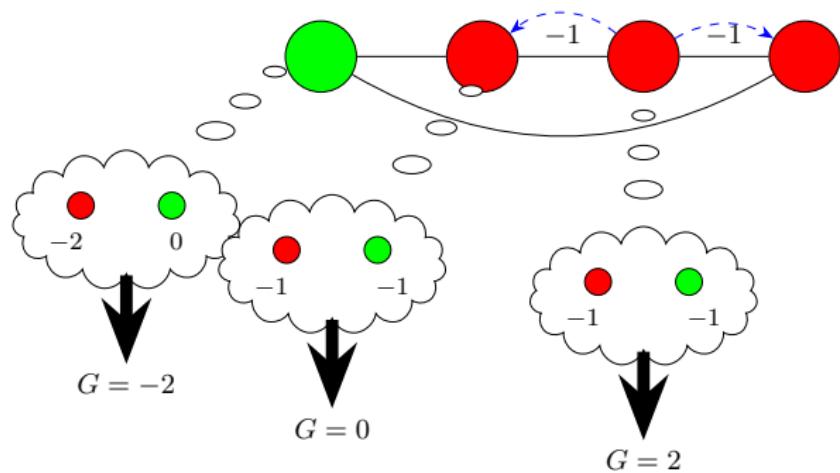
## MGM-1: Example



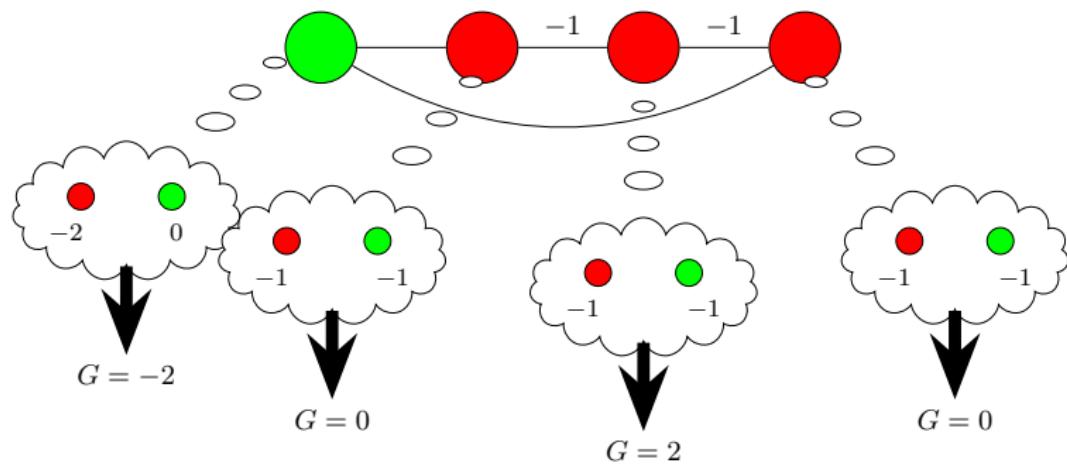
## MGM-1: Example



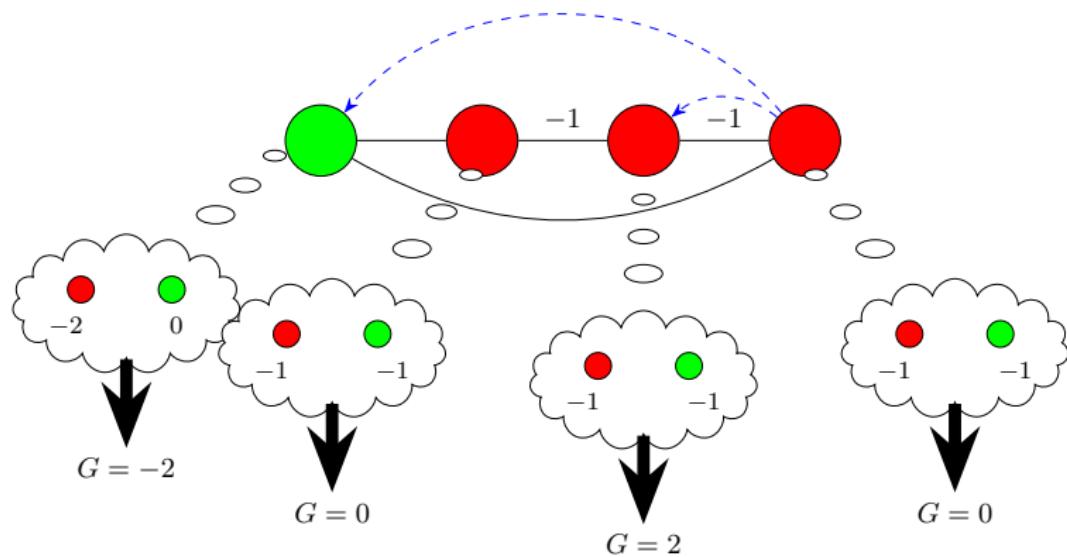
## MGM-1: Example



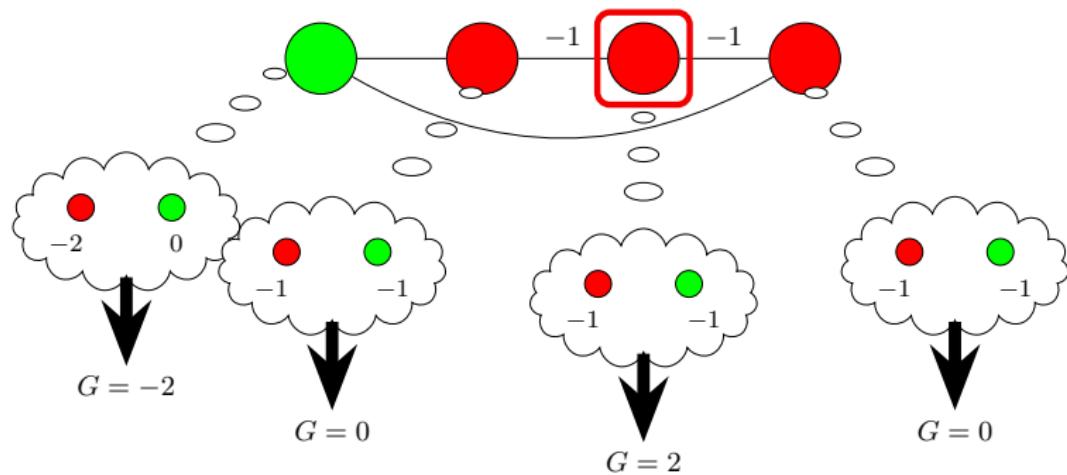
## MGM-1: Example



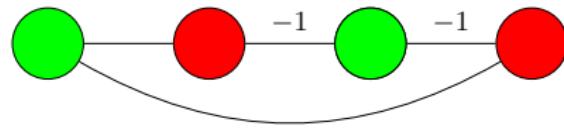
## MGM-1: Example



## MGM-1: Example



## MGM-1: Example



# Menu

DCOP Framework

Focus on Some Solution Methods

Hands on PyDCOP I

Focus on Smart Environment Configuration Problems

Distributing Computations

Hands on PyDCOP II

Dynamic DCOPs

Conclusion

# Hands on PyDCOP I

- Install VirtualBox
- Import the pyDCOP Virtual Machine (<https://bit.ly/2JUrrKP>)

- ▶ It's a Debian image with everything preinstalled:
- ▶ python3, pyDCOP, matplotlib, glpk, etc.

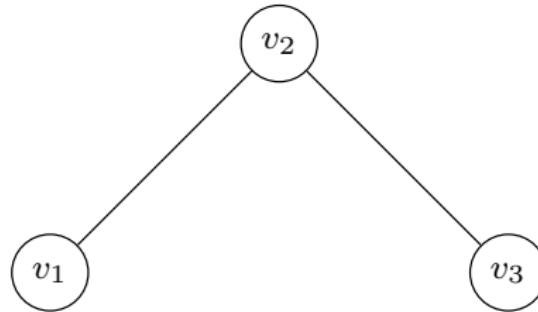
- Alternatively, follow

<https://pydcop.readthedocs.io/en/latest/installation.html>

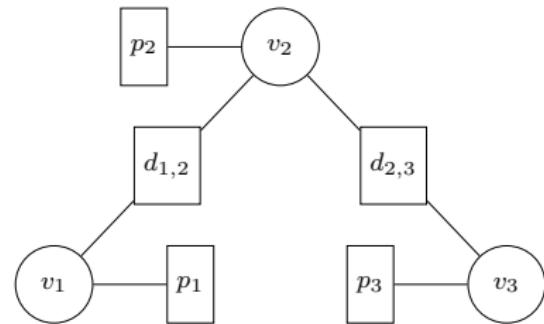
1. [https://pydcop.readthedocs.io/en/latest/tutorials/getting\\_started.html](https://pydcop.readthedocs.io/en/latest/tutorials/getting_started.html)
2. [https://pydcop.readthedocs.io/en/latest/tutorials/analysing\\_results.html](https://pydcop.readthedocs.io/en/latest/tutorials/analysing_results.html)

# Hands on PyDCOP I

## DCOP - Graph Coloring



(a) constraints graph



(b) factor graph

- **Objective:** minimize
- **Domain:** 2 colors  $R$  and  $B$
- **Variables:**  $V_1, V_2, V_3$
- **Constraints:** neighbors must have different colors + preferences
- **Agents:** 3 agents

Yaml representation

# Hands on PyDCOP I

## Solving the Graph Coloring DCOP

### Command:

```
$ pydcop solve --algo dpop graph_coloring.yaml
```

### Output:

```
...
"assignment": {
    "v1": "R",
    "v2": "G",
    "v3": "R"
},
"cost": -0.1,
...
```

### With other algorithms:

```
$ pydcop --timeout 2 solve --algo dsa graph_coloring.yaml
$ pydcop solve --algo mgm --algo_params stop_cycle:20 \
    graph_coloring.yaml
```

# Hands on PyDCOP I

## Results

### Full results :

```
{  
    "agt_metrics": {  
        ...  
    },  
    "assignment": {  
        "v1": "R",  
        "v2": "G",  
        "v3": "R"  
    },  
    "cost": -0.1,  
    "cycle": 20,  
    "msg_count": 158,  
    "msg_size": 158,  
    "status": "FINISHED",  
    "time": 0.03201029699994251,  
    "violation": 0  
}
```

# Hands on PyDCOP I

## Logs

### Simple:

use -v 0..3

```
$ pydcop -v 2 solve --algo dpop graph_coloring.yaml
```

### Precise :

use -log <log.conf>

```
$ pydcop --log log.conf -t 1 solve --algo dsa graph_coloring.yaml
```

Now, look at algo.log

# Hands on PyDCOP I

## Run-time metrics

**periodic:** "--collect\_on period --period <p>"

```
$ pydcop --log log.conf -t 10 solve \  
--collect_on period --period <p>  
--algo dsa graph_coloring.yaml
```

**cycle:** "--collect\_on cycle\_change"

Only supported with synchronous algorithms !

```
$ pydcop solve --algo mgm --algo_params stop_cycle:20 \  
--collect_on cycle_change --run_metric ./metrics.csv \  
graph_coloring_50.yaml
```

**value:** "--collect\_on value\_change"

```
$ pydcop -t 5 solve --algo mgm --collect_on value_change \  
--run_metric ./metrics_on_value.csv \  
graph_coloring_50.yaml
```

# Hands on PyDCOP I

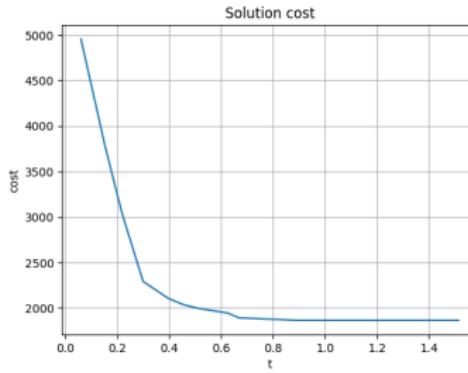
## Run-time metrics

With a bigger graph coloring problem

```
$ pydcop solve --algo mgm --algo_params stop_cycle:20 \
    --collect_on cycle_change \
    --run_metric ./metrics.csv \
    graph_coloring_50.yaml
```

Plotting with matplotlib

```
$ python3 plot_cost.py ./metrics.csv
```



# Menu

DCOP Framework

Focus on Some Solution Methods

Hands on PyDCOP I

Focus on Smart Environment Configuration Problems

Distributing Computations

Hands on PyDCOP II

Dynamic DCOPs

Conclusion

# SECP model

*Smart Environment Configuration Problem (RUST et al., 2016)*

- Example of applying DCOPs to a "real" problem
- Coordinate objects in the building
- Model
  - ▶ objects
  - ▶ relations between objects and environment
  - ▶ user objectives and requirements
- Formulate the problem as an optimization problem



# SECP model

*Smart Environment Configuration Problem (RUST et al., 2016)*

Focus on smart lighting use cases

- **Objects:** anything that can produce light: light bulbs, windows with rolling shutter, etc.
- **User preferences:** having a predefined luminosity level in a room, under some conditions
- **Energy efficiency**

Linking objects and user preferences:

- How to model the luminosity in a room ? **variable**
- How to model the dependency between the light sources and the luminosity ? **function / constraint**

# SECP model

Example application to ambient intelligence scenario



## ■ Actuators

- ▶ Connected light bulbs, TV, Rolling shutters, ...

## ■ Sensors

- ▶ Presence detector, Luminosity Sensor, etc.

## ■ Physical Dependency Models

- ▶ E.g. Living-room light model

## ■ User Preferences

- ▶ Expressed as rules :

IF	presence_living_room	=	1
AND	light_sensor_living_room	<	60
THEN	light_level_living_room	←	60
AND	shutter_living_room	←	0

# SECP model

Example application to ambient intelligence scenario



## ■ Actuators

- ▶ Decision variable  $x_i$ , domain  $\mathcal{D}_{x_i}$
- ▶ Cost function  $c_i : \mathcal{D}_{x_i} \rightarrow \mathbb{R}$

## ■ Sensors

- ▶ Read-only variable  $s_l$ , domain  $\mathcal{D}_{s_l}$

## ■ Physical Dependency Models $\langle y_j, \phi_j \rangle$

- ▶ Give the expected state of the environment from a set of actuator-variables influencing this model
- ▶ Variable  $y_j$  representing the expected state of the environment
- ▶ Function  $\phi_j : \prod_{\varsigma \in \sigma(\phi_j)} \mathcal{D}_\varsigma \rightarrow \mathcal{D}_{y_j}$

## ■ User Preferences

- ▶ Utility function  $u_k$
- ▶ Distance from the current expected state to the target state of the environment

# Formulating SECP as a DCOP

Multi-objective optimization problem

$$\begin{aligned} \min_{x_i \in \nu(\mathfrak{A})} \quad & \sum_{i \in \mathfrak{A}} c_i \quad \text{and} \quad \max_{\substack{x_i \in \nu(\mathfrak{A}) \\ y_j \in \nu(\Phi)}} \sum_{k \in \mathfrak{R}} u_k \\ \text{s.t. } \quad & \phi_j(x_j^1, \dots, x_j^{\overline{\phi_j}}) = y_j \quad \forall y_j \in \nu(\Phi) \end{aligned}$$

Then mono-objective DCOP formulation

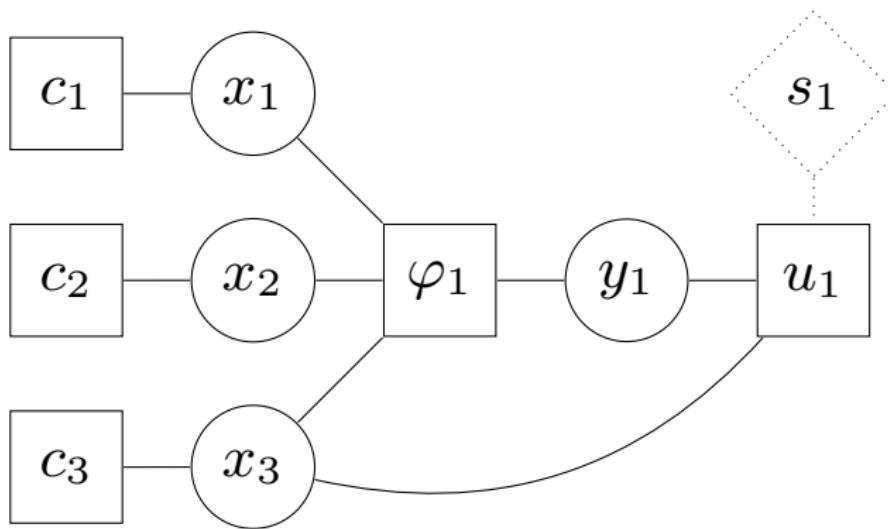
$$\max_{\substack{x_i \in \nu(\mathfrak{A}) \\ y_j \in \nu(\Phi)}} \omega_u \sum_{k \in \mathfrak{R}} u_k - \omega_c \sum_{i \in \mathfrak{A}} c_i + \sum_{\varphi_j \in \Phi} \varphi_j$$

with reformulation of hard constraints  $\phi_j$  into soft ones:

$$\varphi_j(x_j^1, \dots, x_j^{|\sigma(\phi_j)|}, y_j) = \begin{cases} 0 & \text{if } \phi_j(x_j^1, \dots, x_j^{|\sigma(\phi_j)|}) = y_j \\ -\infty & \text{otherwise} \end{cases}$$

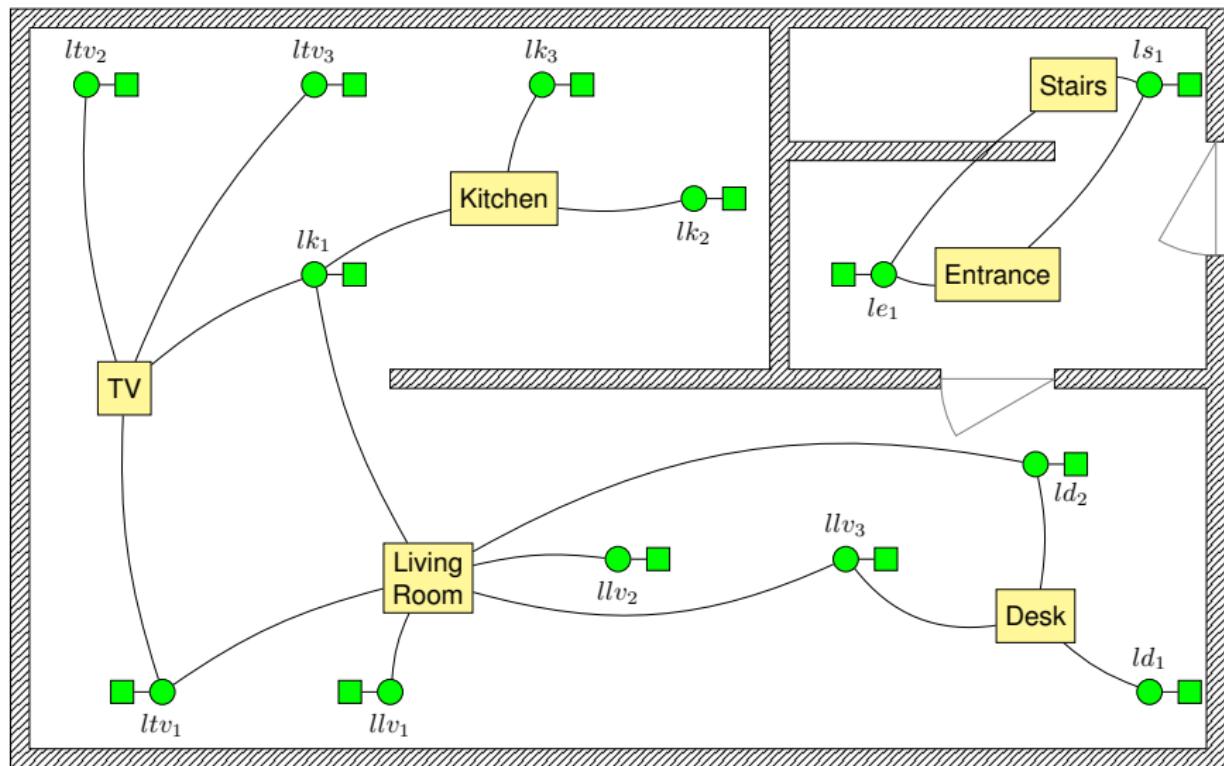
# Formulating SECP as a DCOP

Representing a DCOP as a factor graph



# SECP Factor Graph

in a house (without rules)



# Menu

DCOP Framework

Focus on Some Solution Methods

Hands on PyDCOP I

Focus on Smart Environment Configuration Problems

Distributing Computations

Hands on PyDCOP II

Dynamic DCOPs

Conclusion

# Distribution of computations

## Allocating computations to agents

- DCOP:  $\langle \mathcal{A}, \mathcal{X}, \mathcal{D}, \mathcal{C}, \mu \rangle$
- $\mu$  : function mapping variables to their associated agent

## Why is distribution needed ?

Common assumptions:

- computation  $\equiv$  variable
- each agent controls exactly one variable (bijection)
- binary constraints

Real distributed problems:

- agents must be hosted on real devices
- the set of devices might be given by the problem
- for some variables the relation with an agent is obvious, but not always

# Distribution of computations

## Allocating computations to agents

- Distributing computations
  - ▶ computations depends on the graph model
  - ▶ variables and / or factors
  
- Distribution impacts the system characteristics :
  - ▶ speed,
  - ▶ communication load,
  - ▶ hosting costs, etc.

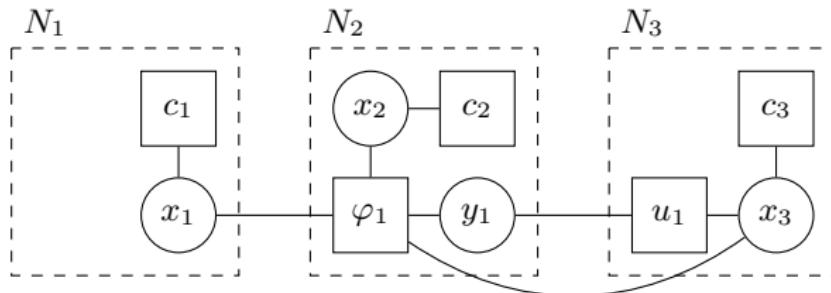
# Distribution of computations

## Allocating computations to agents

### ■ Simple heuristic

- ▶ No computation on sleepy devices (sensors)
- ▶ Computation should be close the the impacted variables
- ▶ Spread the computation load amongst agents

### ■ How good is it ?



# Distribution of computations

## Optimal definition

*Optimal distribution ?*

- Problem dependent
- Optimization problem :  
find the best distribution, for your problem's criteria
- Optimal distribution  $\equiv$  graph partitioning,  
NP-hard in general (BOULLE, 2004)

# Distribution of computations

Better definition

SECP distribution problem

- Devices have limited memory
- Communication is expensive and has limited bandwidth
- Variable related to an actuator are hosted by it
- Objective : **minimize overall communication between agents**

Optimization problem : define an ILP for it !

## Binary ILP for computation distribution

- $x_i^k$ , binary variables that map computations to agents and  $\alpha_{ij}^{mn} = x_i^m \cdot f_j^n$

$$\forall x_i \in X, \quad \sum_{a_m \in \mathbf{A}} x_i^m = 1 \quad (1)$$

- Message's size between variable  $x_i$  and factor  $f_j$ :  $msg(i, j)$

$$\underset{x_i^m}{\text{minimize}} \quad \sum_{(i,j) \in D} \sum_{(m,n) \in \mathbf{A}^2} msg(i, j) \cdot \alpha_{ij}^{mn} \quad (2)$$

- Memory footprint of a computation:  $\text{weight}(e)$ , and memory capacity for a device:  $\text{cap}(a_k)$

$$\forall a_m \in \mathbf{A}, \quad \sum_{x_i \in D} \text{weight}(x_i) \cdot x_i^m \leq \text{cap}(a_m) \quad (3)$$

- and a few linearization constraints

# Binary ILP for computation distribution

More generic case:

- Add route cost:  $\text{com}(i, j, m, n)$

$\forall x_i, x_j \in \mathbf{X}, \forall a_m, a_n \in \mathbf{A},$

$$\text{com}(i, j, m, n) = \begin{cases} \text{msg}(i, j) \cdot \text{route}(m, n) & \text{if } (i, j) \in D, m \neq n \\ 0 & \text{otherwise} \end{cases} \quad (4)$$

$$\underset{x_i^m}{\text{minimize}} \quad \sum_{(i,j) \in D} \sum_{(m,n) \in \mathbf{A}^2} \text{com}(i, j, m, n) \cdot \alpha_{ij}^{mn} \quad (5)$$

- Add hosting costs :  $\text{host}(a_m, x_i)$

$$\underset{x_i^m}{\text{minimize}} \quad \sum_{(x_i, a_m) \in X \times \mathbf{A}} x_i^m \cdot \text{host}(a_m, x_i) \quad (6)$$

# Binary ILP for computation distribution

$$\begin{aligned} \underset{x_i^m}{\text{minimize}} \quad & \omega_{\text{com}} \cdot \sum_{(i,j) \in D} \sum_{(m,n) \in \mathbf{A}^2} \text{com}(i, j, m, n) \cdot \alpha_{ij}^{mn} \\ & + \omega_{\text{host}} \cdot \sum_{(x_i, a_m) \in X \times \mathbf{A}} x_i^m \cdot \text{host}(a_m, x_i) \end{aligned} \quad (7)$$

subject to

$$\forall a_m \in \mathbf{A}, \quad \sum_{x_i \in D} \text{weight}(x_i) \cdot x_i^m \leq \text{cap}(a_m) \quad (8)$$

$$\forall x_i \in X, \quad \sum_{a_m \in \mathbf{A}} x_i^m = 1 \quad (9)$$

$$\forall x_i \in X, \quad \alpha_{ij}^{mn} \leq x_i^m \quad (10)$$

$$\forall x_j \in X, \quad \alpha_{ij}^{mn} \leq x_j^m \quad (11)$$

$$\forall x_i, x_j \in X, a_m \in A, \quad \alpha_{ij}^{mn} \geq x_i^m + x_j^n - 1 \quad (12)$$

## Solving the ILP for computation deployment

- NP-hard, but can be solved with branch-and-cut  
LP solvers are very good at this
- Yet, only possible for small instances
- Gives us a reference for optimality: benchmarking
- When not solvable, still gives us a metrics to compare heuristics

# Menu

DCOP Framework

Focus on Some Solution Methods

Hands on PyDCOP I

Focus on Smart Environment Configuration Problems

Distributing Computations

Hands on PyDCOP II

Dynamic DCOPs

Conclusion

## Hands on PyDCOP II

1. [https://pydcop.readthedocs.io/en/latest/tutorials/deploying\\_on\\_machines.html](https://pydcop.readthedocs.io/en/latest/tutorials/deploying_on_machines.html)
2. <https://pydcop.readthedocs.io/en/latest/usage/cli/run.html>
3. <https://pydcop.readthedocs.io/en/latest/usage/cli/distribute.html>

# Hands on PyDCOP II

SECP

A Very simple SECP: single room

- 3 light bulbs, 1 model and 3 rules
- `/tutorials/hands-on_2/single_room.yaml`
- Solve with

```
pydcop --log log.conf -t 10 solve  
--algo maxsum --algo_params damping:0.8  
--dist adhoc single_room.yaml
```

- Result : "cost": 702.300000000004, ...
  - ▶ not that good ...
  - ▶ Look at the yaml definition
  - ▶ the rules contradict each other !
- Change the yaml definition
  - ▶ comment out rules to keep only one active
  - ▶ could be done with 'read-only' variables
  - ▶ solve it again

# Hands on PyDCOP II

## SECP - Running on several machines

- We used `solve`

- ▶ great for testing
  - ▶ everything run locally, in the same process

- Launching several agents:

- ▶ One agent for each light bulb a1, a2 and a3 (change port for each agent)

```
pydcop -v3 agent -n a1 -p 9001  
--orchestrator 127.0.0.1:9000
```

- ▶ an orchestrator

```
pydcop --log log.conf -t 10 orchestrator \  
--algo maxsum --algo_params damping:0.8 \  
--dist adhoc single_room.yaml
```

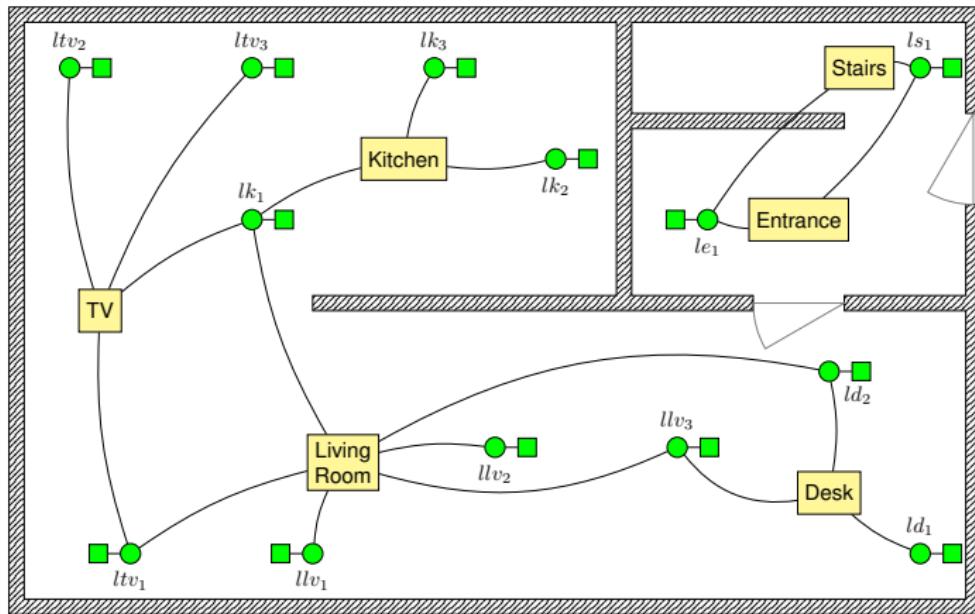
- ▶ run the agents on different Virtual machines, different computers

# Hands on PyDCOP II

A bigger SECP

in `/tutorials/hands-on_2/SimpleHouse.yml`

13 light bulbs, 6 models



# Hands on PyDCOP II

## Distributing a SECP

```
$ pydcop --output dist_house_fg_ilp.yaml distribute -d ilp_compref \  
-a maxsum SimpleHouse.yml
```

Need to specify the algorithm

- computation graph
- weight of computation's
- size of computations' messages

On such a small system, we can compute the optimal distribution !

# Hands on PyDCOP II

## Distributing a SECP

```
cost: 8725.0
distribution:
    a_d1: [mv_desk, mc_desk, l_d1, r_work, mc_livingroom, mv_livingroom]
    a_d2: [l_d2]
    a_e1: [mv_entry, r_entry, mv_stairs, l_e1, mc_entry, mc_stairs]
    a_e2: [l_e2]
    a_k1: [l_k1]
    a_k2: [l_k2]
    a_k3: [l_k3]
    a_lv1: [l_lv1]
    a_lv2: [mc_kitchen, l_lv2]
    a_lv3: [l_lv3]
    a_tv1: [l_tv1]
    a_tv2: [l_tv2]
    a_tv3: [r_lunch, l_tv3, mv_tv, r_cooking, r_homecinema, mc_tv, mv_kitchen]
inputs:
    algo: maxsum
    dcop: [SimpleHouse.yml]
    dist_algo: ilp_comref
    graph: factor_graph
```

# Menu

DCOP Framework

Focus on Some Solution Methods

Hands on PyDCOP I

Focus on Smart Environment Configuration Problems

Distributing Computations

Hands on PyDCOP II

Dynamic DCOPs

Conclusion

## SECP and DCOP

So far we have:

- Designed a model for SECP
- Formulated this model as a DCOP
- Distributed the computation of the DCOP on devices / agents (bootstrap)
- Run our system to get an auto-configuration of the devices

But what happens  
in dynamic environments ?  
if objects appear and disappear?

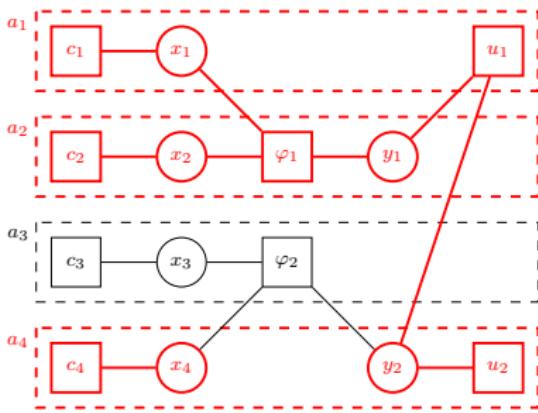
# SECP is a dynamic problem

## Dynamics in the infrastructure

- Devices can disappear
- New devices can be added to the system

At runtime..

- No powerful device available to solve the ILP
- The deployment must be repaired: self-adaptation
- Only consider a portion of the factor graph: the neighborhood



# k-resilience

## Dynamics in the infrastructure

### Definition (*k*-resiliency)

*k*-resiliency is the capacity for a system to repair itself and operate correctly even in the case of the disappearance of up to  $k$  agents

- Two parts:
  - ▶ Do not loose the definition of the computations: **replication**
  - ▶ Migrate the orphaned computations to another agent: **selection / activation**
- Apply to any graph of computations, not only DCOP

# Replication of computations

## Replica distribution

- For each computation, place  $k$  replica on  $k$  other agents  
replica *equiv* definition of the computation
- Must be distributed !
- Optimal replication ? impact the set of available agents when repairing  
which criteria ? too hard (quadratic multiple knapsack problem)...

## Distributed Replica Placement Method (DRPM)

- Heuristic : place replica on agents close (network) the active computation, while respecting capacity
- Distributed version of iterative lengthening (aka uniform cost search based on path costs)

# Replication of computations

iterative lengthening on route and hosting costs

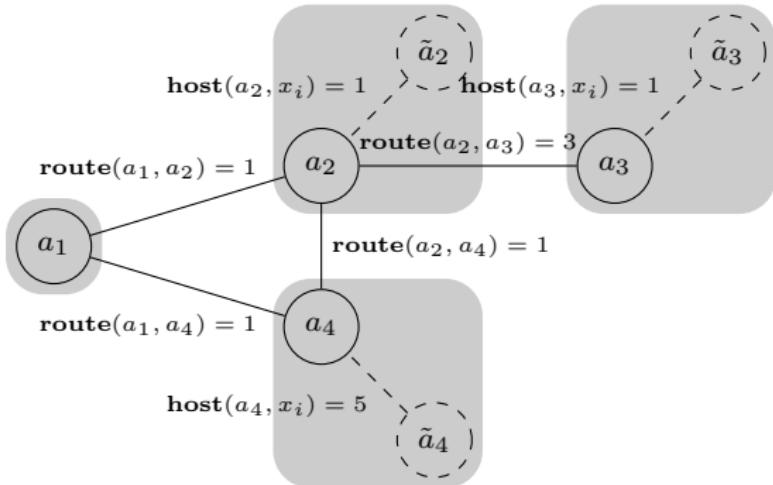


Figure: A sample **route+host**-graph with 4 agents (in gray), where  $a_1$  search for hosting computation  $x_i$ . For  $k = 2$ , DRMP places a replica on  $a_2$  (cost of  $1 + 1 = 2$ ) and another on  $a_3$  (cost of  $1 + 3 + 1 = 5$ ) if enough capacity on these two agents, since the minimum cost path to host on  $a_4$  is higher ( $1 + 5 = 6$ ).

# Migrating computations

## Selecting an agent

Migrating a set of  $x_i$  computations  $X_c$

- set of candidate agents  $A_c$
- migrating the computation must not exceed agent's capacity
- for each computation, select the agent that minimize hosting and communication cost

Same optimization problem than for initial distribution, but on a subset of the graph

Distributed process !

# Migrating computations

## Selecting an agent

Distributed optimization problem  $\Rightarrow$  let's use a DCOP!

- $\mathcal{A}$  is the set of candidate agents  $A_c$
- $\mathcal{X}$  are the binary decision variables  $x_i^m$
- $\mathcal{C}$  are the constraints ensuring that all computations are hosted, agent's capacities are respected and hosting and communication costs are minimized

# Migrating computations

## Selecting an agent

$$\sum_{a_m \in A_c^i} x_i^m = 1 \quad (13)$$

$$\sum_{x_i \in X_c^m} \text{weight}(x_i) \cdot x_i^m + \sum_{x_j \in \mu^{-1}(a_m) \setminus X_c} \text{weight}(x_j) \leq \text{cap}(a_m) \quad (14)$$

$$\sum_{x_i \in X_c^m} \text{host}(a_m, x_i) \cdot x_i^m \quad (15)$$

$$\begin{aligned} & \sum_{(x_i, x_j) \in X_c^m \times N_i \setminus X_c} x_i^m \cdot \text{com}(i, j, m, \mu^{-1}(x_j)) \\ & + \sum_{(x_i, x_j) \in X_c^m \times N_i \cap X_c} x_i^m \cdot \sum_{a_n \in A_c^j} x_j^n \cdot \text{com}(i, j, m, n) \end{aligned} \quad (16)$$

## Decentralized reparation

When agents are removed:

- computation to migrate = computation that where hosted on these agents
- candidate agents = remaining agents that posses a replica of these orphaned computation

# Solving the migration DCOP

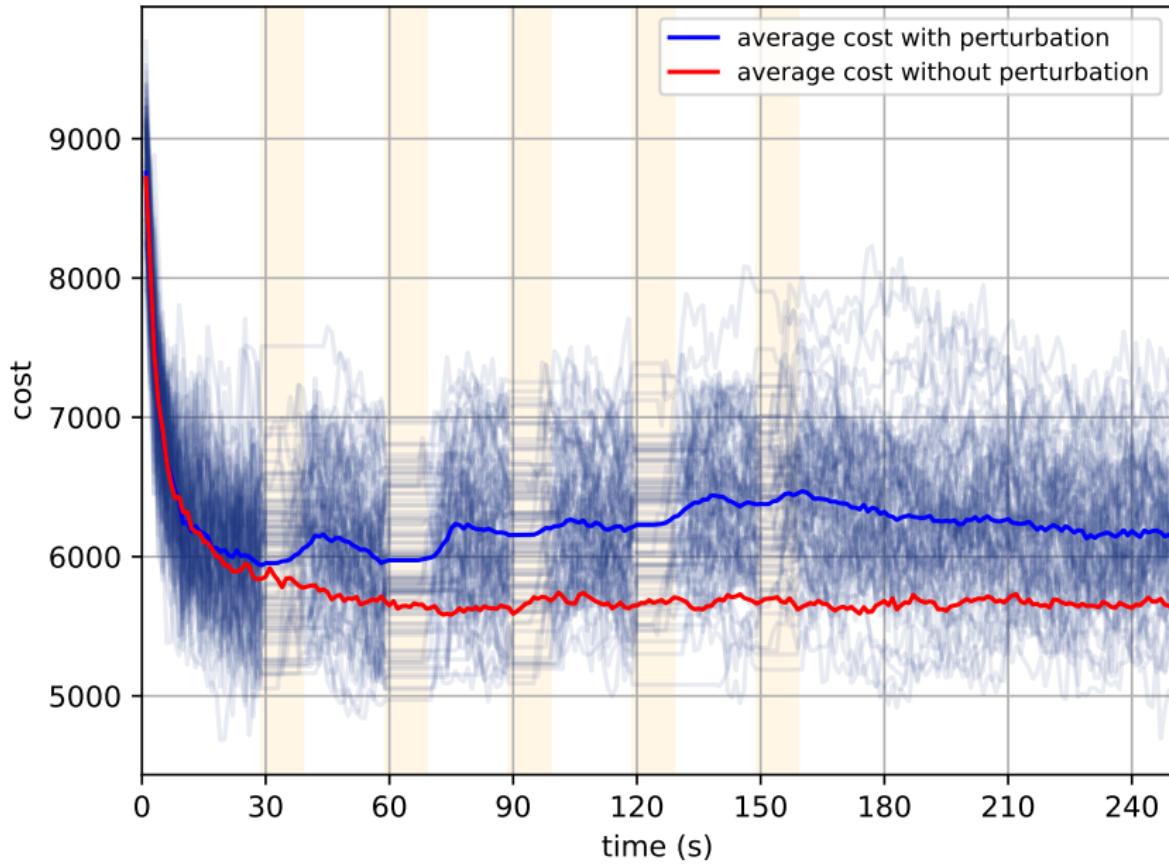
Which algorithm should we use ?

Criteria:

- lightweight
- fast (even if not optimal !)
- monotonic : mix of hard and soft constraints

MGM-2 : like MGM, with 2-coordination

## Experimental results



# Menu

DCOP Framework

Focus on Some Solution Methods

Hands on PyDCOP I

Focus on Smart Environment Configuration Problems

Distributing Computations

Hands on PyDCOP II

Dynamic DCOPs

Conclusion

# To sum up

What we've seen today:

- Some generic concepts
  - ▶ How to model coordination problems using DCOP formalism
  - ▶ Some solution methods (complete and incomplete) to solve DCOP
- Some specificities of IoT-based apps
  - ▶ How to model a specific smart environment configuration problem as a DCOP
  - ▶ How to use PyDCOP to model, run, solve, and distribute DCOP
  - ▶ How to equip a system with resilience using replication and DCOP-based reparation
- Want to go deeper into DCOPs → OPTMAS-DCR workshop series (AAMAS/IJCAI), other tutorials at AAMAS/IJCAI

# The End

# References

-  AJI, S.M. and R.J. McELIECE (2000). "The generalized distributive law". In: *Information Theory, IEEE Transactions on* 46.2, pp. 325–343. ISSN: 0018-9448. DOI: 10.1109/18.825794.
-  BOULLE, M. (2004). "Compact Mathematical Formulation for Graph Partitioning". In: *Optimization and Engineering* 5.3, pp. 315–333. ISSN: 1573-2924. DOI: 10.1023/B:OPTE.0000038889.84284.c7. URL: <http://dx.doi.org/10.1023/B:OPTE.0000038889.84284.c7>.
-  BÜRGER, M., G. NOTARSTEFANO, F. BULLO, and F. ALLGÖWER (2012). "A distributed simplex algorithm for degenerate linear programs and multi-agent assignments". In: *Automatica* 48.9, pp. 2298 –2304. ISSN: 0005-1098. DOI: <http://dx.doi.org/10.1016/j.automatica.2012.06.040>. URL: [//www.sciencedirect.com/science/article/pii/S0005109812002956](http://www.sciencedirect.com/science/article/pii/S0005109812002956).
-  FARINELLI, A., A. ROGERS, A. PETCU, and N. R. JENNINGS (2008). "Decentralised Coordination of Low-power Embedded Devices Using the Max-sum Algorithm". In: *Proceedings of the 7th International Joint Conference on Autonomous Agents and Multiagent Systems - Volume 2*. AAMAS '08. International Foundation for Autonomous Agents and Multiagent Systems, pp. 639–646. ISBN: 978-0-9817381-1-6.
-  FIORETTI, F., E. PONTELLI, and W. YEOH (2018). "Distributed Constraint Optimization Problems and Applications: A Survey". In: *Journal of Artificial Intelligence Research* 61, pp. 623–698.
-  FITZPATRICK, Stephen and Lambert MEERTENS (2003). "Distributed Coordination through Anarchic Optimization". In: *Distributed Sensor Networks: A Multiagent Perspective*. Ed. by Victor LESSER, Charles L. ORTIZ, and Milind TAMBE. Boston, MA: Springer US, pp. 257–295. ISBN: 978-1-4615-0363-7.
-  MAHESWARAN, R.T., J.P PEARCE, and M. TAMBE (2004). "Distributed Algorithms for DCOP: A Graphical-Game-Based Approach". In: *Proceedings of the 17th International Conference on Parallel and Distributed Computing Systems (PDCS)*, San Francisco, CA, pp. 432–439.
-  PETCU, Adrian and Boi FALTINGS (2005). "A scalable method for multiagent constraint optimization". In: *IJCAI International Joint Conference on Artificial Intelligence*, pp. 266–271. ISBN: 1045-0823.

## References (cont.)

-  RUST, P., G. PICARD, and F. RAMPARANY (2016). "Using Message-passing DCOP Algorithms to Solve Energy-efficient Smart Environment Configuration Problems". In: *International Joint Conference on Artificial Intelligence (IJCAI)*. AAAI Press.
-  — (2017). "On the Deployment of Factor Graph Elements to Operate Max-Sum in Dynamic Ambient Environments". In: *Autonomous Agents and Multiagent Systems – AAMAS 2017 Workshops, Best Papers, Sao Paulo, Brazil, May 8–12, 2017, Revised Selected Papers*. Ed. by G. SUKTHANKAR and J.A. RODRIGUEZ-AGUILAR. Vol. 10642. Lecture Notes in Artificial Intelligence (LNAI). Extended Version. Springer, pp. 116–137. DOI: 10.1007/978-3-319-71682-4\_8.
-  VINYALS, Meritxell, Juan A. RODRÍGUEZ-AGUILAR, and Jesus CERQUIDES (2011). "Constructing a unifying theory of dynamic programming DCOP algorithms via the generalized distributive law". In: *Autonomous Agents and Multi-Agent Systems* 3.22, pp. 439–464. ISSN: 1387-2532. DOI: 10.1007/s10458-010-9132-7.
-  WEISS, Yair (2000). "Correctness of Local Probability Propagation in Graphical Models with Loops". In: *Neural Comput.* 12.1, pp. 1–41. ISSN: 0899-7667. DOI: 10.1162/089976600300015880. URL: <http://dx.doi.org/10.1162/089976600300015880>.
-  ZHANG, W., G. WANG, Z. XING, and L. WITTENBURG (2005). "Distributed stochastic search and distributed breakout: properties, comparison and applications to constraint optimization problems in sensor networks.". In: *Journal of Artificial Intelligence Research (JAIR)* 161.1-2, pp. 55–87.
-  ZHANG, Weixiong, Guandong WANG, Zhao XING, and Lars WITTENBURG (2003). "A Comparative Study of Distributed Constraint Algorithms". In: *Distributed Sensor Networks: A Multiagent Perspective*. Ed. by Victor LESSER, Charles L. ORTIZ, and Milind TAMBE. Boston, MA: Springer US, pp. 319–338.