

# Distributed Constraint Optimization

Gauthier Picard

ONERA/DTIS

[gauthier.picard@onera.fr](mailto:gauthier.picard@onera.fr)

— Some contents taken from OPTMAS 2011 and OPTMAS-DCR 2014 Tutorials —

## Contents

## Introduction

## Constraint Optimization Problems

## DCOP Framework

## Application Domains

# Complete Algorithms for DCOP

## Asynchronous Distributed Optimisation (ADOPT)

# Distributed Pseudotree Optimization Procedure (DPOP)

## Approximate Algorithms for DCOP

## Distributed Stochastic Search Algorithm (DSA)

## Maximum Gain Message (MGM-1)

## Synthesis

## Panorama

# Constraint Optimization Problems

Sometimes satisfaction is not possible

- Overconstrained problem
  - Solution is not binary

## Switch from satisfaction to optimization

- Minimizing the number of violated constraints
  - Minimizing the cost of violated constraints
  - Maximizing the overall utility of the system

# DCOP Framework

## Motivations

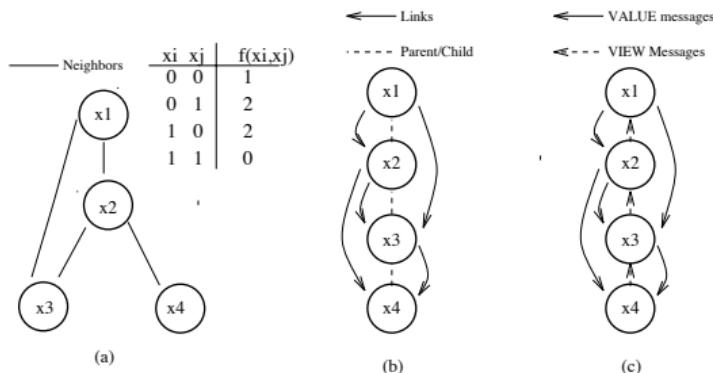
- In dynamic and complex environments not all constraints can be satisfied completely
- Satisfaction → **Optimisation** (combinatorial)
  - ▶ ex: minimizing the number of unchecked constraints, minimizing the sum of the costs of violated constraints, etc.

## Definition (DCOP)

A *DCOP* is a DCSP  $\langle A, X, D, C, \phi \rangle$  with

- a **cost function**  $f_{ij} : D_i \times D_j \mapsto \mathbb{N} \cup \infty$  for each pair  $x_i, x_j$
- an **objective function**  $F : D \mapsto \mathbb{N} \cup \infty$  evaluating an assignment  $\mathcal{A}$  with  $f_{ij}(d_i, d_j)$  for each pair  $x_i, x_j$

## DCOP Framework (cont.)

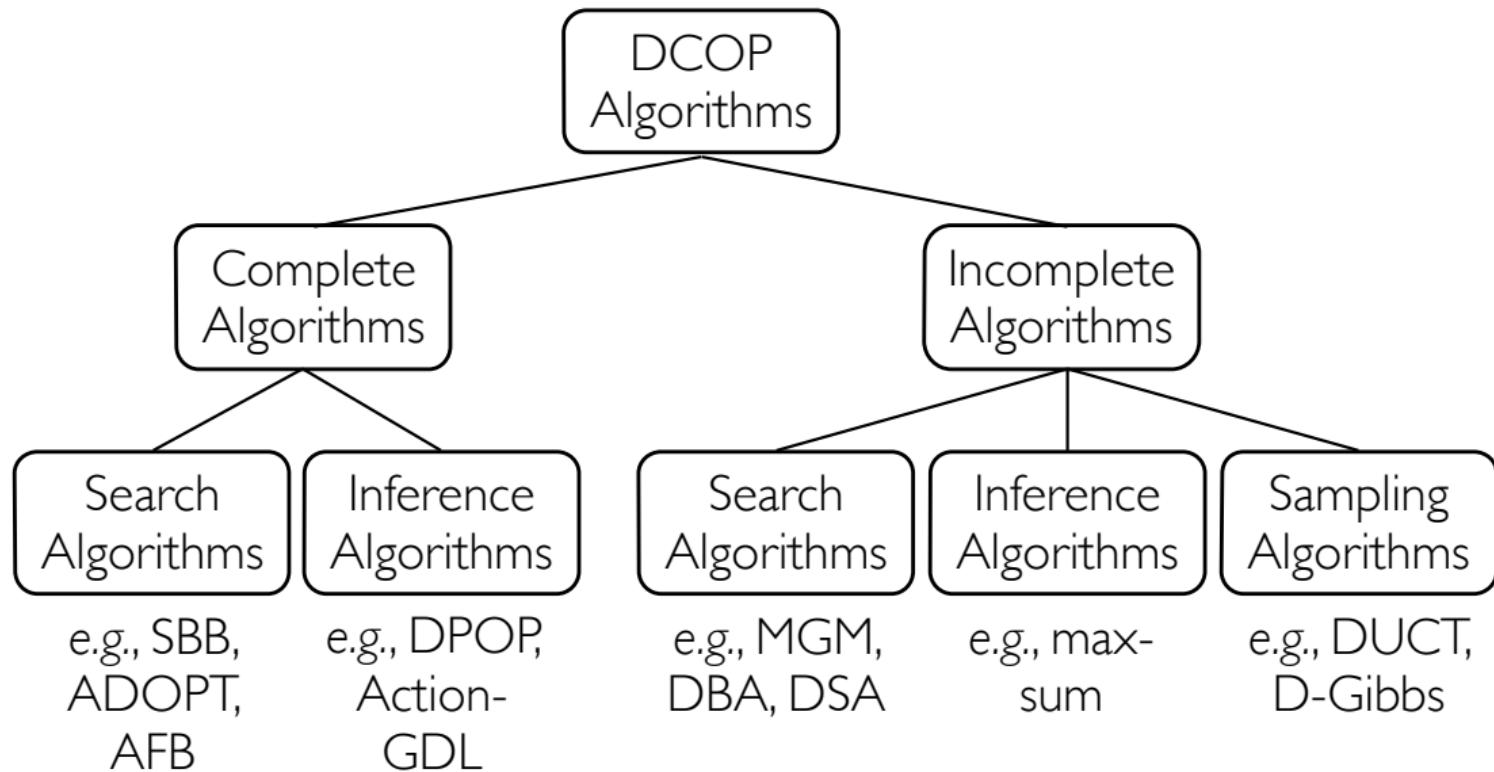


## Objective Function

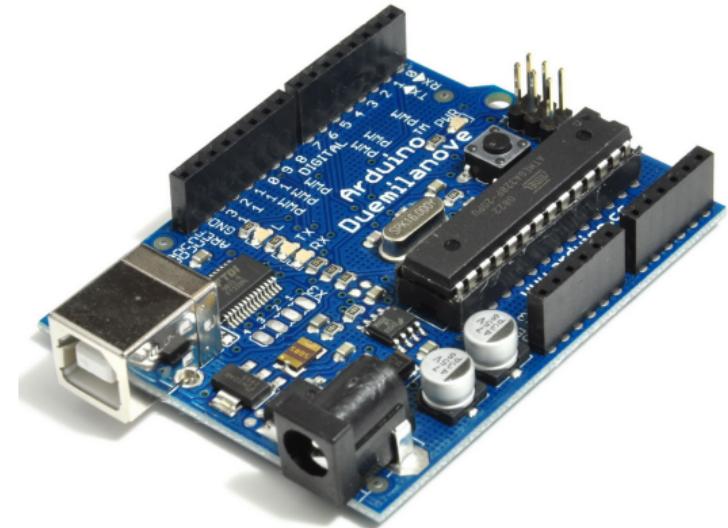
$$F(\mathcal{A}) = \sum_{x_i, x_j \in X} f_{ij}(d_i, d_j) \text{ where } x_i \leftarrow d_i \text{ and } x_i \leftarrow d_i \text{ in } \mathcal{A}$$

- $F(\{(x_1, 0), (x_2, 0), (x_3, 0), (x_4, 0)\}) = 4 \quad \mathcal{A}^* = \{(x_1, 1), (x_2, 1), (x_3, 1), (x_4, 1)\}$
  - $F(\{(x_1, 1), (x_2, 1), (x_3, 1), (x_4, 1)\}) = 0$

# DCOP Algorithms



# Application Domains



Distributed Constraint Optimization



Gauthier Picard

# Contents

Introduction

Complete Algorithms for DCOP

Asynchronous Distributed Optimisation (ADOPT)

Distributed Pseudotree Optimization Procedure (DPOP)

Approximate Algorithms for DCOP

Synthesis

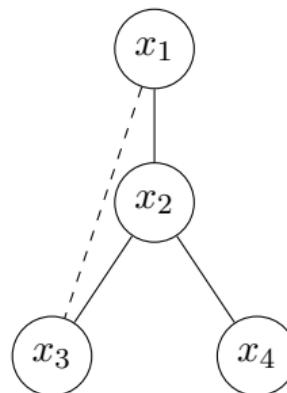
Asynchronous Distributed Optimisation (ADOPT) [MODI et al., 2005]

ADOPT: DFS tree (pseudotree)

ADOPT assumes that agents are arranged in a DFS tree:

- constraint graph → rooted graph (select a node as root)
  - some links form a tree / others are backedges
  - two constrained nodes must be in the same path to the root by tree links (same branch)

Every graph admits a DFS tree: DFS graph traversal



## ADOPT Features

- Asynchronous algorithm
  - Each time an agent receives a message:
    - ▶ Processes it (the agent may take a new value)
    - ▶ Sends VALUE messages to its children and pseudochildren
    - ▶ Sends a COST message to its parent
  - Context: set of (variable value) pairs (as ABT agent view) of ancestor agents (in the same branch)
  - Current context:
    - ▶ Updated by each VALUE message
    - ▶ If current context is not compatible with some child context, the later is initialized (also the child bounds)



# ADOPT Procedures

```

procedure backTrack
(38) if threshold == UB:
(39)   d1 ← d that minimizes UB(d);
(40) else if LB(d) > threshold:
(41)   d1 ← d that minimizes LB(d); endif;
(42)   SEND (VALUE, (x1, d1))
(43)     to each lower priority neighbor;
(44)   maintainAllocationInvariant;
(45) if threshold == UB:
(46)   if TERMINATE received from parent
(47)   or x1 is root:
(48)     SEND (TERMINATE,
(49)     CurrentContext ∪ {(x1, d1)})
(50)     to each child;
(51)   Terminate execution; endif.endif;
(52)   SEND (COST, x1, CurrentContext, LB, UB)
to parent;

when received (THRESHOLD, t, context)
(7)   if context compatible with CurrentContext:
(8)     threshold ← t;
(9)     maintainThresholdInvariant;
(10)    backTrack.endif;

when received (TERMINATE, context)
(11) record TERMINATE received from parent;
(12) CurrentContext ← context;
(13) backTrack;

when received (VALUE, (x1, d1))
(14) if TERMINATE not received from parent:
(15)   add (x1, d1) to CurrentContext;
(16)   forall d ∈ D1, x1 ∈ Children do
(17)     if context(d, x) incompatible with CurrentContext:
(18)       lb(d, x) ← 0; t(d, x) ← 0;
(19)       ub(d, x) ← Inf; context(d, x) ← [];
(20)     maintainThresholdInvariant;
(21)     backTrack.endif;

when received (COST, x1, context, lb, ub)
(22) d ← value of x1 in context;
(23) remove (x1, d) from context;
(24) if TERMINATE not received from parent:
(25)   forall (x1, d1) ∈ context and x1 is not my neighbor do
(26)     add (x1, d1) to CurrentContext.enddo;
(27)   forall d ∈ D1, x1 ∈ Children do
(28)     if context(d', x) incompatible with CurrentContext:
(29)       lb(d', x) ← 0; t(d', x) ← 0;
(30)       ub(d', x) ← Inf; context(d', x) ← [];endif.enddo.endif;
(31) if context compatible with CurrentContext:
(32)   lb(d, x) ← lb;
(33)   ub(d, x) ← ub;
(34)   context(d, x) ← context;
(35)   maintainChildThresholdInvariant;
(36)   maintainThresholdInvariant.endif;
(37) backTrack;

```

**Algorithm 1:** ADOPT Procedures

# ADOPT Messages

- **value**( $parent \rightarrow children \cup pseudochildren, a$ ): parent informs descendants that it has taken value  $a$
  - **cost**( $child \rightarrow parent, lowerbound, upperbound, context$ ): child informs parent of the best cost of its assignment; attached context to detect obsolescence
  - **threshold**( $parent \rightarrow child, t$ ): minimum cost of solution in child is at least  $t$
  - **termination**( $parent \rightarrow children$ ): sent when  $LB = UB$

# ADOPT Data Structures

- 1. Current context** (agent view): values of higher priority constrained agents

$x_i$	$x_j$	$\dots$
$a$	$c$	$\dots$

- ## 2. Bounds (for each value, child)

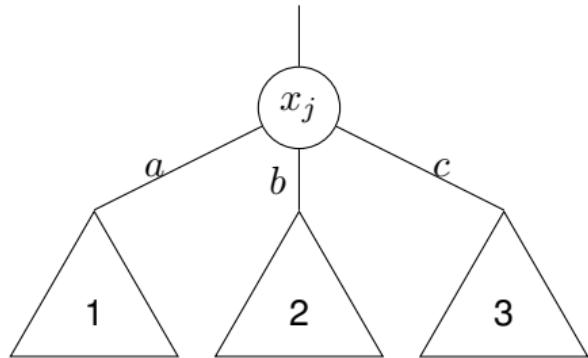
- ▶ lower bounds
  - ▶ upper bounds
  - ▶ thresholds
  - ▶ contexts

$x_j$   
 $lb(x_k)$   
 $ub(x_k)$   
 $th(x_k)$

$a$	$b$	$c$	$d$
3	0	0	0
$\infty$	$\infty$	$\infty$	$\infty$
1	0	0	0

- Stored contexts must be active:  $context \in currentcontext$
  - If a context becomes no active, it is removed ( $lb \leftarrow 0, th \leftarrow 0, ub \leftarrow \infty$ )

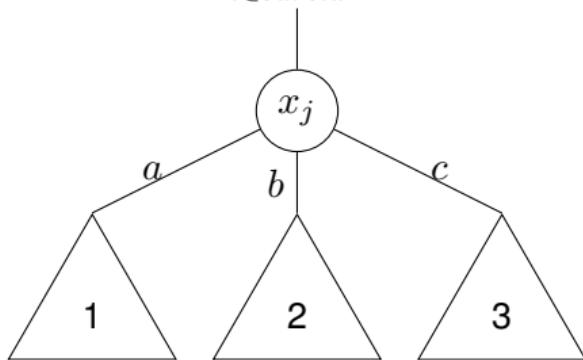
## ADOPT Bounds



## ADOPT Bounds

$\delta(\text{value}) = \text{cost with higher agents}$

$$\delta(b) = \sum_{i \in curctx} c_{ij}(a, b)$$

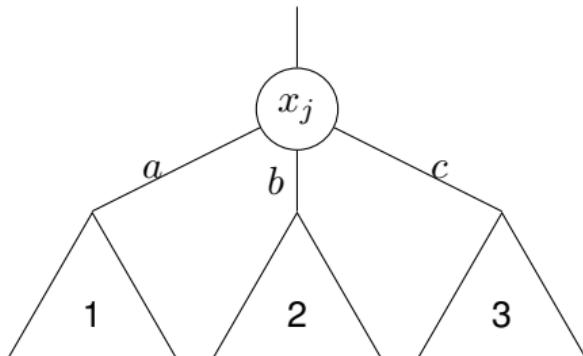


# ADOPT Bounds

$\delta(\text{value}) = \text{cost with higher agents}$

$$\delta(b) = \sum_{i \in curctx} c_{ij}(a, b)$$

$$OPT(x_j, ctx) = \min_{d \in d_j} \delta(d) + \sum_{x_k \in children} OPT(x_k, ctx \cup (x_j, d))$$

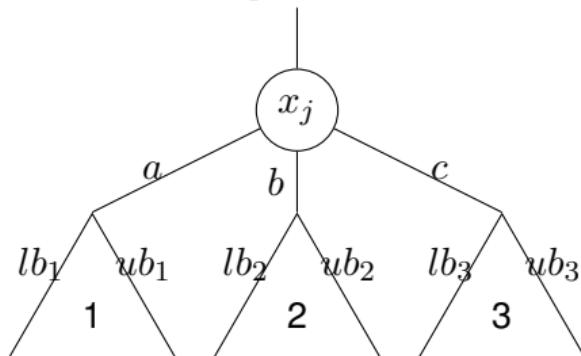


## ADOPT Bounds

$\delta(\text{value}) = \text{cost with higher agents}$

$$\delta(b) = \sum_{i \in curctx} c_{ij}(a, b)$$

$$OPT(x_j, ctx) = \min_{d \in d_j} \delta(d) + \sum_{x_k \in children} OPT(x_k, ctx \cup (x_j, d))$$

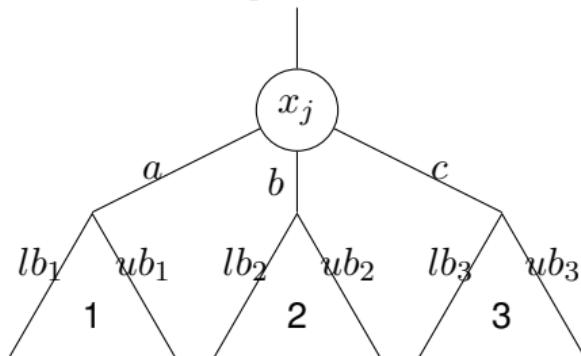


# ADOPT Bounds

$\delta(\text{value}) = \text{cost with higher agents}$

$$\delta(b) = \sum_{i \in curctx} c_{ij}(a, b)$$

$$OPT(x_j, ctx) = \min_{d \in d_j} \delta(d) + \sum_{x_k \in children} OPT(x_k, ctx \cup (x_j, d))$$

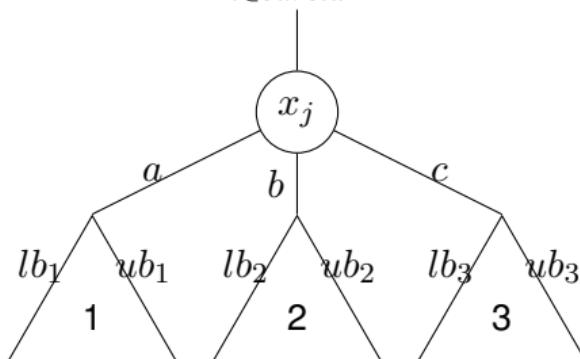


$[lb_k, ub_k]$  = cost of lower agents

## ADOPT Bounds

$\delta(\text{value}) = \text{cost with higher agents}$

$$\delta(b) = \sum_{i \in curctx} c_{ij}(a, b)$$



$[lb_k, ub_k] = \text{cost of lower agents}$

$$OPT(x_j, ctx) = \min_{d \in d_j} \delta(d) + \sum_{x_k \in children} OPT(x_k, ctx \cup (x_j, d))$$

$$LB(b) = \delta(b) + \sum_{x_k \in children} lb(b, x_k)$$

$$LB = \min_{b \in d_j} LB(b)$$

$$UB(b) = \delta(b) + \sum_{x_k \in children} ub(b, x_k)$$

$$UB = \min_{b \in d_j} UB(b)$$

# ADOPT Value Assignment

- An ADOPT agent takes the value with minimum LB
- Eager behavior:
  - ▶ Agents may constantly change value
  - ▶ Generates many context changes
- Threshold:
  - ▶ lower bound of the cost that children have from previous search
  - ▶ parent distributes threshold among children
  - ▶ incorrect distribution does not cause problems: the child with minor allocation would send a COST to the parent later, and the parent will rebalance the threshold distribution

## ADOPT Properties

- For any  $x_i$ ,  $LB \leq OPT(x_l, ctx) \leq UB$
  - For any  $x_i$ , its threshold reaches UB
  - For any  $x_i$ , its final threshold is equal to  $OPT(x_l, ctx)$

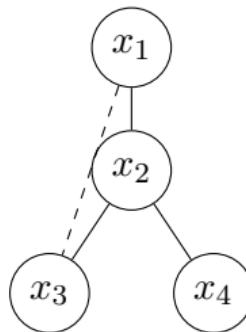
→ ADOPT terminates with the optimal solution

## ADOPT Example

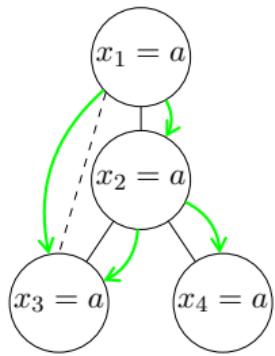
- 4 variables (4 agents)  $x_1, x_2, x_3$  and  $x_4$  with  $D = \{a, b\}$
  - 4 binary identical cost functions

$x_i$	$x_j$	cost
a	a	1
a	b	2
b	a	2
b	b	0

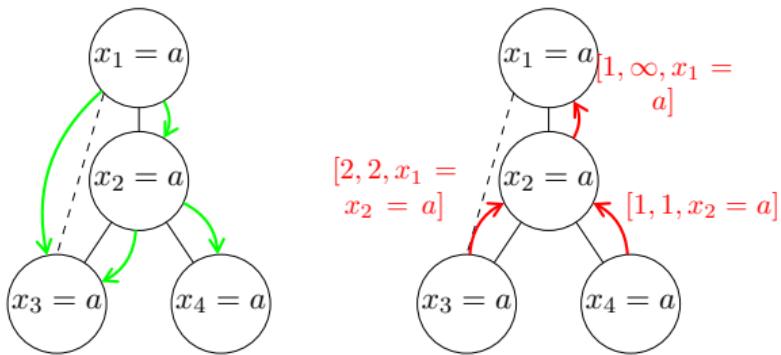
- ## ■ Constraint graph



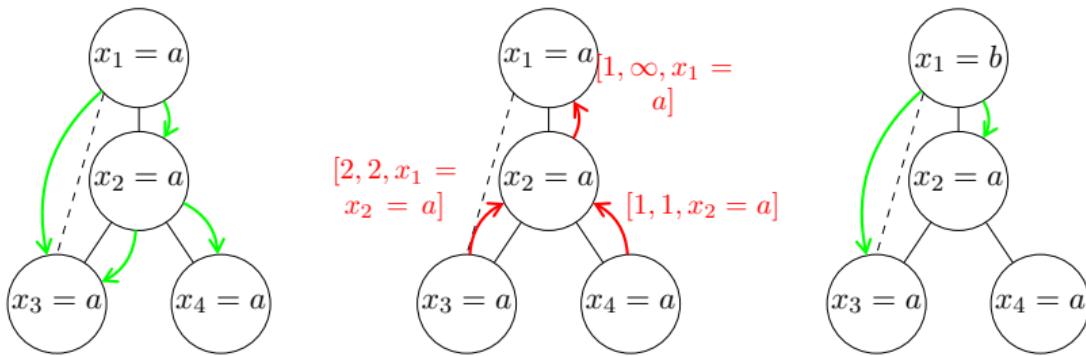
## ADOPT Example (cont.)



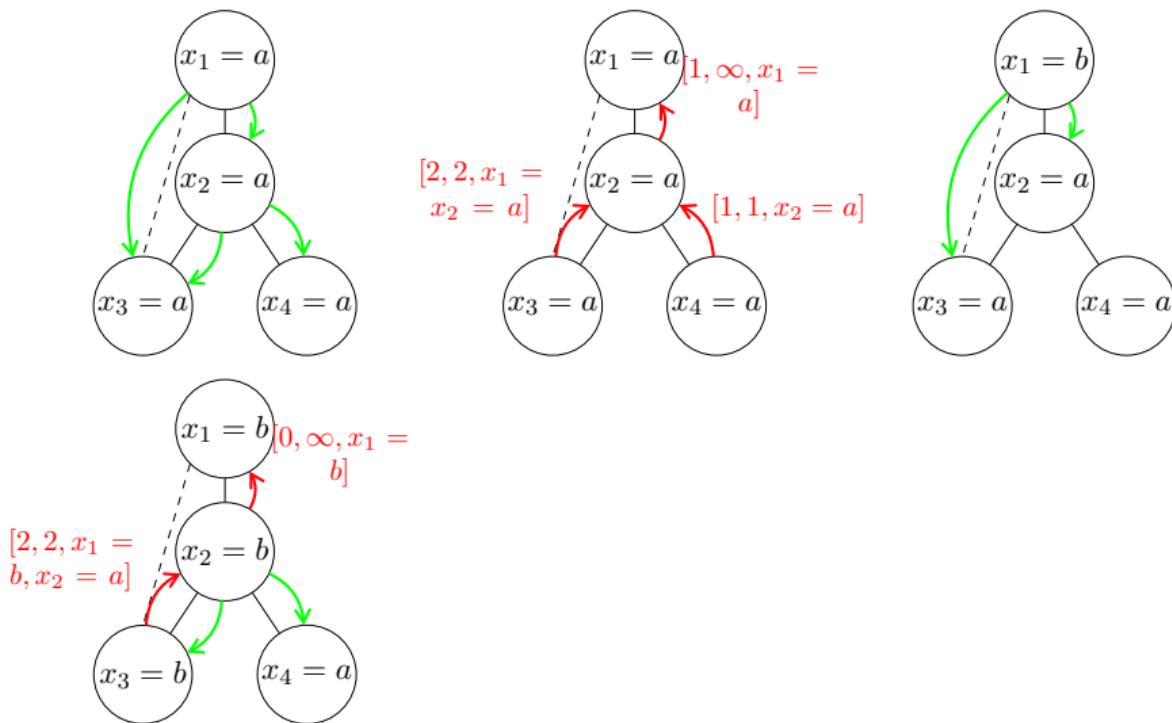
## ADOPT Example (cont.)



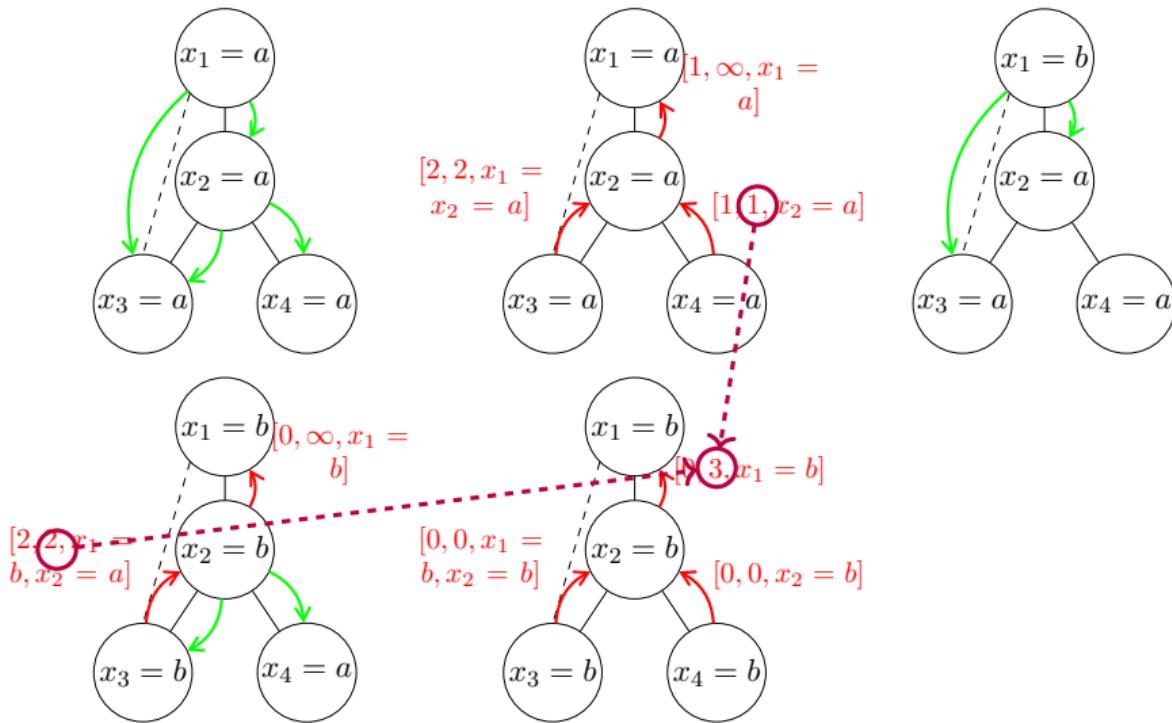
## ADOPT Example (cont.)



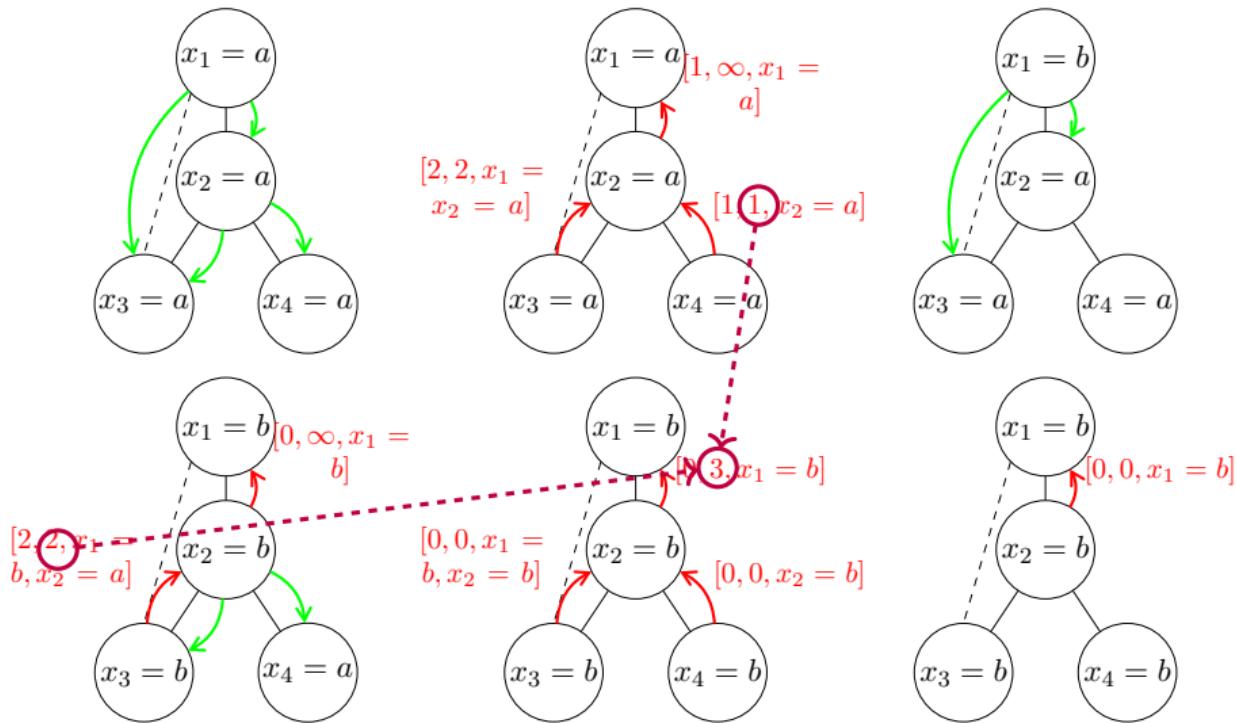
## ADOPT Example (cont.)



## ADOPT Example (cont.)



## ADOPT Example (cont.)



# Distributed Pseudotree Optimization Procedure (DPOP) [PETCU and FALTINGS, 2005]

## 3-phase distributed algorithm

PHASES	MESSAGES
1. DFS Tree construction	token passing
2. Utility phase: from leaves to root	<b>util</b> (child → parent, constraint table [-child])
3. Value phase: from root to leaves	<b>value</b> (parent → children ∪ pseudochildren, parent value)

## DFS Tree Phase

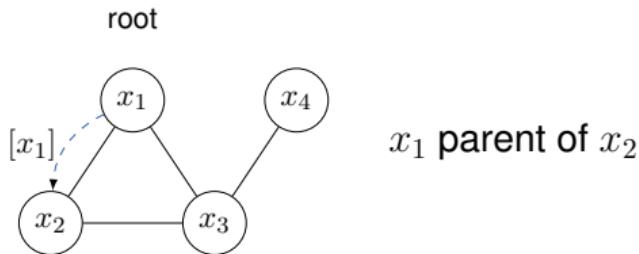
### ■ Distributed DFS graph traversal: token, ID, $\text{neighbors}(X)$

1.  $X$  owns the token: adds its own  $ID$  and sends it in turn to each of its neighbors, which become children
2.  $Y$  receives the token from  $X$ : it marks  $X$  as visited. First time  $Y$  receives the token then  $\text{parent}(Y) = X$ . Other IDs in token which are also  $\text{neighbors}(Y)$  are **pseudoparent**. If  $Y$  receives token from neighbor  $W$  to which it was never sent,  $W$  is pseudochild.
3. When all  $\text{neighbors}(X)$  visited,  $X$  removes its ID from token and sends it to  $\text{parent}(X)$ .

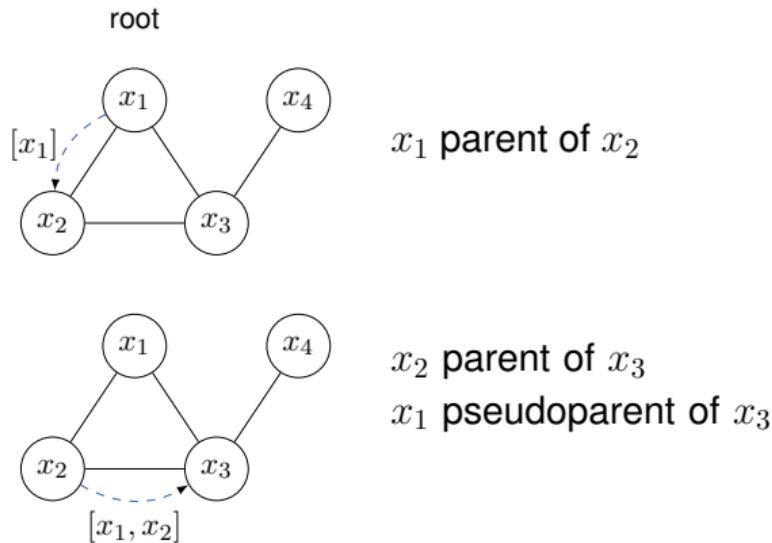
### ■ A node is selected as root, which starts

### ■ When all neighbors of root are visited, the DFS traversal ends

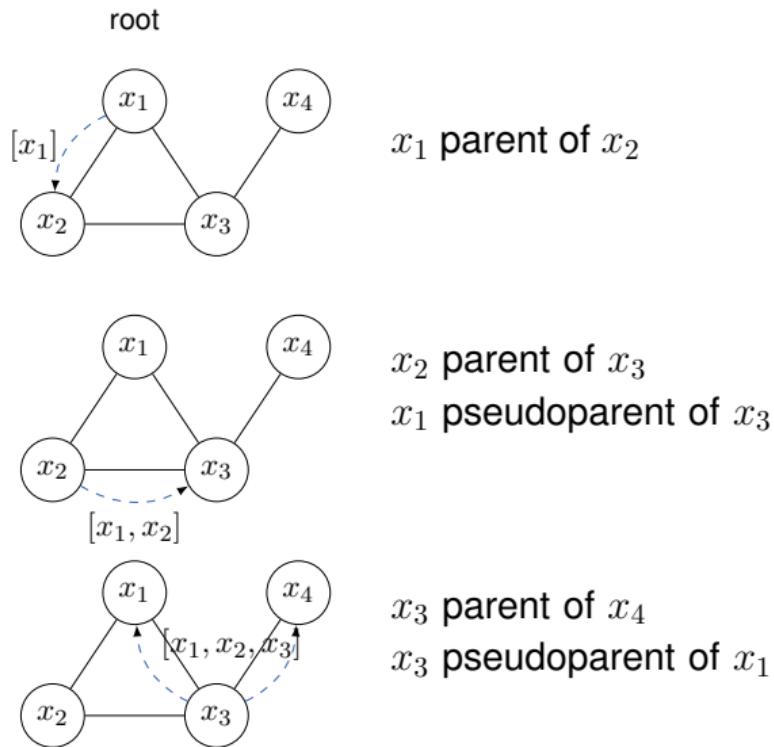
## DFS Tree Phase: Example



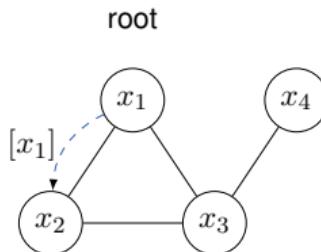
## DFS Tree Phase: Example



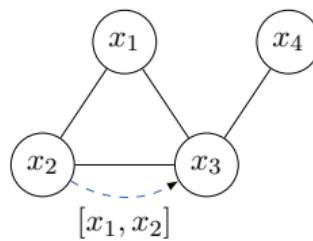
## DFS Tree Phase: Example



## DFS Tree Phase: Example

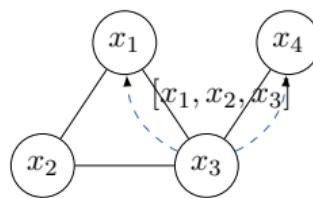


$x_1$  parent of  $x_2$



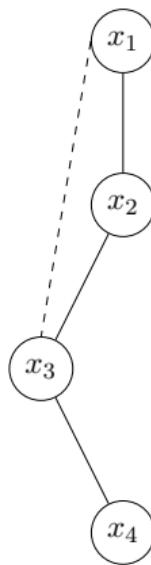
$x_2$  parent of  $x_3$

$x_1$  pseudoparent of  $x_3$



$x_3$  parent of  $x_4$

$x_3$  pseudoparent of  $x_1$



# Util Phase

Agent  $X$ :

- receives from each child  $Y_i$  a cost function:  $C(Y_i)$
- combines (adds, joins) all these cost functions with the cost functions with  $parent(X)$  and  $pseudoparents(X)$
- projects  $X$  out of the resulting cost function, and sends it to  $parent(X)$

From the leaves to the root

## Util Phase: Example



## Util Phase: Example

$X$	$T$	
a	a	1
a	b	2
b	a	2
b	b	0

*parent*

X

*children*

$X$	$Y$	
a	a	1
a	b	2
b	a	2
b	b	0

$X$	$Z$	
a	a	1
a	b	2
b	a	2
b	b	0

## Util Phase: Example

The diagram illustrates a DCOP network with four variables: X, Y, Z, and T. Variable X is the root node, labeled 'parent' with a vertical line above it and 'children' with lines connecting to Y, Z, and T. A red dashed arrow labeled 'add' points from the X domain to the Y domain. The domains are represented as tables:

X	T	
a	a	1
a	b	2
b	a	2
b	b	0

X	Y	Z	T	
a	a	a	a	3
a	a	a	b	4
a	a	b	a	4
a	a	b	b	5
a	b	a	a	4
a	b	a	b	5
a	b	b	a	5
a	b	b	b	6
b	a	a	a	6
b	a	a	b	4
b	a	b	a	4
b	a	b	b	2
b	b	a	a	4
b	b	a	b	2
b	b	b	a	2
b	b	b	b	0

X	Y	
a	a	1
a	b	2
b	a	2
b	b	0

X	Z	
a	a	1
a	b	2
b	a	2
b	b	0

All value combinations  
Costs are the sum of  
applicable costs

## Util Phase: Example

*parent*

*children*

X	T
a	a
a	b
b	a
b	b

X	Y
a	a
a	b
b	a
b	b

X	Z
a	a
a	b
b	a
b	b

X	Y	Z	T
a	a	a	3
a	a	b	4
a	b	a	4
a	b	b	5
a	b	a	4
a	b	a	5
a	b	b	6
b	a	a	6
b	a	b	4
b	a	a	4
b	a	b	2
b	b	a	4
b	b	b	2
b	b	a	2
b	b	b	0

All value combinations  
Costs are the sum of applicable costs

*Project out X*

a	a	a	a	3
a	a	a	b	4
a	a	b	a	4
a	a	b	b	5
a	b	a	a	4
a	b	a	b	5
a	b	b	a	5
a	b	b	b	6
b	a	a	a	6
b	a	a	b	4
b	a	b	a	4
b	a	b	b	2
b	b	a	a	4
b	b	a	b	2
b	b	b	a	2
b	b	b	b	0

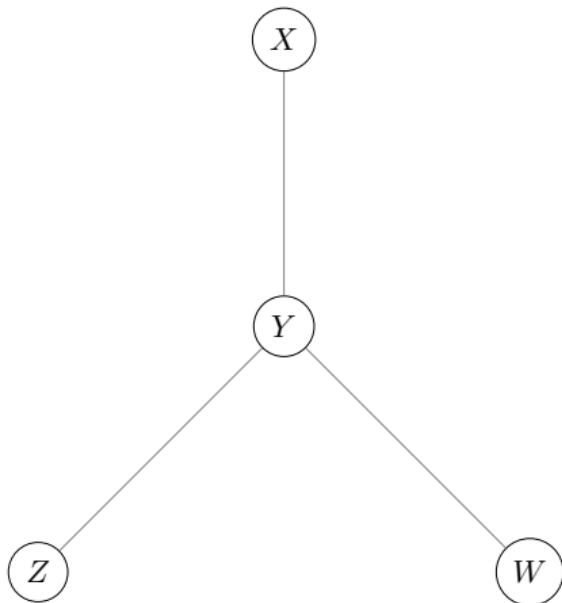
Remove X  
Remove duplicates  
Keep the min cost

## Value Phase

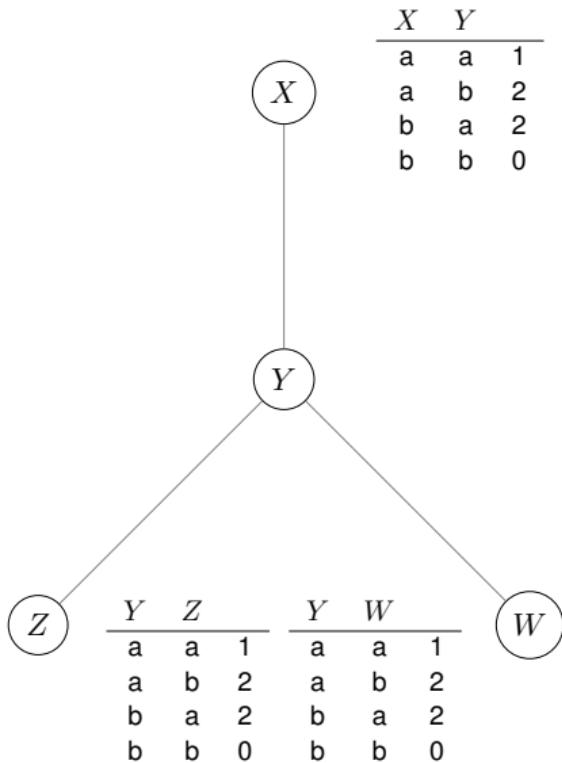
1. The root finds the **value that minimizes the received cost function** in the util phase, and informs its descendants ( $\text{children} \cup \text{pseudochildren}$ )
2. Each agent **waits to receive** the value of its parent / pseudoparents
3. Keeping fixed the value of parent/pseudoparents, finds the **value that minimizes the received cost function** in the Util phase
4. Informs of this value to its children/pseudochildren

This process starts at the root and ends at the leaves

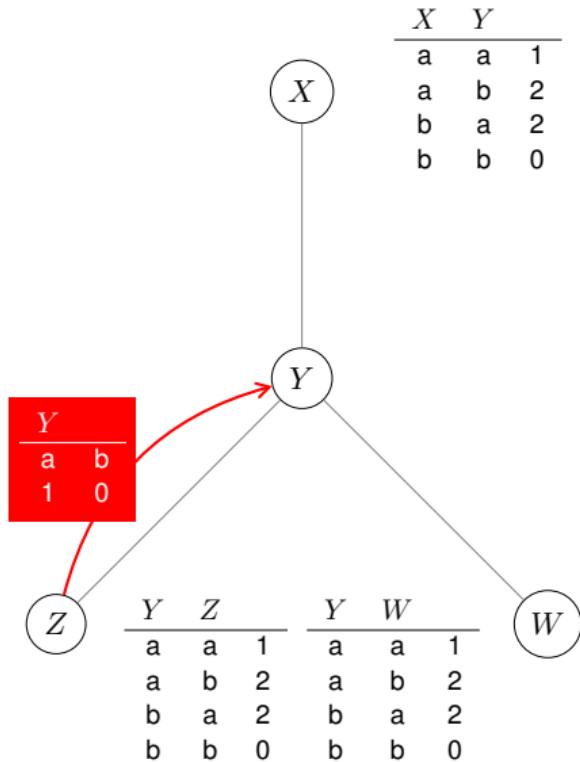
## DTREE : DPOP for DCOPs without backedges



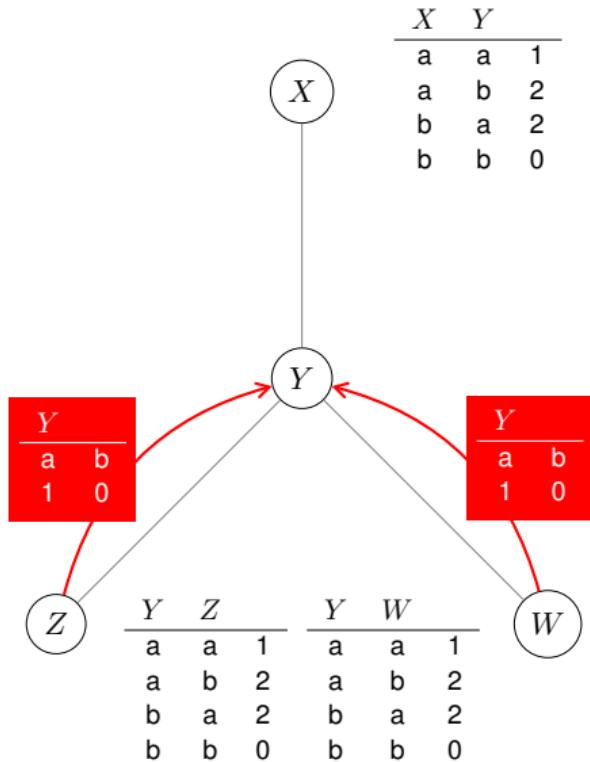
## DTREE : DPOP for DCOPs without backedges



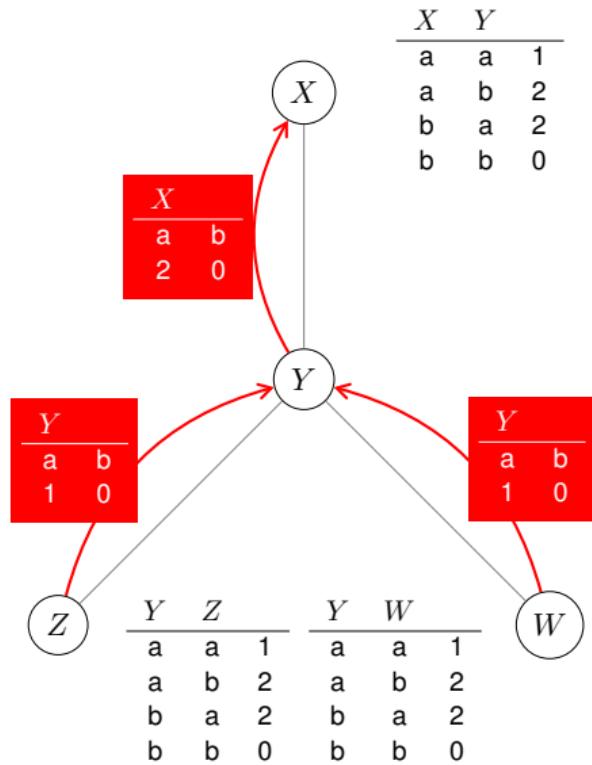
## DTREE : DPOP for DCOPs without backedges



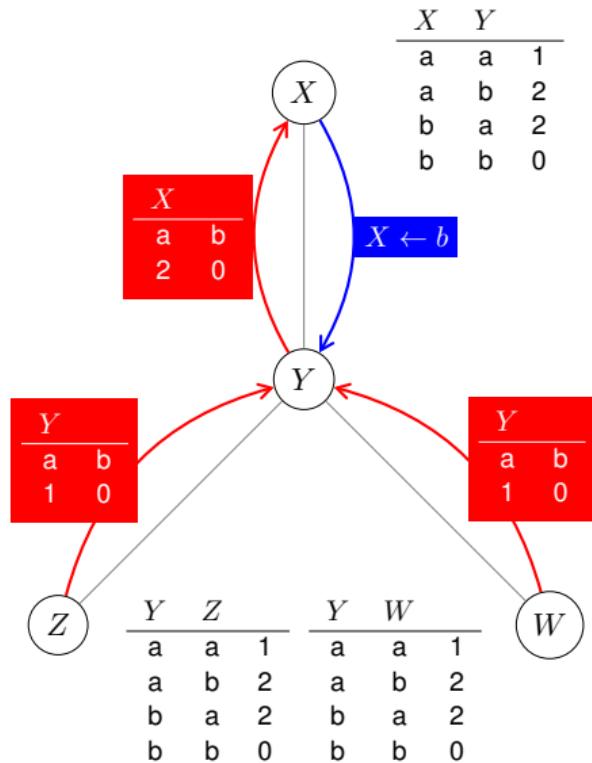
## DTREE : DPOP for DCOPs without backedges



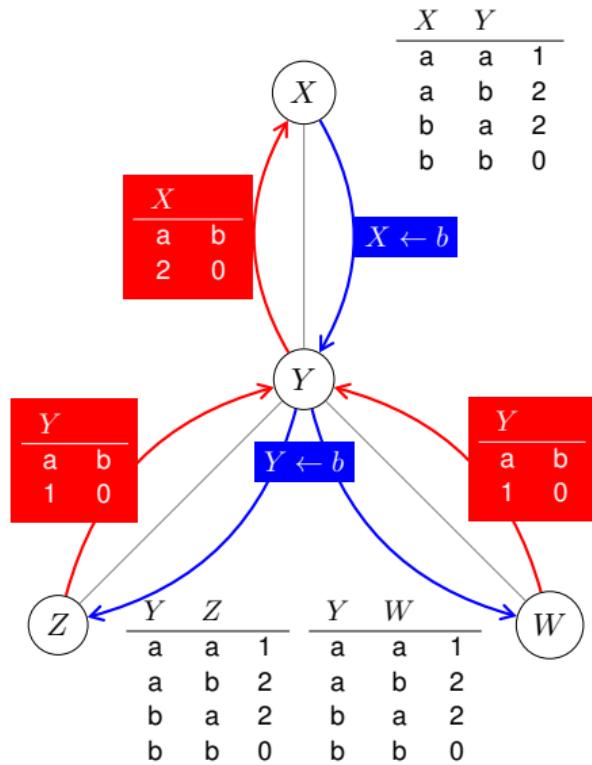
## DTREE : DPOP for DCOPs without backedges



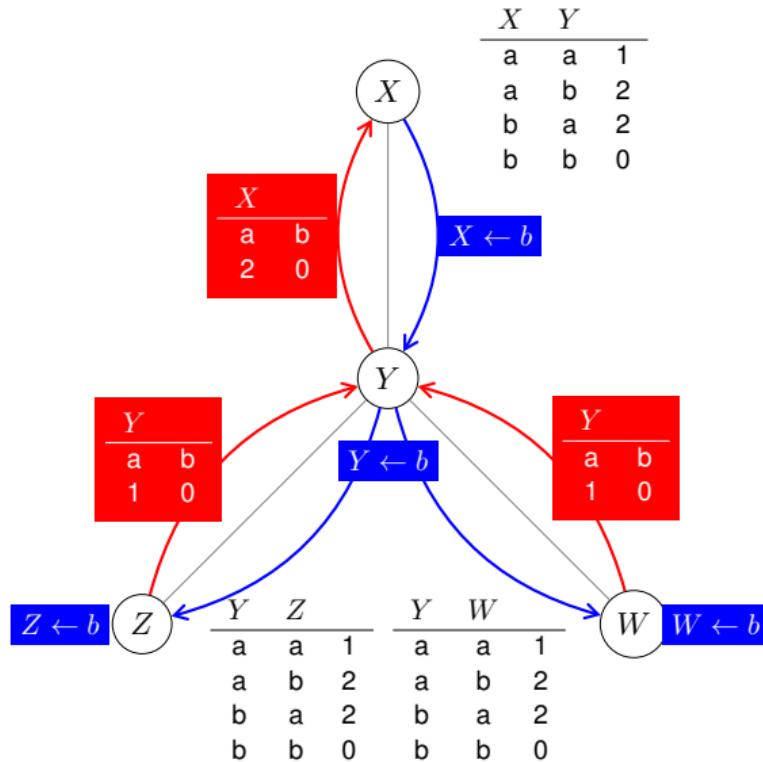
## DTREE : DPOP for DCOPs without backedges



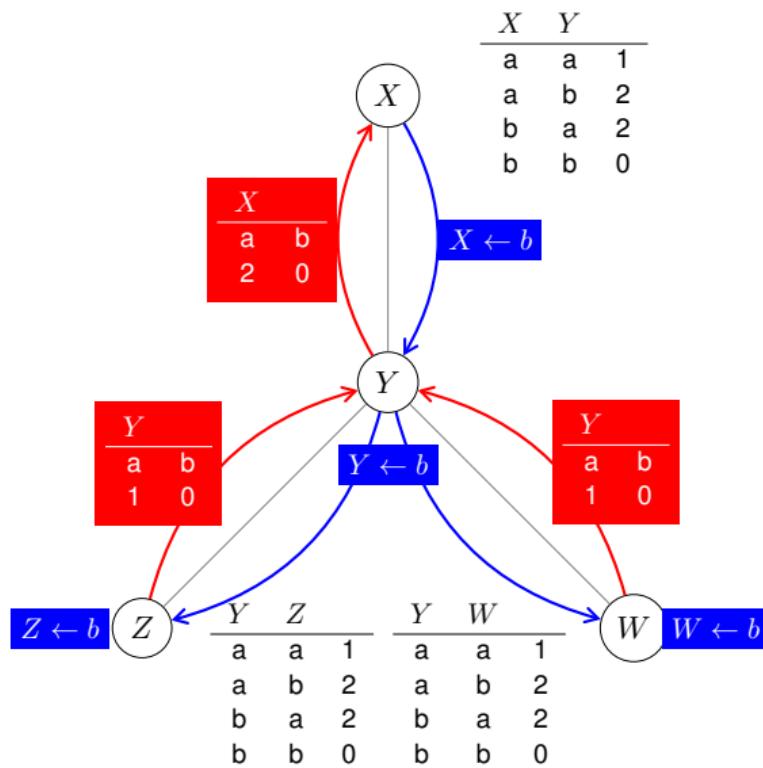
## DTREE : DPOP for DCOPs without backedges



## DTREE : DPOP for DCOPs without backedges



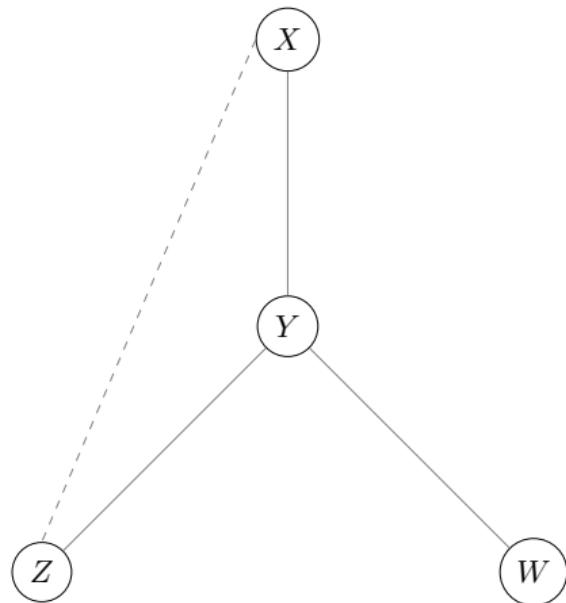
## DTREE : DPOP for DCOPs without backedges



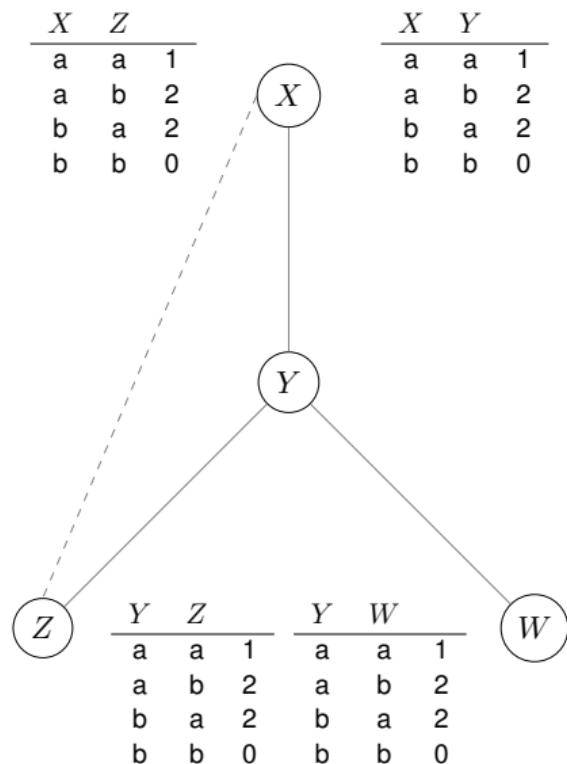
## Optimal solution:

- linear number of messages
  - message size: linear

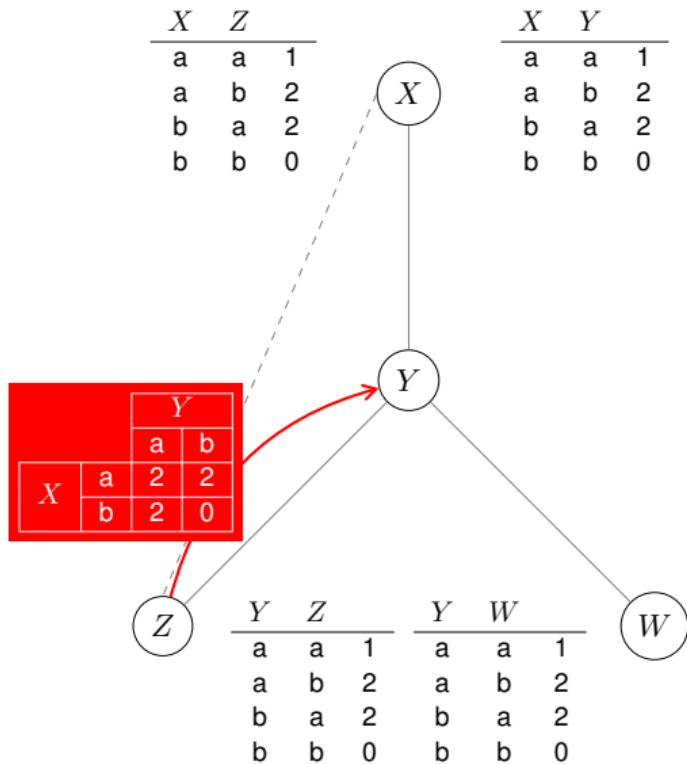
## DPOP for any DCOP



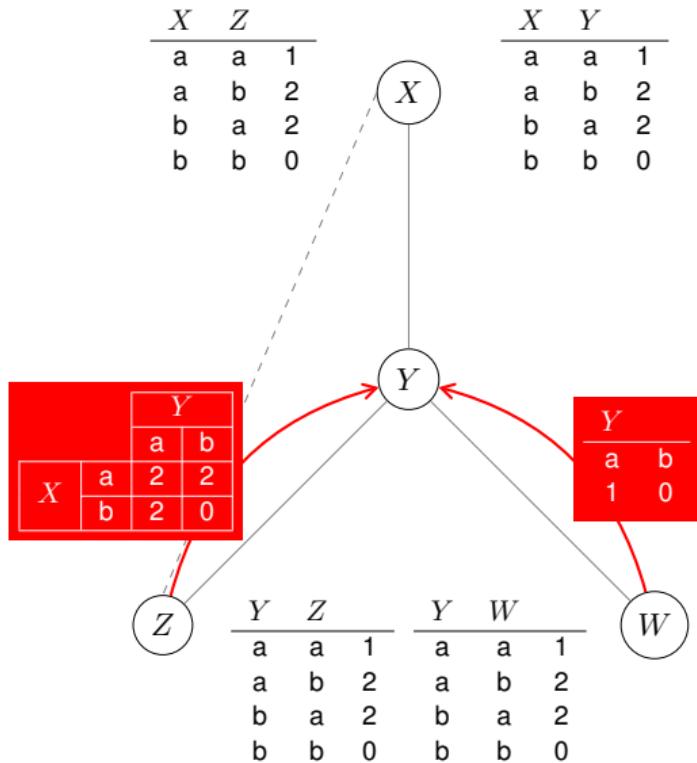
## DPOP for any DCOP



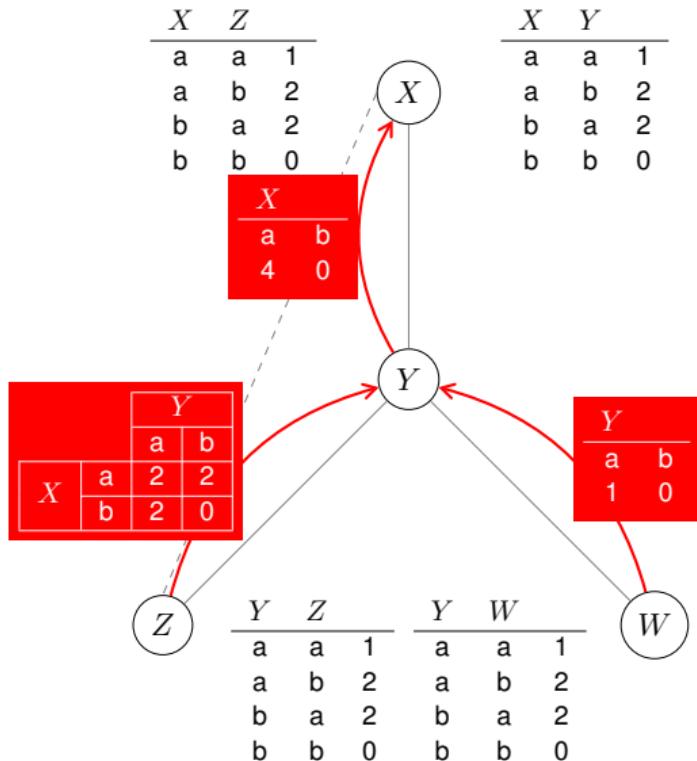
## DPOP for any DCOP



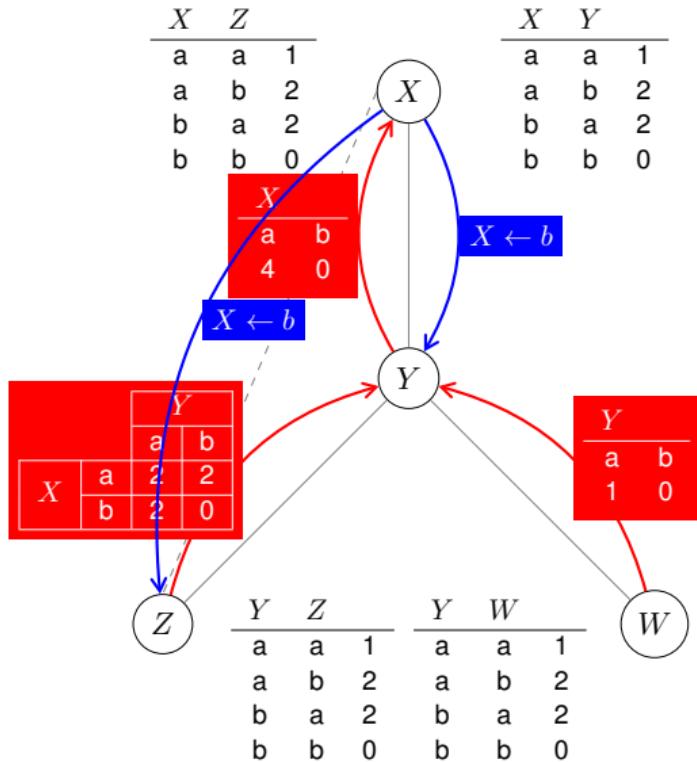
## DPOP for any DCOP



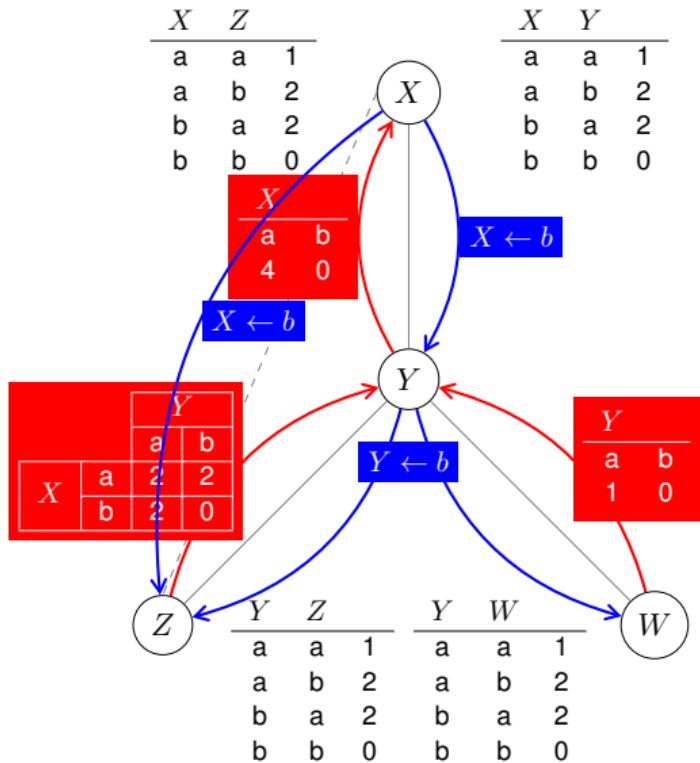
## DPOP for any DCOP



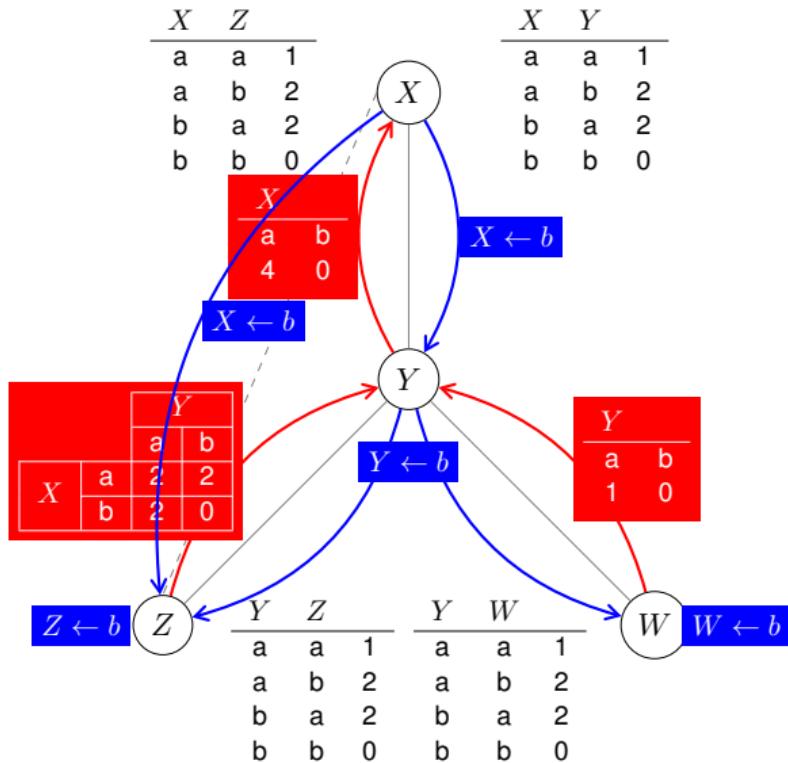
## DPOP for any DCOP



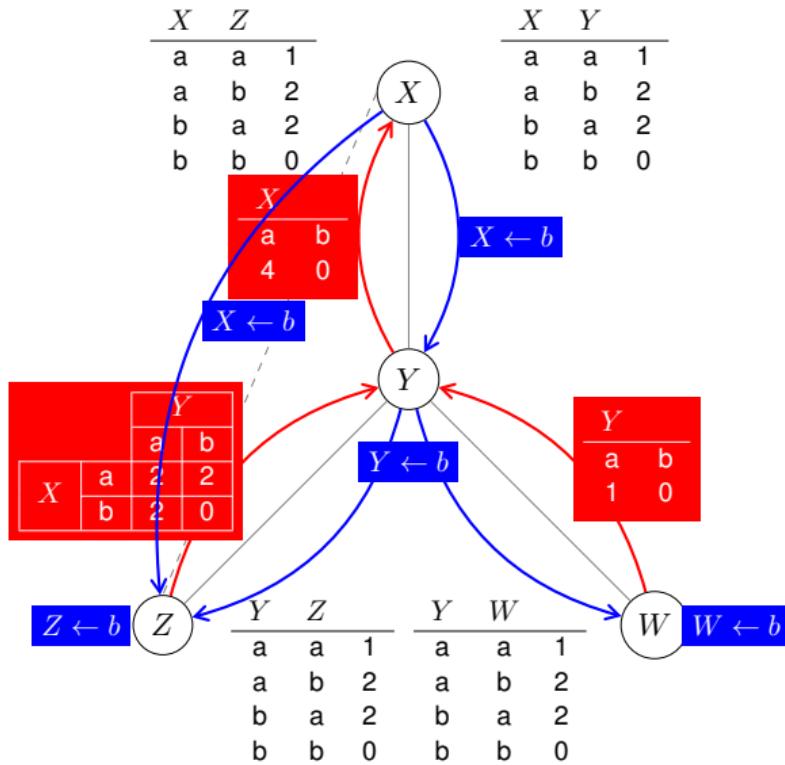
## DPOP for any DCOP



## DPOP for any DCOP



## DPOP for any DCOP



Optimal solution:

- linear number of messages
- message size: exponential

# Contents

Introduction

Complete Algorithms for DCOP

Approximate Algorithms for DCOP

Distributed Stochastic Search Algorithm (DSA)

Maximum Gain Message (MGM-1)

Synthesis

# Approximate Algorithms for DCOPs

## Complete algorithms

- e.g. ADOPT [MODI et al., 2005] and DPOP [PETCU and FALTINGS, 2005]
  - ✓ complete
  - ✗ slow

Approximate algorithms exist (fast, but sub-optimal in many cases)

- Search algorithms
  - ▶ DBA [YOKOO, 2001], DSA [ZHANG et al., 2005], MGM [MAHESWARAN et al., 2004]
- Inference algorithms
  - ▶ Max-sum [FARINELLI et al., 2008]

## Why Approximate Algorithms

## ■ Motivations

- ▶ Often optimality in practical applications is not achievable
  - ▶ Fast good enough solutions are all we can have

## ■ Example – Graph coloring

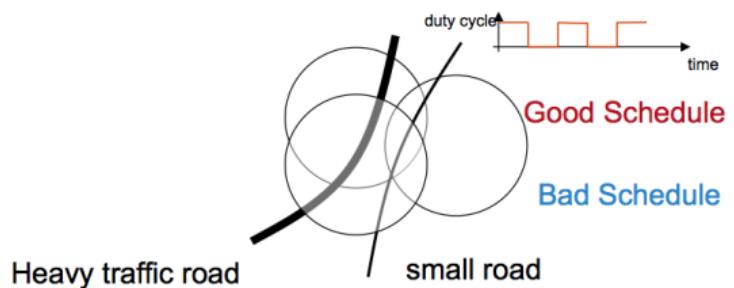
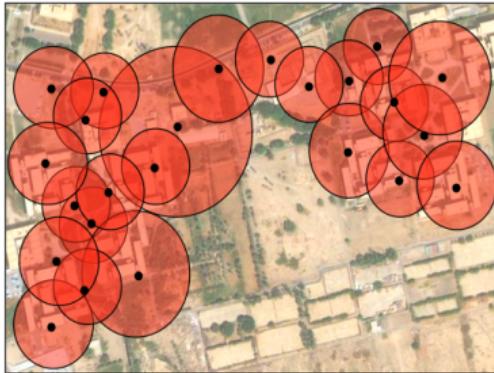
- ▶ Medium size problem (about 20 nodes, three colors per node)
  - ▶ Number of states to visit for optimal solution in the worst case  $3^{20} = 3M$  states

## ■ Key problem

- ▶ Provides guarantees on solution quality

# Exemplar Application: Surveillance

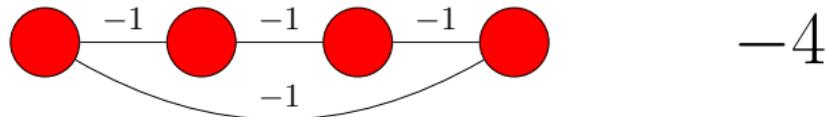
- Event Detection
  - ▶ Vehicles passing on a road
- Energy Constraints
  - ▶ Sense/Sleep modes
  - ▶ Recharge when sleeping
- Coordination
  - ▶ Activity can be detected by single sensor
  - ▶ Roads have different traffic loads
- Aim
  - ▶ Focus on road with more traffic load



## Centralized Local Greedy approaches

## ■ Greedy local search

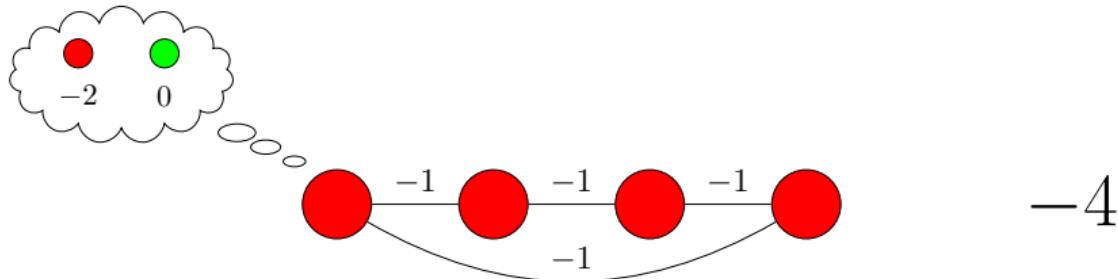
- ▶ Start from random solution
  - ▶ Do **local** changes if global solution improves
  - ▶ **Local**: change the value of a subset of variables, usually one



# Centralized Local Greedy approaches

## ■ Greedy local search

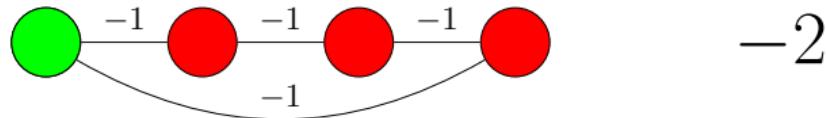
- ▶ Start from random solution
- ▶ Do **local** changes if global solution improves
- ▶ **Local**: change the value of a subset of variables, usually one



# Centralized Local Greedy approaches

## ■ Greedy local search

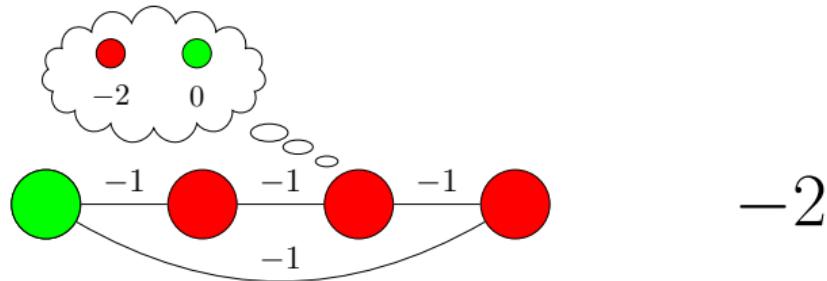
- ▶ Start from random solution
- ▶ Do **local** changes if global solution improves
- ▶ **Local:** change the value of a subset of variables, usually one



# Centralized Local Greedy approaches

## ■ Greedy local search

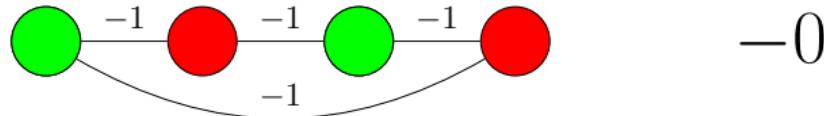
- ▶ Start from random solution
- ▶ Do **local** changes if global solution improves
- ▶ **Local**: change the value of a subset of variables, usually one



# Centralized Local Greedy approaches

## ■ Greedy local search

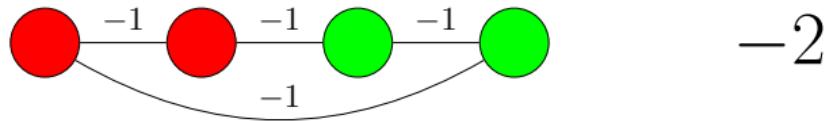
- ▶ Start from random solution
- ▶ Do **local** changes if global solution improves
- ▶ **Local**: change the value of a subset of variables, usually one



## Centralized Local Greedy approaches

## ■ Problems

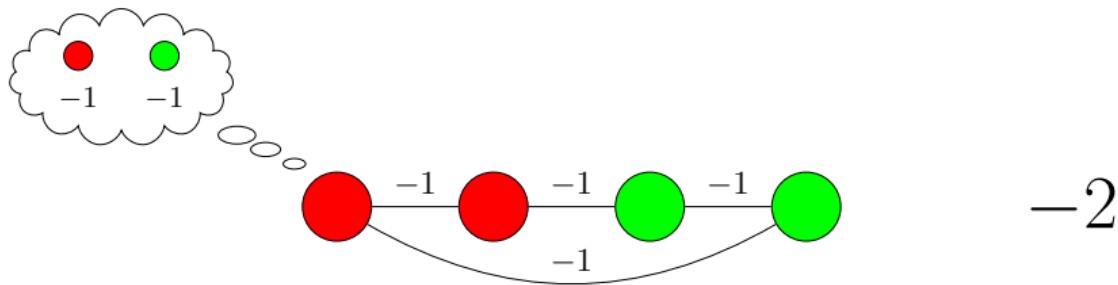
- ▶ Local minima
  - ▶ Standard solutions: Random Walk, Simulated Annealing



## Centralized Local Greedy approaches

## ■ Problems

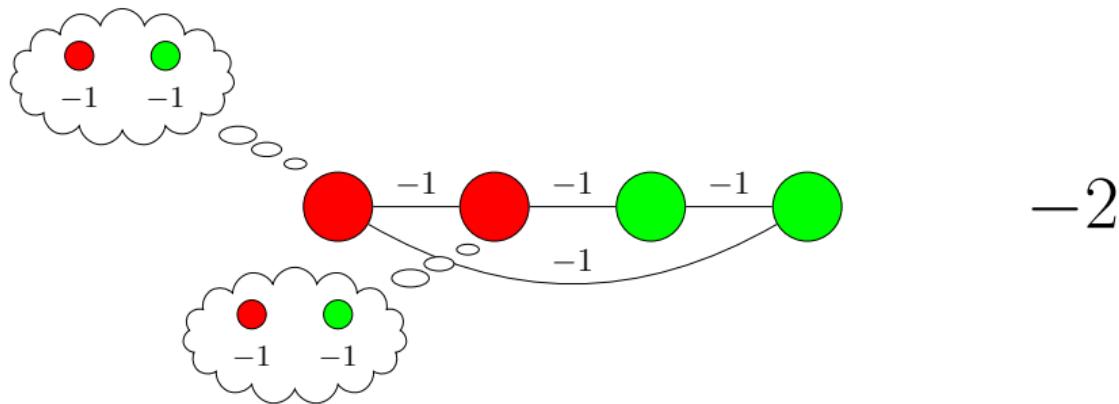
- ▶ Local minima
  - ▶ Standard solutions: Random Walk, Simulated Annealing



# Centralized Local Greedy approaches

## ■ Problems

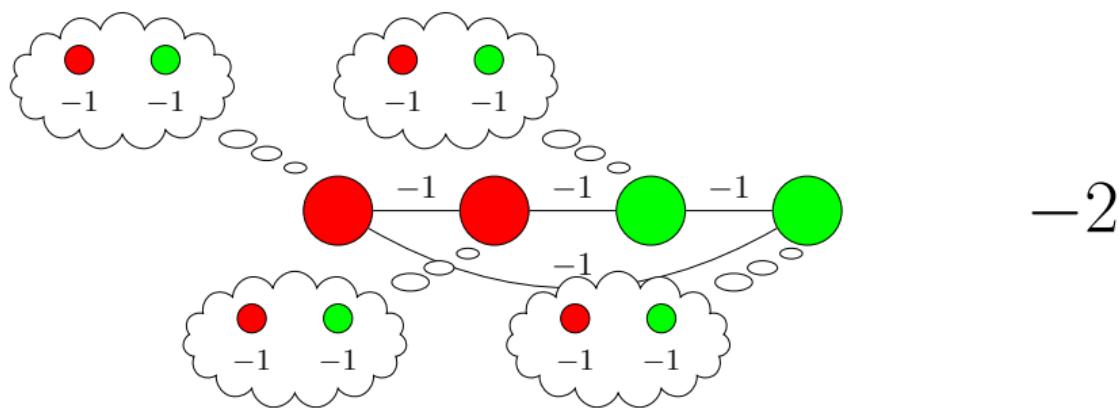
- ▶ Local minima
- ▶ Standard solutions: Random Walk, Simulated Annealing



# Centralized Local Greedy approaches

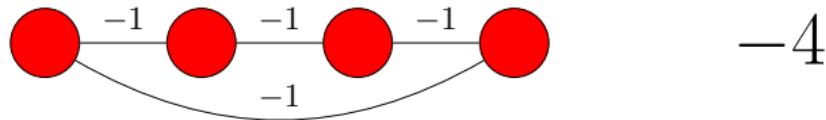
## ■ Problems

- ▶ Local minima
- ▶ Standard solutions: Random Walk, Simulated Annealing



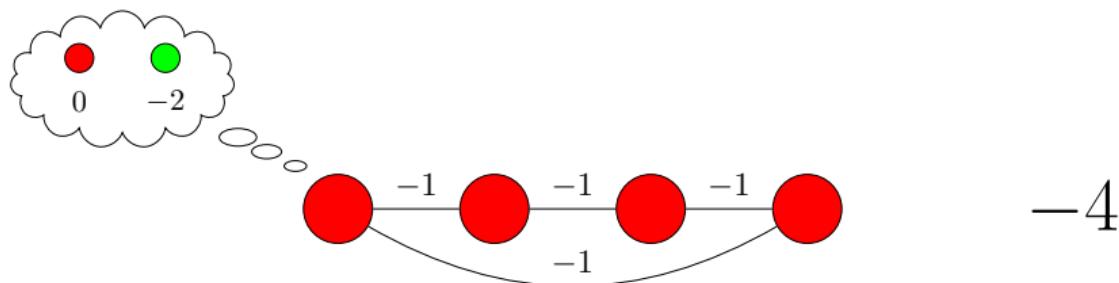
## Distributed Local Greedy approaches

- Local knowledge
- Parallel execution
  - ▶ A greedy local move might be harmful/useless
  - ▶ Need coordination



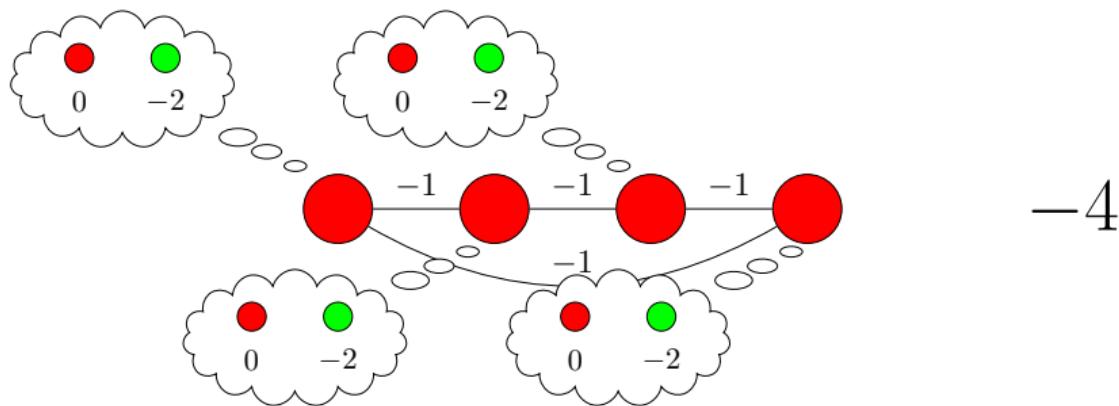
## Distributed Local Greedy approaches

- Local knowledge
- Parallel execution
  - ▶ A greedy local move might be harmful/useless
  - ▶ Need coordination



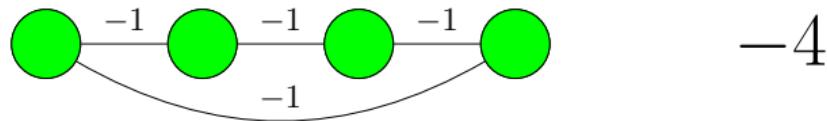
# Distributed Local Greedy approaches

- Local knowledge
- Parallel execution
  - ▶ A greedy local move might be harmful/useless
  - ▶ Need coordination



# Distributed Local Greedy approaches

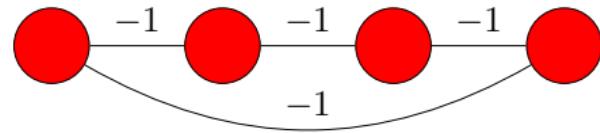
- Local knowledge
- Parallel execution
  - ▶ A greedy local move might be harmful/useless
  - ▶ Need coordination



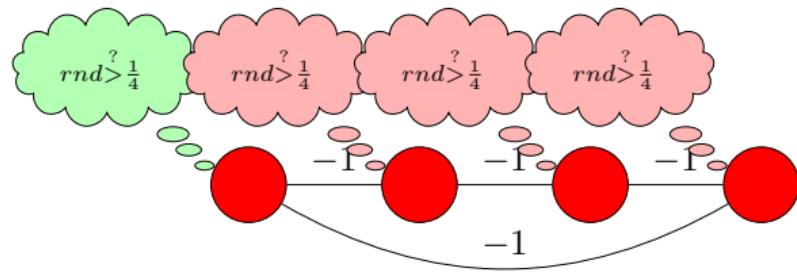
# Distributed Stochastic Search Algorithm (DSA) [ZHANG et al., 2005]

- Greedy local search with activation probability to mitigate issues with parallel executions
- DSA-1: change value of one variable at time
- Initialize agents with a random assignment and communicate values to neighbors
- Each agent:
  - ▶ Generates a random number and execute only if rnd less than activation probability
  - ▶ When executing changes value maximizing local gain
  - ▶ Communicate possible variable change to neighbors

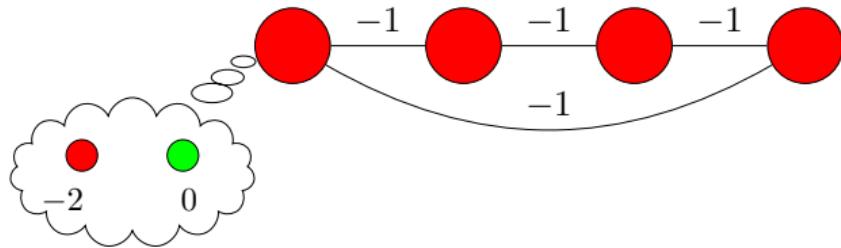
## DSA-1: Execution Example



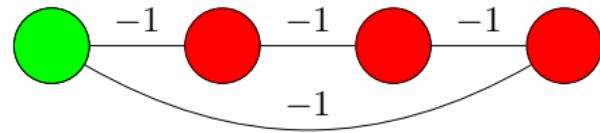
## DSA-1: Execution Example



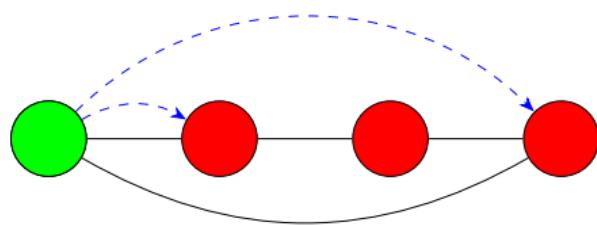
## DSA-1: Execution Example



## DSA-1: Execution Example



## DSA-1: Execution Example



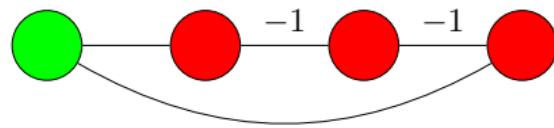
## DSA-1: Discussion

- Extremely “cheap” (computation/communication)
- Good performance in various domains
  - ▶ e.g. target tracking [FITZPATRICK and MEERTENS, 2003; ZHANG et al., 2003]
  - ▶ Shows an anytime property (not guaranteed)
  - ▶ Benchmarking technique for coordination
- Problems
  - ▶ Activation probability must be tuned [ZHANG et al., 2003]
  - ▶ No general rule, hard to characterise results across domains

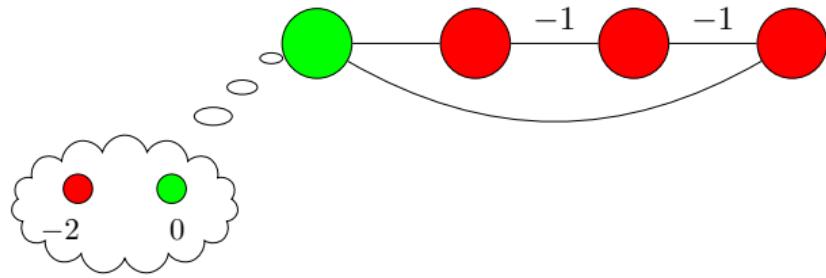
## Maximum Gain Message (MGM-1) [MAHESWARAN et al., 2004]

- Coordinate to decide who is going to move
    - ▶ Compute and exchange possible gains
    - ▶ Agent with maximum (positive) gain executes
  - Analysis
    - ▶ Empirically, similar to DSA
    - ▶ More communication (but still linear)
    - ▶ **No Threshold to set**
    - ▶ **Guaranteed to be monotonic** (Anytime behavior)

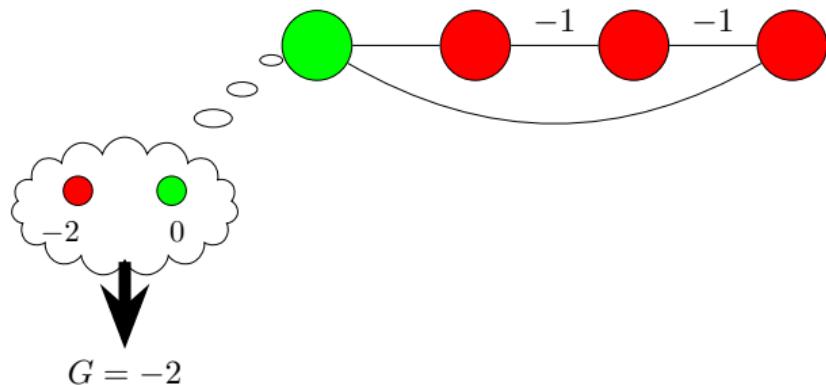
## MGM-1: Example



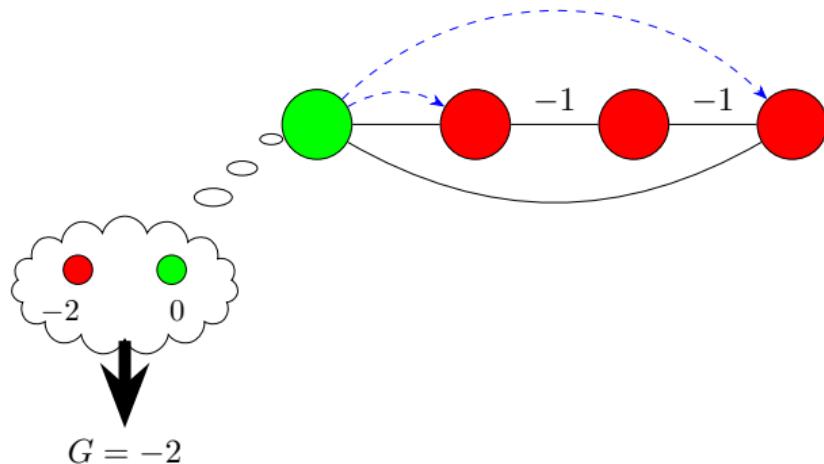
## MGM-1: Example



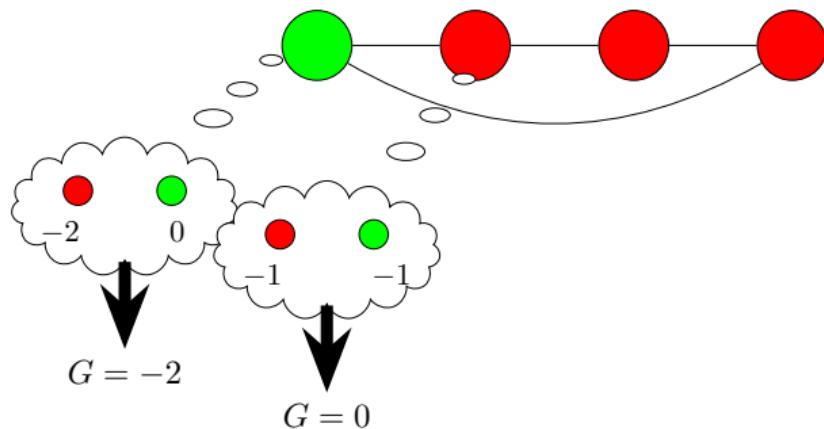
## MGM-1: Example



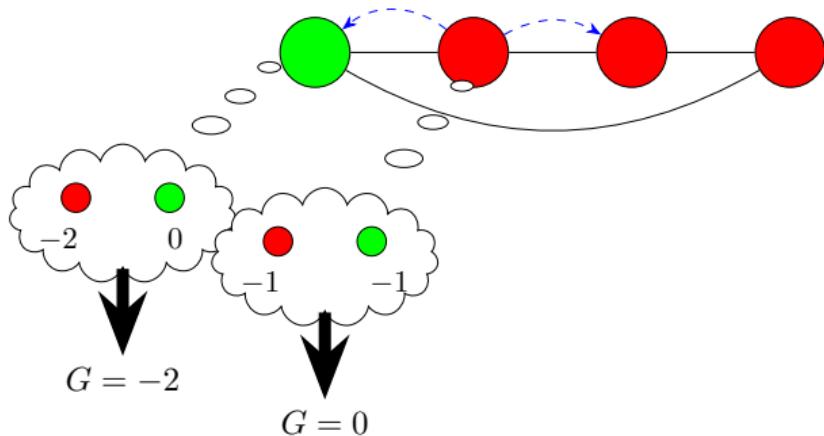
## MGM-1: Example



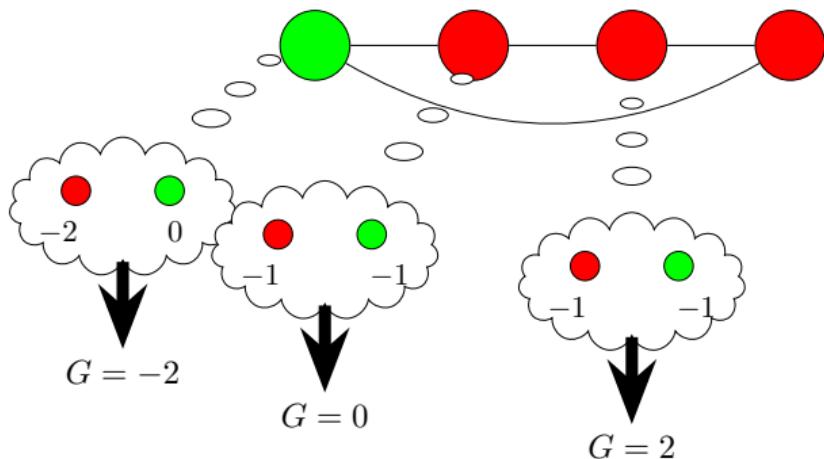
## MGM-1: Example



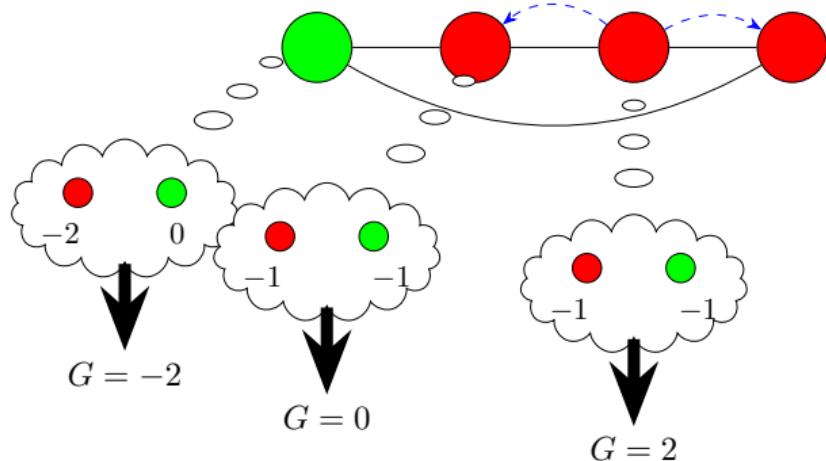
## MGM-1: Example



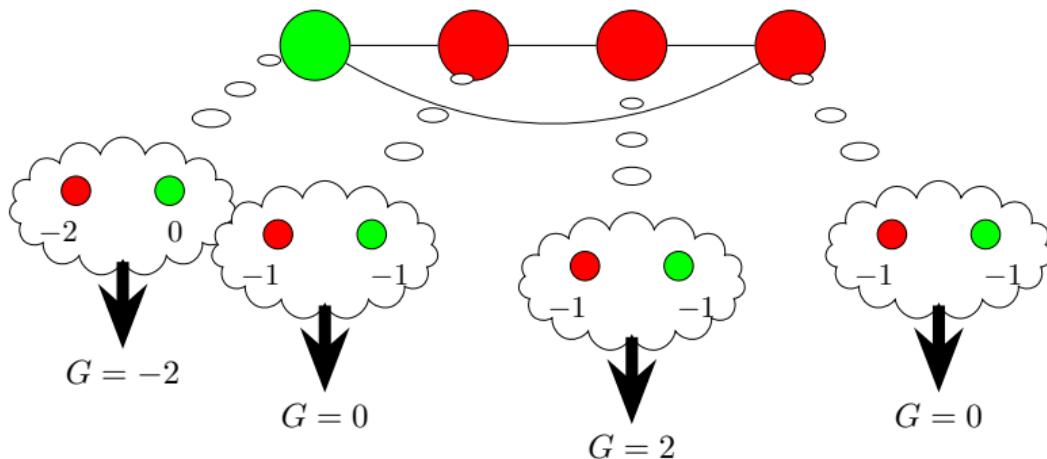
## MGM-1: Example



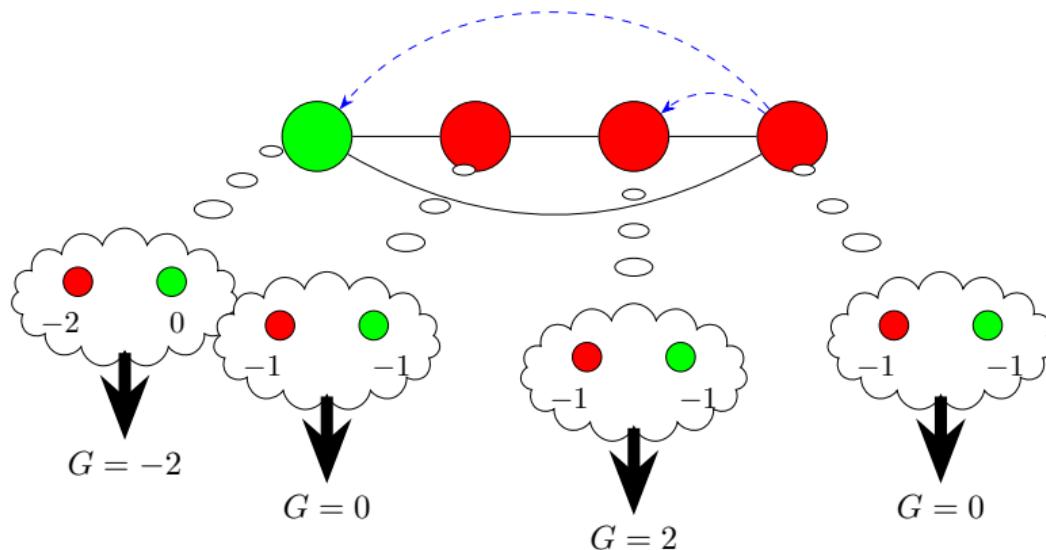
## MGM-1: Example



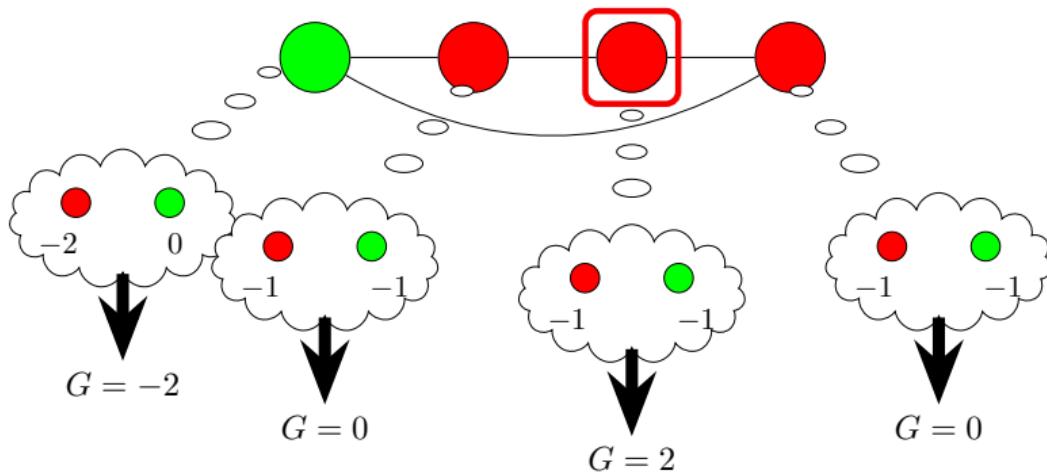
## MGM-1: Example



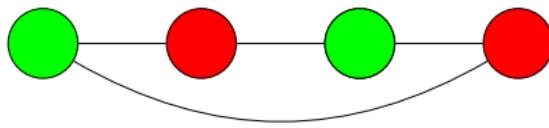
## MGM-1: Example



## MGM-1: Example



## MGM-1: Example



## To sum up on local greedy approaches

- Exchange local values for variables
    - ▶ Similar to search based methods (e.g. ADOPT)
  - Consider only local information when maximizing
    - ▶ Values of neighbors
  - Anytime behaviors
  - Could result in very bad solutions

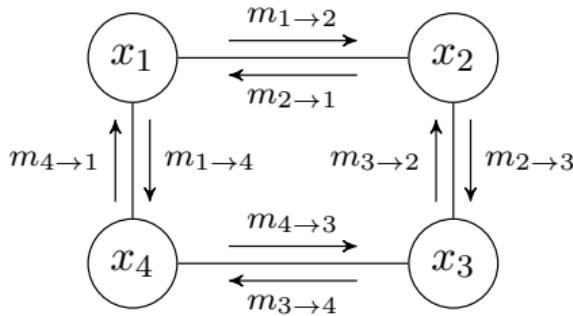
## GDL-based approaches

- Generalized Distributive Law [AJI and McELIECE, 2000]
  - ▶ Unifying framework for inference in Graphical models
  - ▶ Builds on basic mathematical properties of semi-rings
  - ▶ Widely used in Info theory, Statistical physics, Probabilistic models
- Max-sum
  - ▶ DCOP settings: maximise social welfare

$K$	“(+, 0)”	“(·, 1)”	short name
1. $A$	(+, 0)	(·, 1)	
2. $A[x]$	(+, 0)	(·, 1)	
3. $A[x, y, \dots]$	(+, 0)	(·, 1)	
4. $[0, \infty)$	(+, 0)	(·, 1)	sum-product
5. $(0, \infty]$	(min, ∞)	(·, 1)	min-product
6. $[0, \infty)$	(max, 0)	(·, 1)	max-product
7. $(-\infty, \infty]$	(min, ∞)	(+, 0)	min-sum
8. $[-\infty, \infty)$	(max, -∞)	(+, 0)	max-sum
9. $\{0, 1\}$	(OR, 0)	(AND, 1)	Boolean
10. $2^S$	( $\cup$ , $\emptyset$ )	( $\cap$ , $S$ )	
11. $\Lambda$	( $\vee$ , 0)	( $\wedge$ , 1)	
12. $\Lambda$	( $\wedge$ , 1)	( $\vee$ , 0)	

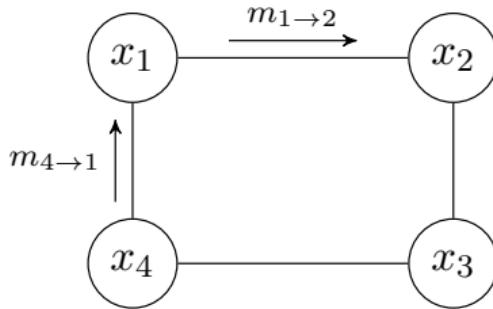
## Max-Sum

Agents iteratively computes local functions that depend only on the variable they control



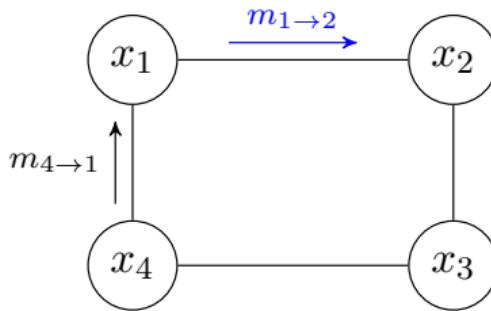
## Max-Sum

Agents iteratively computes local functions that depend only on the variable they control



## Max-Sum

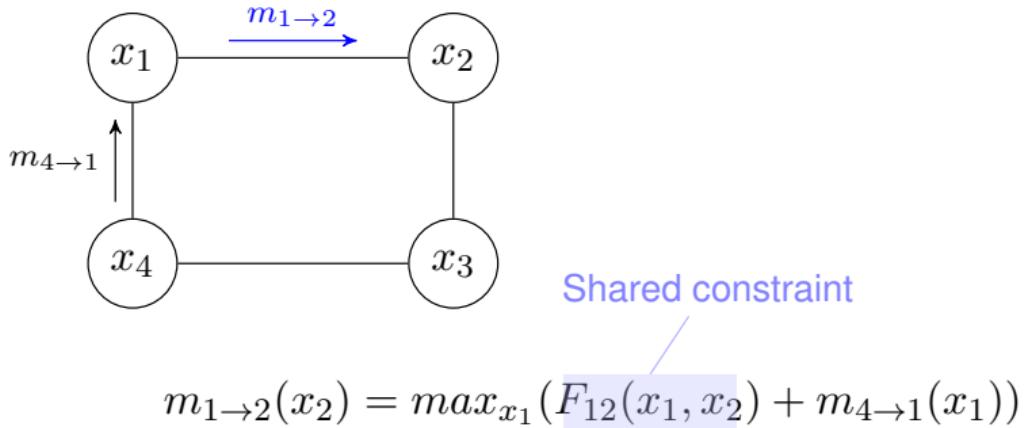
Agents iteratively computes local functions that depend only on the variable they control



$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (F_{12}(x_1, x_2) + m_{4 \rightarrow 1}(x_1))$$

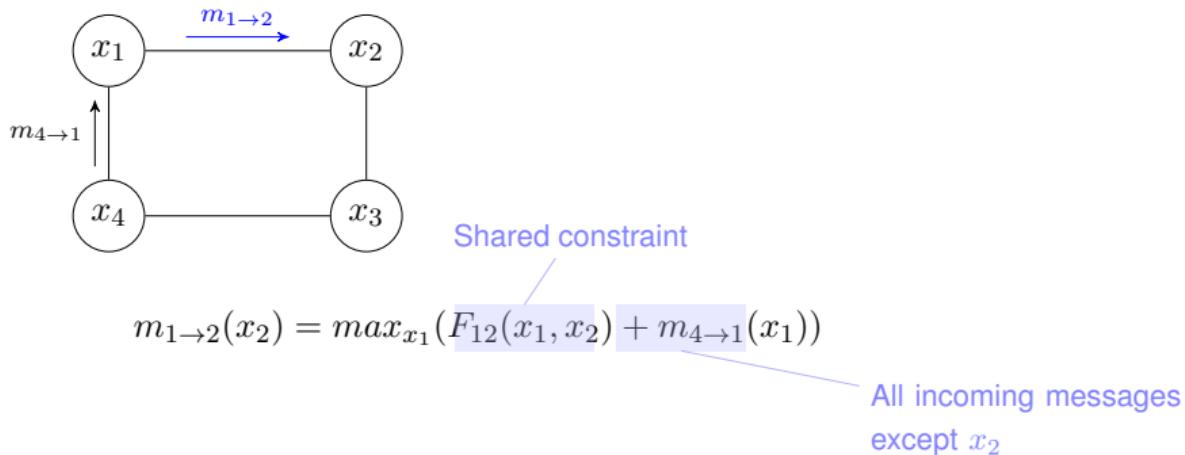
## Max-Sum

Agents iteratively computes local functions that depend only on the variable they control



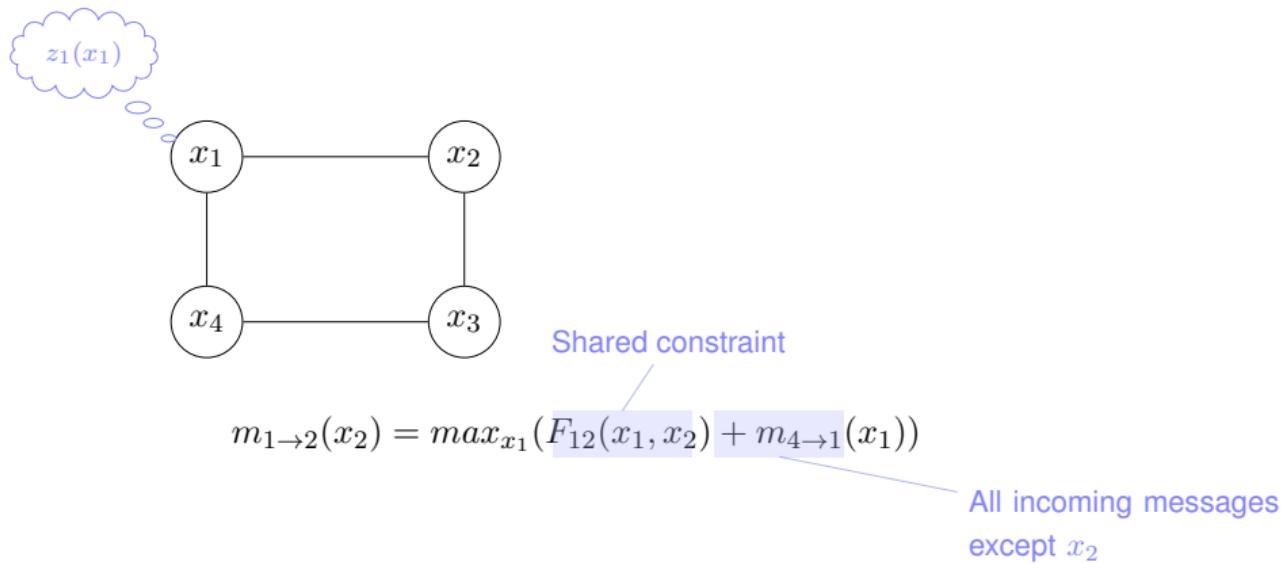
## Max-Sum

Agents iteratively computes local functions that depend only on the variable they control



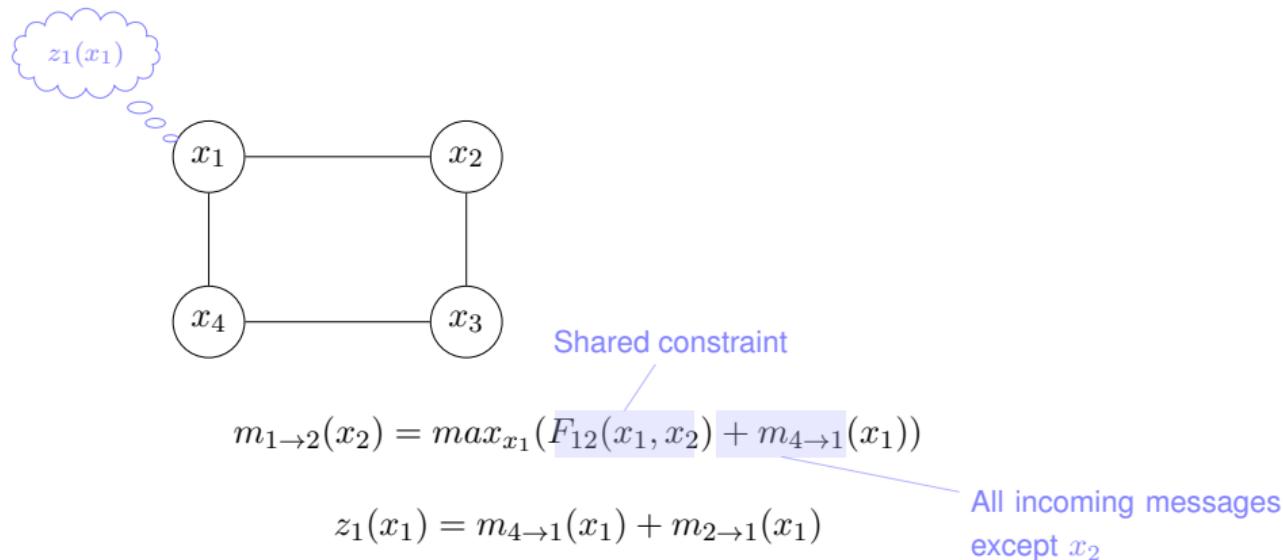
## Max-Sum

Agents iteratively computes local functions that depend only on the variable they control



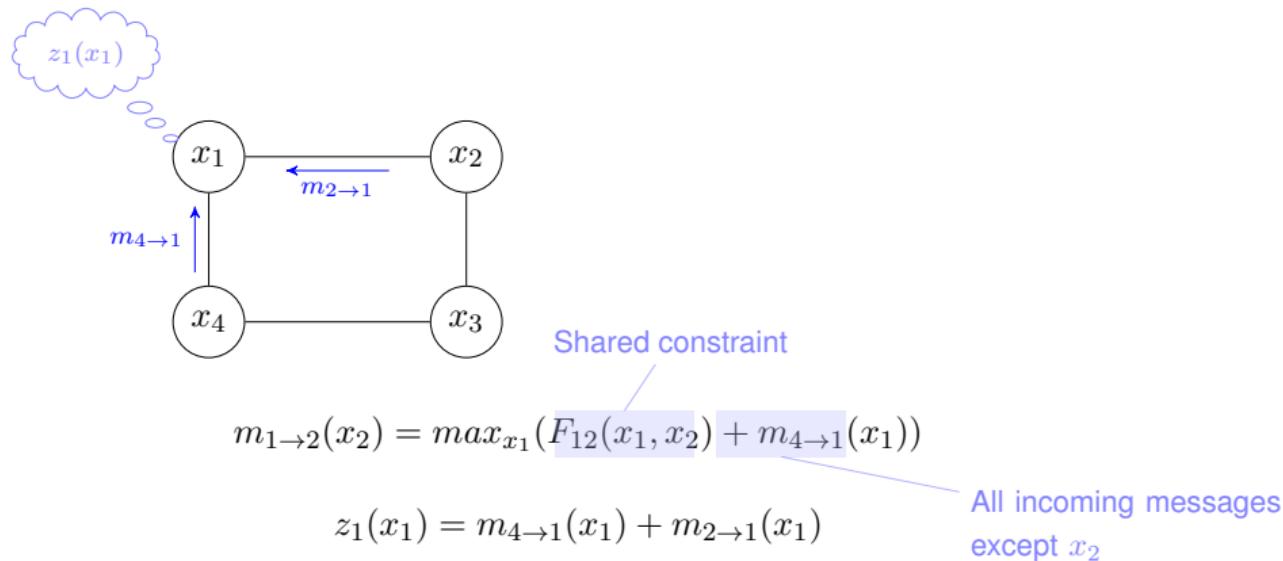
## Max-Sum

Agents iteratively computes local functions that depend only on the variable they control



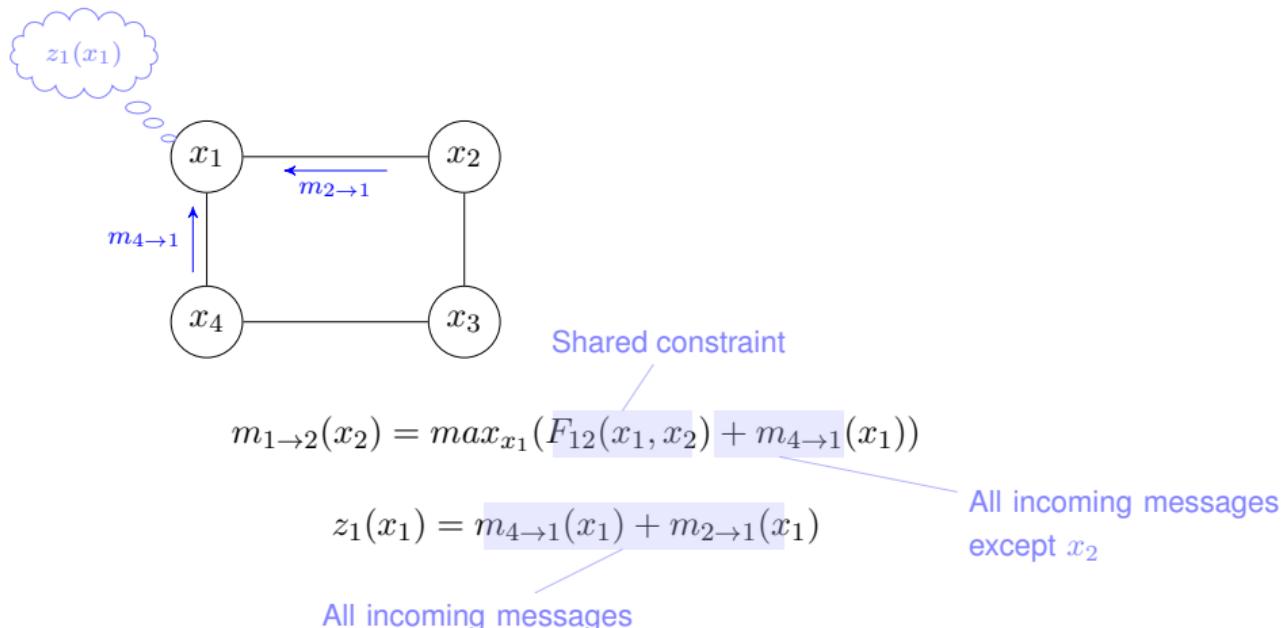
# Max-Sum

Agents iteratively computes local functions that depend only on the variable they control



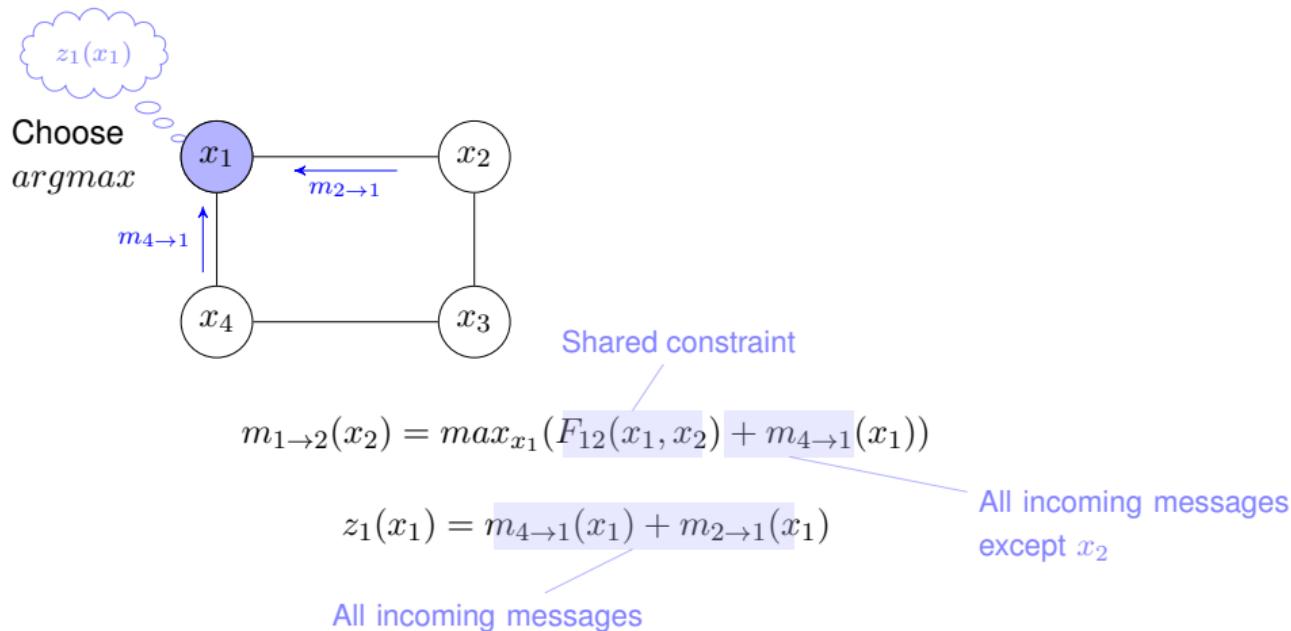
# Max-Sum

Agents iteratively computes local functions that depend only on the variable they control



# Max-Sum

Agents iteratively computes local functions that depend only on the variable they control



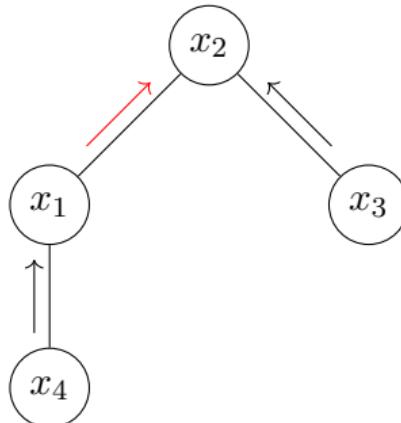
# Max-Sum on acyclic graphs

- Max-sum Optimal on acyclic graphs

- ▶ Different branches are independent
- ▶ Each agent can build a correct estimation of its contribution to the global problem ( $z$  functions)

- Message equations very similar to Util messages in DPOP

- ▶ Sum messages from children and shared constraint
  - ▶ Maximize out agent variable
  - ▶ GDL generalizes DPOP
- [VINYALS et al., 2011]

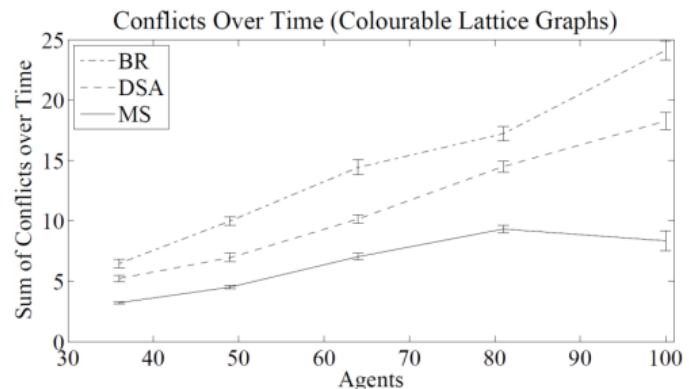
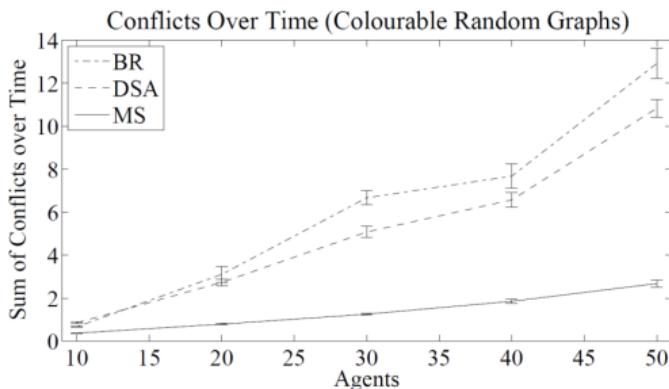


$$m_{1 \rightarrow 2}(x_2) = \max_{x_1} (F_{12}(x_1, x_2) + m_{4 \rightarrow 1}(x_1))$$

# Max-sum Performance

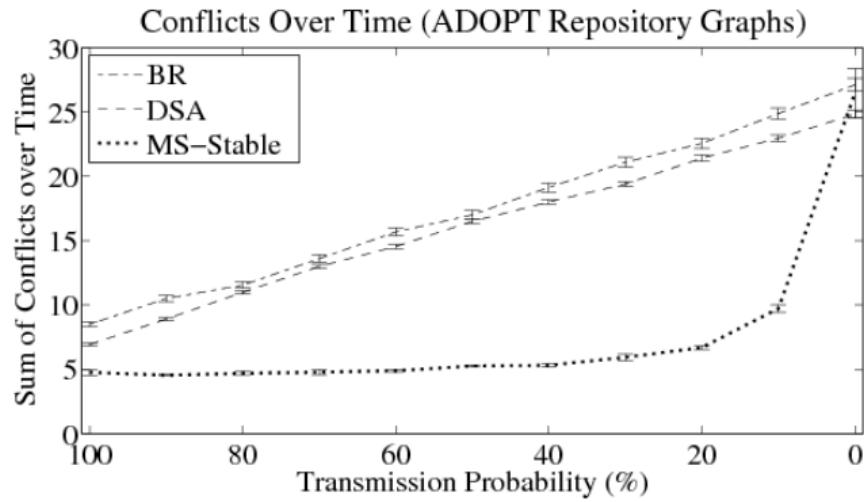
- Good performance on loopy networks [FARINELLI et al., 2008]

- ▶ When it converges very good results
    - ▶ Interesting results when only one cycle [WEISS, 2000]
  - ▶ We could remove cycle but pay an exponential price (see DPOP)



## Max-Sum for low power devices

- Low overhead
  - ▶ Msgs number/size
- Asynchronous computation
  - ▶ Agents take decisions whenever new messages arrive
- Robust to message loss



# Contents

Introduction

Complete Algorithms for DCOP

Approximate Algorithms for DCOP

Synthesis

Panorama

## Panorama

<b>Algorithm</b>	<b>Type</b>	<b>Memory</b>	<b>Messages</b>	<b>Remarks</b>
ADOPT	COP	Polynomial	Exponential	Complete
DPOP	COP	Exponential	Linear	Complete
DSA	COP	Linear	?	Not complete
MGM	COP	Linear	?	Not complete
Max-Sum	COP	Exponential	Linear on acyclic	Complete on trees

## Table: DCOP algorithms

## References

-  AJI, S.M. and R.J. McELIECE (2000). "The generalized distributive law". In: *Information Theory, IEEE Transactions on* 46.2, pp. 325–343. ISSN: 0018-9448. DOI: 10.1109/18.825794.
-  FARINELLI, A., A. ROGERS, A. PETCU, and N. R. JENNINGS (2008). "Decentralised Coordination of Low-power Embedded Devices Using the Max-sum Algorithm". In: *Proceedings of the 7th International Joint Conference on Autonomous Agents and Multiagent Systems - Volume 2*. AAMAS '08. International Foundation for Autonomous Agents and Multiagent Systems, pp. 639–646. ISBN: 978-0-9817381-1-6.
-  FITZPATRICK, Stephen and Lambert MEERTENS (2003). "Distributed Coordination through Anarchic Optimization". In: *Distributed Sensor Networks: A Multiagent Perspective*. Ed. by Victor LESSER, Charles L. ORTIZ, and Milind TAMBE. Boston, MA: Springer US, pp. 257–295. ISBN: 978-1-4615-0363-7.
-  MAHESWARAN, R.T., J.P. PEARCE, and M. TAMBE (2004). "Distributed Algorithms for DCOP: A Graphical-Game-Based Approach". In: *Proceedings of the 17th International Conference on Parallel and Distributed Computing Systems (PDCS), San Francisco, CA*, pp. 432–439.
-  MODI, P. J., W. SHEN, M. TAMBE, and M. YOKOO (2005). "ADOPT: Asynchronous Distributed Constraint Optimization with Quality Guarantees". In: *Artificial Intelligence* 161.2, pp. 149–180.
-  PETCU, Adrian and Boi FALTINGS (2005). "A scalable method for multiagent constraint optimization". In: *IJCAI International Joint Conference on Artificial Intelligence*, pp. 266–271. ISBN: 1045-0823.
-  VINYALS, Meritxell, Juan A. RODRÍGUEZ-AGUILAR, and Jesus CERQUIDES (2011). "Constructing a unifying theory of dynamic programming DCOP algorithms via the generalized distributive law". In: *Autonomous Agents and Multi-Agent Systems* 3.22, pp. 439–464. ISSN: 1387-2532. DOI: 10.1007/s10458-010-9132-7.

## References (cont.)

-  WEISS, Yair (Jan. 2000). "Correctness of Local Probability Propagation in Graphical Models with Loops". In: *Neural Comput.* 12.1, pp. 1–41. ISSN: 0899-7667. DOI: 10.1162/089976600300015880. URL: <http://dx.doi.org/10.1162/089976600300015880>.
-  YOKOO, M. (2001). *Distributed Constraint Satisfaction: Foundations of Cooperation in Multi-Agent Systems*. Springer.
-  ZHANG, W., G. WANG, Z. XING, and L. WITTENBURG (2005). "Distributed stochastic search and distributed breakout: properties, comparison and applications to constraint optimization problems in sensor networks.". In: *Journal of Artificial Intelligence Research (JAIR)* 16.1-2, pp. 55–87.
-  ZHANG, Weixiong, Guandong WANG, Zhao XING, and Lars WITTENBURG (2003). "A Comparative Study of Distributed Constraint Algorithms". In: *Distributed Sensor Networks: A Multiagent Perspective*. Ed. by Victor LESSER, Charles L. ORTIZ, and Milind TAMBE. Boston, MA: Springer US, pp. 319–338.