

Distributed Constraint Processing

An Introduction

Gauthier Picard

ONERA/DTIS
gauthier.picard@onera.fr

Contents

Introduction

- Motivating Examples

- Preliminaries

- DCOP Model

- Constraint Satisfaction Problems

- Multi-Agent Approaches to DisCSP

ABT and Extensions

- Asynchronous Algorithms for DisCSP

- ABT

- AWCS

Distributed Local Search

- Classical Centralised LS Algorithms

- Distributed Breakout Algorithm (DBA)

- Environment, Reactive rules and Agents (ERA)


Synthesis

- Panorama

- Using Distributed Problem Solving


Motivating example

Sensor networks




Motivating example

Sensor networks




Motivating example

Sensor networks




Motivating example

Sensor networks




Motivating example

Sensor networks



Motivating example

Sensor networks



x_1	x_3	x_5	Sat?
N	N	N	X
N	N	E	X
...			X
S	W	N	✓
...			X
W	W	W	X

Model the problem
as a CSP!


CSP

Constraint Satisfaction

- Variables $X = \{x_1, \dots, x_n\}$
 - Domains $D = \{D_1, \dots, D_n\}$
 - Constraints $C = \{c_1, \dots, c_m\}$
where a constraint $c_i \subseteq D_{i_1} \times D_{i_2} \times \dots \times D_{i_n}$ denotes the possible valid joint assignments for the variables $x_{i_1}, x_{i_2}, \dots, x_{i_n}$ it involves
 - **Goal:** Find an assignment to all variables that satisfies all the constraints

CSP

Constraint Satisfaction




x_1	x_3	x_5	Sat?
N	N	N	X
N	N	E	X
...			X
S	W	N	✓
...			X
W	W	W	X

Model the problem
as a CSP!

Max-CSP

Max Constraint Satisfaction



x_1	x_3	x_5	Sat?
N	N	N	X
N	N	E	X
...			X
S	W	N	✓
...			X
W	W	W	X

Model the problem
as a Max-CSP!


Max-CSP

Max Constraint Satisfaction

- Variables $X = \{x_1, \dots, x_n\}$
- Domains $D = \{D_1, \dots, D_n\}$
- Constraints $C = \{c_1, \dots, c_m\}$
where a constraint $c_i \subseteq D_{i_1} \times D_{i_2} \times \dots \times D_{i_n}$ denotes the possible valid joint assignments for the variables $x_{i_1}, x_{i_2}, \dots, x_{i_n}$ it involves
- **Goal:** Find an assignment to all variables that **satisfies a maximum number of constraints**

Max-CSP

Max Constraint Satisfaction




x_1	x_3	x_5	Sat?
N	N	N	X
N	N	E	X
...			X
S	W	N	✓
...			X
W	W	W	X

Model the problem
as a Max-CSP!

WCSP (or COP)

Constraint Optimization



x_1	x_3	x_5	Cost
N	N	N	∞
N	N	E	∞
...			∞
S	W	N	10
...			∞
W	W	W	∞


Model the problem
as a COP!

WCSP (or COP)

Constraint Optimization


- Variables $X = \{x_1, \dots, x_n\}$
- Domains $D = \{D_1, \dots, D_n\}$
- Constraints $C = \{c_1, \dots, c_m\}$
where a constraint $c_i : D_{i_1} \times D_{i_2} \times \dots \times D_{i_n} \rightarrow \mathbb{R}_+ \cup \{\infty\}$ expresses the degree of constraint violation
- **Goal:** Find an assignment to all variables that **minimizes the sum of all the constraints**

Constraint Reasoning



WCSP (or COP)

Constraint Optimization




Imagine that each sensor is an autonomous agent

How should this problem be modeled and solved in a decentralized manner?

DCOP

Distributed Constraint Optimization [MODI et al., 2005]




Imagine that each sensor is an autonomous agent

How should this problem be modeled and solved in a decentralized manner?

DCOP

Distributed Constraint Optimization [MODI et al., 2005]



DCOP


Distributed Constraint Optimization [MODI et al., 2005]

- Agents $X = \{a_1, \dots, a_l\}$
- Variables $X = \{x_1, \dots, x_n\}$
- Domains $D = \{D_1, \dots, D_n\}$
- Constraints $C = \{c_1, \dots, c_m\}$
- Mapping of variables to agents

- **Goal:** Find an assignment to all variables that minimizes the sum of all the constraints


DCOP

Distributed Constraint Optimization [MODI et al., 2005]



DCOP

Distributed Constraint Optimization [MODI et al., 2005]



- Variables are controlled by agents
- Communication model
- Local knowledge


Introduction

Multiagent Systems

- **Agent:** An entity that behaves autonomously in the pursuit of goals
- **Multi-agent system:** A system of multiple interacting agents

An agent is...

- **Autonomous:** Is of full control of itself
- **Interactive:** May communicate with other agents
- **Reactive:** Responds to changes in the environment or requests by other agents
- **Proactive:** Takes initiatives to achieve its goals



Introduction

Motivations

- Multi-agent systems are a way to model decentralised problem solving (privacy, distribution)
- Agents, having personal goals and constraints, negotiate as to reach a global equilibrium
⇒ distributed problem solving using agents

Approaches

- Classical CSP solver extensions
- Classical local search solver extensions

Cooperative Decentralized Decision Making

- Decentralised Decision Making
 - ▶ Agents have to coordinate to perform best actions
- Cooperative settings
 - ▶ Agents form a **team** → best actions for the **team**
- Why DDM in cooperative settings is important
 - ▶ Surveillance (target tracking, coverage)
 - ▶ Robotics (cooperative exploration)
 - ▶ Autonomous cars (cooperative traffic management)
 - ▶ Scheduling (meeting scheduling)
 - ▶ Rescue Operation (task assignment)


Distributed Constraint Optimisation Problems (DCOPs) for DDM

Why DCOPs for Cooperative DDM?


- Well defined problem
 - ▶ Clear formulation that captures most important aspects
 - ▶ Many solution techniques
 - ▶ Optimal: ABT, ADOPT, DPOP, ...
 - ▶ Approximate: DSA, MGM, Max-Sum, ...
- Solution techniques can handle large problems
 - ▶ compared for example to sequential decision making (MDP, POMDP)

Modeling Problems as DCOP

Target Tracking



Meeting Scheduling



■ Why decentralize


- ▶ Robustness to failure and message

■ Why decentralize


- ▶ Privacy

Target Tracking as a DCOP


- Variables → Cameras
- Domains → Camera actions
 - ▶ look left, look right
- Constraints
 - ▶ Overlapping cameras
 - ▶ Related to targets
 - ▶ Diabolik, Eva
- Maximise sum of constraints



Meeting Scheduling as a DCOP



Meeting Scheduling as a DCOP



Benchmarking problems

Motivations


- Analysis of complexity and optimality is not enough
- Need to empirically evaluate algorithms on the same problem

Graph coloring

- Simple to formalise very hard to solve
 - ▶ Well known parameters that influence complexity
 - ▶ Number of nodes, number of colors, density (number of link/number of nodes)
- Many versions of the problem
 - ▶ CSP, MaxCSP, COP

Graph Coloring


- Network of nodes
- Nodes can take on various colors
- Adjacent nodes should not have the same color
 - ▶ If it happens this is a conflict




CSP

Graph Coloring


- Network of nodes
- Nodes can take on various colors
- Adjacent nodes should not have the same color
 - ▶ If it happens this is a conflict



CSP




Yes




No

Graph Coloring - MaxCSP


- Optimization Problem
- Natural extension of CSP
- Minimise number of conflicts



0




-1




-4

Weighted Graph Coloring - COP


- Optimization Problem
- Conflicts have a weight
- Maximise the sum of weights of violated constraints



COP



-2



-1

Constraint Satisfaction Problems [DECHTER, 2003]

Definition (CSP)

A CSP is a triplet $\langle X, D, C \rangle$ such as:

- $X = \{x_1, \dots, x_n\}$ is the set of *variables* to instantiate.
- $D = \{D_1, \dots, D_m\}$ is the set of *domains*. Each variable x_i is related to a domain of value.
- $C = \{c_1, \dots, c_k\}$ is the set of *constraints*, which are relations between some variables from X that constrain the values the variables can be simultaneously instantiated to.

Definition (Solution to a CSP)

A solution to a CSP is a complete assignment of values from D to variables from X such that every constraint in C is satisfied.

Issues in CSP

Classical CSPs

- Constraint satisfaction is NP-complete in general
- Constraints are generally expressed as binary constraints
- The topology of a constraint-based problem can be represented by a *constraint network*, in which vertexes represent variables and edges represent binary constraints between variables

Extensions

- Distribution : variables, constraints
 - ▶ ex.: constraint c_i belongs to stakeholder j , $\phi(c_i) = j$ (or $belongs(c_i, j)$)
- Dynamics : adding removing variables and/or constraints at runtime

Multi-Agent Approaches to CSP

- **Complete and asynchronous solvers** for combinatorial problems, within the DisCSP framework, such as Asynchronous Backtracking (ABT) or Asynchronous Weak-Commitment Search (AWCS)
- **Distributed local search** methods, such as Distributed Breakout Algorithm (DBA) or Environment, Reactive rules and Agents (ERA) approach

Asynchronous Algorithms for DisCSP

Idea

- Inspired by classical centralised algorithms to solve CSP
- Each agent is responsible for assigning one (or several) variables
- Agents propose values to some other agents (depending on the organisation i.e. constraint network)

Main algorithm: Asynchronous backtracking (ABT) [Yokoo, 2001]

- Agents will perform a distributed version of the backtracking procedure
- ABT is complete
- Extensions exist to handle dynamics

Definition (DisCSP or DCSP)

A *DisCSP* (or *DCSP*) is a 5-uplet $\langle A, X, D, C, \phi \rangle$ where $\langle X, D, C \rangle$ is a CSP, A is a set of agents and $\phi : X \mapsto A$ is a function assigning variables from X to agents from A .

Centralised Backtracking

```

 $i \leftarrow 0$ 
 $D'_i \leftarrow D_i$ 
while  $0 \leq i < n$  do
     $x_i \leftarrow \text{null}$ 
     $ok? \leftarrow \text{false}$ 
    while not  $ok?$  and  $D'_i$  not empty do
         $a \leftarrow \text{a value from } D'_i$ 
        remove  $a$  from  $D'_i$ 
        if  $a$  is in conflict with  $\{x_0, \dots, x_{i-1}\}$  then
             $x_i \leftarrow a$ 
             $ok? \leftarrow \text{true}$ 
        end
    end
    if  $x_i$  is null then backtrack
         $i \leftarrow i - 1$ 
    else
         $i \leftarrow i + 1$ 
         $D'_i \leftarrow D_i$ 
    end
end

```


Algorithm 1: A classical centralised backtracking search method

Asynchronous Backtracking (ABT) [ibid.]

- First complete asynchronous algorithm for DisCSP solving
- Asynchronous:
 - ▶ All agents active, take a value and inform
 - ▶ No agent has to wait for other agents
- Total order among agents: to avoid cycles
 - ▶ $i < j < k$ means that: i more priority than j , j more priority than k
- Constraints are directed, following total order
- ABT plays in asynchronous distributed context the same role as backtracking in centralized


ABT: Directed Constraints

- Directed: from higher to lower priority agents
- Higher priority agent (j) informs the lower one (k) of its assignment
- Lower priority agent (k) evaluates the constraint with its own assignment
 - ▶ If permitted, no action
 - ▶ else it looks for a value consistent with j
 - ▶ If it exists, k takes that value
 - ▶ else, the agent view of k is a nogood, backtrack



ABT: Directed Constraints

- Directed: from higher to lower priority agents
- Higher priority agent (j) informs the lower one (k) of its assignment
- Lower priority agent (k) evaluates the constraint with its own assignment
 - ▶ If permitted, no action
 - ▶ else it looks for a value consistent with j
 - ▶ If it exists, k takes that value
 - ▶ else, the agent view of k is a nogood, backtrack



generates nogoods: eliminate values of k

ABT: Nogoods

Definition (Nogood)

Conjunction of (variable, value) pairs of higher priority agents, that removes a value of the current one

Example

- $x \neq y, d_x = d_y = \{a, b\}, x$ higher than y
- When $[x \leftarrow a]$ arrives to y , this agent generates the nogood $[x = a \Rightarrow y \neq a]$ that removes value a of d_y
- If x changes value, when $[x \leftarrow b]$ arrives to y , the nogood $[x = a \Rightarrow y \neq a]$ is eliminated, value a is again available and a new nogood removing b is generated

ABT: Nogood Resolution

- When all values of variable y are removed, the conjunction of the left-hand sides of its nogoods is also a nogood
- Resolution:** the process of generating the new nogood

Example

- $x \neq y, z \neq y, d_x = d_y = d_z = \{a, b\}, x, z$ higher than y

$$x = a \Rightarrow y \neq a$$

$x = a \wedge z = b$ is a nogood


$$z = b \Rightarrow y \neq b$$

$x = a \Rightarrow z \neq b$ (assuming x higher than z)

How ABT works


- **ABT agents:** asynchronous action, spontaneous assignment
- **Assignment:** j takes value a , j informs lower priority agents
- **Backtrack:** k has no consistent values with high priority agents, k resolves nogoods and sends a backtrack message
- **New links:** j receives a nogood mentioning i , unconnected with j ; j asks i to set up a link
- **Stop:** “no solution” detected by an agent, stop
- **Solution:** when agents are silent for a while (quiescence), every constraint is satisfied
→ solution; detected by specialized algorithms

ABT: Messages




ABT: Messages

- $Ok?(i \rightarrow k, a)$:
 - ▶ i informs k that it takes value a



ABT: Messages

- $Ok?(i \rightarrow k, a)$:
 - ▶ i informs k that it takes value a
- $Ngd(k \rightarrow j, i = a \Rightarrow j \neq b)$
 - ▶ all k values are forbidden
 - ▶ k requests j to backtrack
 - ▶ k forgets j value
 - ▶ k takes some value
 - ▶ j may detect obsolescence



ABT: Messages

- *Ok?*($i \rightarrow k, a$):


- ▶ i informs k that it takes value a

- *Ngd*($k \rightarrow j, i = a \Rightarrow j \neq b$)

- ▶ all k values are forbidden
 - ▶ k requests j to backtrack
 - ▶ k forgets j value
 - ▶ k takes some value
 - ▶ j may detect obsolescence


- *Addl*($j \rightarrow i$):

- ▶ set a link from i to j , to know i value



ABT: Messages

- *Ok?*($i \rightarrow k, a$):
 - ▶ i informs k that it takes value a
- *Ngd*($k \rightarrow j, i = a \Rightarrow j \neq b$)
 - ▶ all k values are forbidden
 - ▶ k requests j to backtrack
 - ▶ k forgets j value
 - ▶ k takes some value
 - ▶ j may detect obsolescence
- *Addl*($j \rightarrow i$):
 - ▶ set a link from i to j , to know i value
- Stop:
 - ▶ there is no solution



ABT Procedures

```

when received (ok?, (xj, dj)) do — (i)
  revise agent_view;
  check_agent_view;
end do;




when received (nogood, xj, nogood) do — (ii)
  record nogood as a new constraint;
  when nogood contains an agent xk that is not its neighbor
    do request xk to add xi as a neighbor,
       and add xk to its neighbors; end do;
  old_value  $\leftarrow$  current_value; check_agent_view;
  when old_value = current_value do
    send (ok?, (xj, current_value)) to xj; end do; end do;

procedure check_agent_view
  when agent_view and current_value are not consistent do
    if no value in Di is consistent with agent_view then backtrack;
    else select d  $\in$  Di where agent.view and d are consistent;
      current_value  $\leftarrow$  d;
      send (ok?, (xi, d)) to neighbors; end if; end do;

procedure backtrack
  generate a nogood V — (iii)
  when V is an empty nogood do
    broadcast to other agents that there is no solution,
    terminate this algorithm; end do;
  select (xj, dj) where xj has the lowest priority in a nogood;
  send (nogood, xj, V) to xj;
  remove (xj, dj) from agent.view;
  check_agent_view;

```

Algorithm 2: ABT Procedures



ABT: Correctness and Completeness

■ Correctness

- ▶ silent network \Leftrightarrow all constraints are satisfied

■ Completeness

- ▶ ABT performs an exhaustive traversal of the search space
- ▶ Parts not searched: those eliminated by nogoods
- ▶ Nogoods are legal: logical consequences of constraints
- ▶ Therefore, either there is no solution \Rightarrow ABT generates the empty nogood, or it finds a solution if exists

ABT: Remarks

- Fixed ordered organisation
 - ▶ Agents only communicate with agents with lower priority for *ok*?
 - ▶ Agents only communicate with the agent with direct higher priority for *nogood*
- No termination procedure is given (but it is easily implemented using Dijkstra's tokens)
- Really distributable
- What if x_0 disappears?...

Extensions and Filiation

- **Changing ordering** in every conflict with AWCS [ibid.]
- Satisfaction → **Optimisation** with ADOPT (Asynchronous B&B) [MODI et al., 2005] or APO [MAILLER and LESSER, 2006]
- **Adding new agents** at runtime in DynAPO [MAILLER, 2005]

Asynchronous Weak-Commitment Search (AWCS) [Yokoo, 2001]


```

procedure check_agent_view
  when agent.view and current.value are not consistent do
    if no value in  $D_i$  is consistent with agent.view then backtrack;
    else select  $d \in D_i$  where agent.view and  $d$  are consistent
      and  $d$  minimizes the number of constraint violations
      with lower priority agents; — (i)
      current.value  $\leftarrow d$ ;
      send (ok?,  $(x_i, d, current.priority)$ ) to neighbors;
  end if; end do;
```


```

procedure backtrack
  generate a nogood  $V$ ;
  when  $V$  is an empty nogood do
    broadcast to other agents that there is no solution,
    terminate this algorithm; end do;
  when  $V$  is a new nogood do — (ii)
    send  $V$  to the agents in the nogood;
    current.priority  $\leftarrow 1 + p_{max}$ ,
    where  $p_{max}$  is the maximal priority value of neighbors;
    select  $d \in D_i$  where agent.view and  $d$  are consistent,
    and  $d$  minimizes the number of constraint violations
    with lower priority agents;
    current.value  $\leftarrow d$ ;
    send (ok?,  $(x_i, d, current.priority)$ ) to neighbors; end do;
```


Algorithm 3: AWCS Procedures




(a)



(b)



(c)



(d)

Distributed Local Search Approaches

Local Search (LS)

- LS algorithms explore the search space from state to state
- Always tend to improve the current state of the system
- Can naturally handle dynamics (adding constraints, changing values)
- Time efficient
- Not complete and require some subtle parameter tuning

```
choose an initial assignment  $s(0)$ 
while  $s(t)$  not terminal do
    select an acceptable move  $m(t)$  to another assignment
    apply move  $m(t)$  to reach  $s(t + 1)$ 
     $t := t + 1$ 
end
```

Algorithm 4: A generic centralised local search algorithm

Classical Centralised LS Algorithms

Common points

- Initial point (ex: randomly chosen)
- Termination criterion (ex: limit time, δ improvement)
- Acceptable move (ex: $+ \epsilon$)

Famous LS Methods

- **Tabu search** [GLOVER and LAGUNA, 1997]
- **Simulated annealing** [KIRKPATRICK et al., 1983]
- **Iterative Breakout method** [MORRIS, 1993]

Distributed Breakout Algorithm (DBA)

```
wait_ok? mode — (i)
when received (ok?,  $x_j$ ,  $d_j$ ) do
```

add (x_j, d_j) to *agent_view*;
when received ok? messages from all neighbors do
 send_improve;
 goto wait_improve mode; end do;
 goto wait_ok mode; end do;

```
procedure send_improve
```

current_eval \leftarrow evaluation value of *current_value*;
my_improve \leftarrow possible maximal improvement;
new_value \leftarrow the value which gives the maximal improvement;
 send (*improve*, x_i , *my_improve*, *current_eval*) to neighbors;

```
wait_improve? mode — (ii)
```


```
when received (improve,  $x_j$ , improve, eval) do
```

record this message;
when received improve? messages from all neighbors do
 send_ok; clear *agent_view*;
 goto wait_ok mode; end do;
 goto wait_improve mode; end do;


```
procedure send_ok
```

when its improvement is largest among neighbors do
current_value \leftarrow *new_value*; end do;
when it is in a quasi-local-minimum do
 increase the weights of constraint violations; end do;
 send (ok?, x_i , *current_value*) to neighbors;


Algorithm 5: DBA Message Handler




(a)



(b)



(c)



(d)

Distributed Breakout Algorithm (DBA) (cont.)

Principles of DBA [Yokoo, 2001]

- Distribution difficulties:
 - (i) if two neighbouring agents concurrently change their value, the *system may oscillate*
 - (ii) detecting the fact that the whole system is trapped in local minimum requires the agents to *globally exchange data*
- DBA answers:
 - (i) for a given neighbourhood, only the agent that can *maximally improve* the evaluation value is given the right to change its value
 - (ii) agents only detects *quasi-local-minimum*, which is a weaker local-minimum that can be detected only by local interactions

Distributed Breakout Algorithm (DBA) (cont.)

Remarks

- Distributed version of the iterative breakout algorithm
- Two-mode behaviour alternating between exchange of potential improvement and exchange of assignments
- ✓ There is no order over the agents society → neighbourhoods
- The system halts if a solution is found or if the weight of constraints have reached a *predefined upper bound*
 - the **only** difficult parameter to set
- ✗ DBA is not complete
- ✓ DBA is able to detect the termination or a global solution only by reasoning on local data.

Environment, Reactive rules and Agents (ERA) [LIU et al., 2002]

Components

- A discrete grid **environment**, that is used as a communication medium
- **Agents** that evolves in some regions of the grid (their domain)
 - ▶ Agents move *synchronously*
 - ▶ Agents cannot move in the domain of other agents, but can mark it with the number of potential conflicts
 - ▶ These marks represents therefore the number of violated constraints if an agent chooses the marked cell
- **Rules (moves)** that agent follow to reach an equilibrium
 - ▶ 3 possible actions
 - ▶ *least-move*: the next cell is the one with minimum cost
 - ▶ *better-move*: the next cell is randomly chosen and if it has less conflicts than the actual one the agent moves else the agent rests
 - ▶ *random-move*: the next cell is randomly chosen
 - ▶ A decision consists in a random Monte-Carlo choice of the action to perform

Environment, Reactive rules and Agents (ERA) [ibid.] (cont.)


```

 $t \leftarrow 0$ 
initialise the grid to 0 violation in each cell; foreach agent  $i$  do
| randomly move to a cell of row  $i$ 
end
while  $t < t_{max}$  and no solution do
| foreach agent  $i$  do
| | select a move behaviour
| | compute new position
| | decrease markers in all cells with past violations
| | increase markers in all cells with new violations
| end
|  $t \leftarrow t + 1$ 
end


```

Algorithm 6: ERA Outline

X_1	1	(2)	1	1	0	0
X_2	0	1	0	0		
X_3	0	1	1	(1)	1	



X_1	1	2	1	1	(0)	0
X_2	0	(0)	0	0		
X_3	0	0	0	(0)	1	



X_1	1	1	(1)	1	0	0
X_2	0	(0)	1	0		
X_3	0	0	1	(1)	1	

X_1	1	(2)	1	1	0	0
X_2	0	1	0	0		
X_3	0	1	1	(1)	1	



X_1	(1)	2	1	1	0	0
X_2	1	(0)	0	0		
X_3	1	1	1	(1)	1	

Environment, Reactive rules and Agents (ERA) [ibid.] (cont.)

Remarks

- The environment is the communication medium
 - ✓ There is no asynchronous mechanisms and message handling
 - ✗ Synchronisation point: high synchronous solving process with no benefit from distribution, in case of high connected constraint networks
- ✓ ERA quickly finds assignments close to the solution → repairing issues
- ✗ Redundant usage of random choices: non-guided method, close to random walk, and non complete
- ✗ Termination: ERA requires a time limit (t_{max}) (problem-dependant)

Panorama

Algorithm	Type	Memory	Messages	Remarks
ABT	CSP	Exponential	–	Complete, Static ordering
AWCS	CSP	Exponential	–	Complete (only with exponential space), Reordering, fast
DBA	Max-CSP	Linear	Bounded	Incomplete, Fast
ERA	Max-CSP	Polynomial	n/a	Incomplete, randomness

Table: DCSP and DCOP algorithms

Using Distributed Problem Solving

Problem and Environment Characteristics

- Geographic distribution
 - ▶ ex: agents are physically distributed, and solving the whole problem is not possible in a centralised manner
- Constraint network topology
 - ▶ ex: bounded vertex degrees or large constraint graph diameter
- Knowledge encapsulation
 - ▶ ex: *privacy* preserving, limited knowledge
- Dynamics
 - ▶ ex: rather than solving the whole problem again, only repair sub-problems

Some Applications

- Internet of things
- Scheduling
- Resource allocation, Manufacturing control

References

-  DECHTER, R. (2003). *Constraint Processing*. Morgan Kaufmann.
-  GLOVER, F. and M. LAGUNA (1997). *Tabu Search*. Kluwer.
-  KIRKPATRICK, S., C. GELLAT, and M. VECCHI (1983). "Optimization by Simulated Annealing". In: *Science* 220.4598, pp. 671–680.
-  LIU, J., H. JING, and Y. Y. TANG (2002). "Multi-agent Oriented Constraint Satisfaction". In: *Artificial Intelligence* 136.1, pp. 101–144.
-  MAILLER, R. (2005). "Comparing two approaches to dynamic, distributed constraint satisfaction". In: *Proceedings of the Fourth International Joint Conference on Autonomous Agents and Multiagent Systems (AAMAS'05)*. ACM Press, pp. 1049–1056.
-  MAILLER, R. and V. R. LESSER (2006). "Asynchronous Partial Overlay: A New Algorithm for Solving Distributed Constraint Satisfaction Problems". In: *Journal of Artificial Intelligence Research* 25, pp. 529–576.
-  MODI, P. J., W. SHEN, M. TAMBE, and M. YOKOO (2005). "ADOPT: Asynchronous Distributed Constraint Optimization with Quality Guarantees". In: *Artificial Intelligence* 161.2, pp. 149–180.
-  MORRIS, P. (1993). "The Breakout Method for Escaping from Local Minima". In: *AAAI*, pp. 40–45.
-  YOKOO, M. (2001). *Distributed Constraint Satisfaction: Foundations of Cooperation in Multi-Agent Systems*. Springer.